

Student's name and surname: Grzegorz Koszczał

ID: 175405

Cycle of studies: postgraduate

Mode of study: Full-time studies

Field of study: Informatics

Specialization: Distributed Applications and Internet Services

MASTER'S THESIS

Title of thesis: Research on comparison of various methods for energy measurement for computational systems

Title of thesis (in Polish): Badanie porównawcze różnych metod pomiaru zużycia energii systemów obliczeniowych

Supervisor: dr hab. inż. Paweł Czarnul

**DECLARATION regarding the diploma thesis titled:
Research on comparison of various methods for energy measurement for
computational systems**

First name and surname of student: Grzegorz Koszczał

Date and place of birth: 11.02.1999, Olsztyn

ID: 175405

Faculty: Faculty of Electronics, Telecommunications and Informatics

Field of study: informatics

Cycle of studies: postgraduate

Mode of study: Full-time studies

Type of the diploma thesis: master's thesis

Aware of criminal liability for violations of the Act of 4th February 1994 on Copyright and Related Rights (Journal of Laws 2021, item 1062 with later amendments) and disciplinary actions set out in the Act of 20th July 2018 on the Law on Higher Education and Science (Journal of Laws 2022 item 574 with later amendments),¹ as well as civil liability, I hereby declare that the submitted diploma thesis is my own work.

This diploma thesis has never before been the basis of an official procedure associated with the awarding of a professional title.

All the information contained in the above diploma thesis which is derived from written and electronic sources is documented in a list of relevant literature in accordance with art. 34 of the Copyright and Related Rights Act.

15.08.2023, Grzegorz Koszczał

Date and signature of the student or authentication on the university portal Moja PG

**) The document was drawn up in the IT system, on the basis of paragraph 15 clause 3b of the Decree of the Ministry of Science and Higher Education of 12 May 2020, amending the decree concerning university studies (Journal of Laws of 2020, item 853). No signature or stamp required.*

¹ The Act of 20th July 2018 on the Law on Higher Education and Science:

Art. 312, section 3. Should a student be suspected of committing an act referred to in Art. 287 section 2 items 1–5, the rector shall forthwith order an enquiry.

Art. 312, section 4. If the evidence collected during an enquiry confirms that the act referred to in section 5 has been committed, the rector shall suspend the procedure for the awarding of a professional title pending a judgement of the disciplinary committee and submit formal notification on suspicion of committing a criminal offence.

ABSTRACT

The goal of the project is to conduct experiments and compare various methods of power draw measurement of parallel applications. The used measurement tools are both software (Intel RAPL, NVIDIA NVML) and hardware (Yokogawa WT310E). Moreover, the conditions of implementing mentioned solutions are evaluated, as well as their limitations for newest CPUs and GPUs.

Keywords: Informatics, Parallel Programming, High-Performance Computing Systems, Green Computing, NAS Parallel Benchmark, Linux Perf, Intel RAPL, NVML, Internal Power Sensors, External Power Meters.

Field of science and technology in accordance with OECD requirements:

Electrical engineering, Electronic engineering, Information engineering

CONTENTS

List of important symbols and abbreviations	3
1. Introduction	5
1.1. Purpose and research question	6
1.2. Scope and limitation	6
1.3. Project requirements	6
2. Theoretical background	7
2.1. Measurement software	7
2.1.1. Intel RAPL	7
2.1.2. NVIDIA NVML	7
2.2. Measurement hardware	8
2.2.1. Yokogawa WT310E	8
2.3. Benchmark applications	9
2.3.1. NPB for CPU, C++ with OMP	10
2.3.2. NPB for CPU, Fortran with MPI	11
2.3.3. NPB for GPU, CUDA	11
2.3.4. Custom Deep Neural Networks model	12
2.4. Related works	13
2.4.1. 'A Comparative Study of Methods for Measurement of Energy of Computing'	13
2.4.2. 'Verified Instruction-Level Energy Consumption Measurement for NVIDIA GPUs'	15
2.4.3. 'Measuring GPU Power with the K20 Built-in Sensor'	17
3. Definition of experiments	20
3.1. Choosing the correct configurations	20
3.2. Preliminary tests	21
3.3. Overview on all final configurations	24
3.3.1. Servers	24
3.3.2. Devices	25
3.3.3. Implementations	25
3.3.4. Benchmarks	26
3.3.5. Configurations	27
4. Experiments	30
4.1. Proposed workflow and methodology	30
4.2. Main tests	31
4.2.1. Overview on the scheduler script	31
4.2.2. Main automation function	32
4.2.3. Threads pinning and kernels execution	35
4.2.4. GPUs and threads management	36
4.2.5. Measurements with Yokotool software	38
4.2.6. Measurements with Linux Perf software	39
4.2.7. Measurements with NVML handling function	40

4.2.8. Termination of benchmarks in Hybrid configuration	41
4.2.9. Cleanup of measurements daemons	41
4.3. Analysis of the results and discussion	41
5. Summary and future work	89
5.1. Summary	89
5.2. Future work	90
References	92
List of figures	96
List of tables	99

LIST OF IMPORTANT SYMBOLS AND ABBREVIATIONS

OS	–	Operating System
CPU	–	Central Processing Unit
GPU	–	Graphics Processing Unit
PSU	–	Power Supply Unit
PCI-E	–	Peripheral Component Interconnect Express
RAPL	–	Running Average Power Limit
NVML	–	NVIDIA Management Library
NASA	–	National Aeronautics and Space Administration
NAS	–	NASA Advanced Supercomputing
NPB	–	NAS Parallel Benchmarks
ICT	–	Information and Communication Technology
MSR	–	Model-Specific Register
RAM	–	Random-Access Memory
DRAM	–	Dynamic Random-Access Memory
VRAM	–	Video Random-Access Memory
API	–	Application Programming Interface
SDK	–	Software Development Kit
SMA	–	Sampling Monitoring Approach
MTSM	–	Multi-Threaded Synchronized Monitoring
PAPI	–	Performance Application Programming Interface
CUDA	–	Compute Unified Device Architecture
NCCL	–	NVIDIA Collective Communication Library
OMP	–	Open Multi-Processing
MPI	–	Message Passing Interface
IS	–	Integer Sort
EP	–	Embarassingly Parallel
CG	–	Conjugate Gradient
MG	–	Multi-Grid
BT	–	Block Tri-diagonal solver
SP	–	Scalar Penta-diagonal solver
LU	–	Lower-Upper Gauss-Seidel solver
FFT	–	Fast Fourier Transform
FFTW	–	Fastest Fourier Transform in the West
MKL-FFT	–	Math Kernel Library Fast Fourier Transform

HPC	–	High Performance Computing
PIP	–	Python Index Package
OpenBLAS	–	Optimized Basic Linear Algebra Subprograms
DGEMM	–	Double Precision General Matrix Multiply
PTX	–	Parallel Thread Execution
MAPE	–	Mean Absolute Percentage Error
RMSE	–	Root Mean Square Error

1. INTRODUCTION

The Information and Communication Technology sector is responsible for a significant share of global electricity use. In 2020, the data centers, communication networks and user devices accounted for an estimated 4–6% of global electricity use [41]. This value has grown exponentially over the last years, mainly due to technological advances such as cloud computing and the rapid growth of the use of Internet services [4]. Moreover, it is predicted that ICT energy use is likely to increase until 2030 and may reach approximately 13% of global electricity use. According to the increased energy demand, vast efforts have been put in order to improve the energy efficiency across the ICT sector with a success. As a result the overall energy use remains mostly flat according to some estimates.

The goal of energy efficiency increase is reached using various methods. Some of them involve the kernels optimizations and using the energy methods of computing [19]. Such methods focus mainly on load imbalance, mixed precision in floating-point operations. Other methods of increasing computational efficiency are related to power-capping [24]. Such an approach assumes setting a certain power limit level on CPU or GPU in order to achieve a power/performance trade-off. Appropriate limitations of power draw results in slightly longer execution times and significant savings, which render this method a viable option.

In order to make such implementation meaningful, we must make sure that the tools used to validate them are also as precise as possible. For the measurement of the power draw of a CPU, Intel provides its own interface, called Running Average Power Limit (RAPL) [18]. On the GPU side, NVIDIA provides the NVIDIA Management Library (NVML) [29]. Unfortunately, concerns arise on the precision of such softwares, mainly because their providers don't share any information about estimated error of power draw readings, leaving researchers questioning their practical use. In order to verify the precision of software measurement tools, the external measurement tool is used – the Yokogawa WT310E [46] [47]. Its high precision, backed up by certificates, makes it an excellent tool to benchmark the precision of Intel RAPL and NVML.

1.1. PURPOSE AND RESEARCH QUESTION

The goal of the project is to verify the accuracy and reliability of the CPU and GPU power draw measurement tools during the computational-straining benchmark applications.

1.2. SCOPE AND LIMITATION

The scope of the project is to verify whether the software power measurement tools are precise, based upon the results of the certified external measurement tool and to specify their error range in case of inaccuracy.

In case of benchmark applications, a hypothesis worth considering is whether the change of computational data impacts the power draw measurements between hardware and software tools. There are experiments that proved that increasing the data used in the benchmarks influences the increasingly different results from the tools. Another hypothesis, however, makes the claim that the application of power capping does not impact the results. In order to make a conclusion, all those configurations should be investigated independently. Another aspect worth investigating is the use of various power supply units, both the server and consumer grade. Those tests would give us insight, whether they differ in total power draw or is it more likely associated with the certificates they have.

The project is limited to testing CPUs and software released by Intel Corporation and to testing the GPUs and software released by NVIDIA Corporation.

1.3. PROJECT REQUIREMENTS

The project requires a reliable testbed in order to run the computations, verified benchmark applications and certified external measurement tool. Moreover, the computational station must be exclusively reserved for the time of research in order to prevent other user's applications from interfering in the tests results, as well as the tests itself should be repeated several times in order to maintain credibility.

The workstations on which the experiments will be conducted are the two different university computational nodes. One of them consists of two server grade CPUs and eight GPUs, while the second one is equipped with two server grade CPUs and four GPUs. Both nodes are adapted to use a custom power strip that supports the use of external measurement tool.

The tests benchmark applications are considered suitable for the purpose of the project, when they are able to strain the workstation's hardware, using their entire computational resources. The chosen applications for this task are mainly NAS Parallel Benchmarks [5] [27] – a set of benchmarks designed to evaluate the performance of parallel computing systems. The NPB suite contains a variety of benchmarks, including linear algebra, FFT, stencil computations and others. The benchmarks are intended to be representative of some important real-world application problems and can be used to assess the performance of various systems under different conditions.

The external measurement tool is Yokogawa WT310E. It is a precise and reliable equipment that will serve as the ground truth in verification of reading from the software tools.

2. THEORETICAL BACKGROUND

2.1. MEASUREMENT SOFTWARE

2.1.1. INTEL RAPL

Intel RAPL (Running Average Power Limit) [40] is an interface [11], which allows software to set a power limit that hardware ensures and any power control system takes it as an input and tunes behavior to ensure that this operating limit is respected. That way, it is possible to set and monitor power limits both on processor and DRAM, and by controlling the maximum average power, it matches the expected power and cooling budget. RAPL exposes its energy counters through model-specific registers (MSRs) It updates these counters once in every 1 ms. The energy is calculated as a multiple of model specific energy units. For Sandy Bridge, the energy unit is 15.3 μ J, whereas it is 61 μ J for Haswell and Skylake.

Moreover, the Intel RAPL divides the platform into four domains, which consists of:

- PP0 (Core Devices) – Includes the energy consumption by all the CPU cores in the socket(s).
- PP1 (Uncore Devices) – Includes the power consumption of integrated graphics processing unit (unavailable on the server platforms)
- DRAM – The energy consumption of the main memory.
- Package – The energy consumption of the entire socket including core and uncore.

2.1.2. NVIDIA NVML

NVIDIA Management Library (NVML) [29] – A C-based API for monitoring and managing various states of the NVIDIA GPU devices. It provides direct access to the queries and commands exposed via `nvidia-smi` [33]. The runtime version of NVML ships with the NVIDIA display driver, and the SDK provides the appropriate header, stub libraries and sample applications. Each new version of NVML is backwards compatible and is intended to be a platform for building third party applications.

There are various techniques of computing the energy consumption using the NVIDIA Management Library, which query the onboard sensors and read the power usage of the device. Such techniques are either from the native NVML API, like Sampling Monitoring Approach (SMA) or Multi-Threaded Synchronized Monitoring (MTSM) or from CUDA component and it's called Performance Application Programming Interface (PAPI) [31].

Sampling Monitoring Approach (SMA) – The C-based API provided by NVML that can query the power usage of the device and provide an instantaneous power measurement. Therefore, it can be programmed to keep reading the hardware sensor with a certain frequency. The `nvmlDeviceGetPowerUsage()` function is used to retrieve the power usage reading for the device, in milliwatts. This function is called and executed by the CPU. The highest frequency possible is 66.7 Hz, which means the measurements are done every 15 ms.

Performance Application Programming Interface (PAPI) provides an API to access the hardware performance counters found on modern processors. The various performance metrics can be read through simple programming interfaces from C or Fortran. It could be used as a middleware in different profiling and tracing tools. PAPI can work as a high-level wrapper for different components. Previously it used only Intel RAPL's interface to report the power usage and energy

consumption for Intel CPUs, but recent updates added the NVML component, which supports both measuring and capping power usage on modern NVIDIA GPU architectures. The major advantage of using PAPI is that the measurements are by default synchronized with the kernel execution. The NVML component implemented in PAPI uses the function, *getPowerUsage()* which query *nvmDeviceGetPowerUsage()* function. According to the documentation, this function is called only once when the command “papi end” is called. Thus, the power returned using this method is an instantaneous power when the kernel finishes execution.

Multi-Threaded Synchronized Monitoring (MTSM) – this method differs from SMA approach in the measurement period, a specific, exact window of the kernel execution is identified which results in recording of only the power reading of the kernel solely. Monitoring part is performed by the host CPU in that way the master thread calls and then monitors the kernel and other threads records the power, therefore it requires the use of parallel programming execution models, such as Pthreads [21] or OpenMP [34]. This approach at first initializes the volatile variable (at master thread) that is used later in recording of power readings. Then, the remaining threads execute the monitoring function in parallel and start measuring the time and power draw as the benchmark kernel starts doing the computation. After its work is done, the timing is ended and the stored measurements are synchronized, giving the precise logs of power consumption during the test period.

2.2. MEASUREMENT HARDWARE

2.2.1. YOKOGAWA WT310E

In order to perform comparison and to check the reliability of software power measurement methods, such as mentioned above Intel RAPL or NVIDIA NVML, it is mandatory to use a certified tool that would serve as the ground truth in such tests. Such a tool is Yokogawa WT310E – an external power meter that will serve this purpose in tests in this paper. It is a digital power analyzer that provides extremely low current measurement capability down to 50 micro-Amps, and a maximum of up to 26 Amps RMS. This device follows standards and certificates such as Energy Star® [36], SPECpower [42] and IEC62301 [10] / EN50564 [43] testing. This model comes from a WT300E’s family of appliances that offer a wide range of functions and enhanced specifications, allowing handling of all the measurement applications from low frequency to high frequency inverters using a single power meter. The WT300E series with the fast display update rate of 100ms, offer’s engineers a short tact time in their testing procedures. The basic accuracy for all input ranges is 0.1% rdg + 0.05% rng (50/60Hz) and DC 0.1% rdg + 0.2% rng.

To use the Yokogawa WT310E power meter, a special software has been written for easy use – the Yokotool [6]. Yokotool is a command-line tool for controlling Yokogawa power meters in Linux. The tool is written in Python and works with both Python 2.7 and Python 3. The tool comes with the ‘yokolib.PowerMeter’ module which can be used from Python scripts.

2.3. BENCHMARK APPLICATIONS

For the purpose of the experiments, benchmark applications should fully utilize the resources of the tested CPUs and GPUs, as well as be able to run on various configurations parameters. Such parameters are: being able to run in parallel on various numbers of logical processors, being able to run on one or more GPUs and being able to execute on various input parameters class sizes.

There are four benchmark sets, that satisfy mentioned goals:

- NAS Parallel Benchmarks (C++ with OMP)
- NAS Parallel Benchmarks (Fortran with MPI)
- NAS Parallel Benchmarks (CUDA)
- Custom deep learning model, based on Xceptionnet with MPI communication

In a general sense, the NAS Parallel Benchmarks are a set of programs designed and created in order to help evaluate the performance of parallel supercomputers. They are based on computational fluid dynamics applications and originally consisted of five kernels and three pseudo-applications. Later on, the benchmark suite has been extended with more kernels, such as adaptive meshes, parallel I/O, multi-zone applications, and computational grids. Every application comes with predefined and indicated problem size, labeled as class size. Moreover, the benchmark kernels are available in widely-used programming models like MPI and OpenMP, which allows for easy configuration of use with various number of CPU threads (NPB-OMP and NPB-MPI) [26], as well as execution on two or more computational nodes (NPB-MPI only). For the computations on GPUs, different set has been created, which excels in tests on a single devices (NPB-CUDA) [2] [3]. In order To run GPUs benchmarks in multi-GPUs and multi-node environments, a custom deep learning model has been created solely for the purpose of this task.

Below is listed the original set of eight NPB benchmarks, which are tested later in the experiments conducted for the purpose of this work.

Five kernels:

- **IS** – Integer Sort (random memory access)
- **EP** – Embarassingly Parallel
- **CG** – Conjugate Gradient (irregular memory access and communication)
- **MG** – Multi-Grid on a sequence of meshes
- **FT** – Discrete 3D fast Fourier Transform (all-to-all communication)

Three pseudo applications:

- **BT** – Block Tri-diagonal solver
- **SP** – Scalar Penta-diagonal solver
- **LU** – Lower-Upper Gauss-Seidel solver

The three pseudo applications mentioned earlier also come with the multi-zone versions, designed to exploit multiple levels of parallelism. Moreover, NPB suite also offers benchmarks for unstructured computation, parallel I/O and data movement. For the purpose of this work only the

original single-zone kernels and applications were used, therefore these benchmarks suites are mentioned only.

In addition to solving different computational problems, each kernel or application operates on various sizes of input data, determined during compilation, known as classes. Those classes helps choosing the right benchmark in term of execution time, which helps in measurements. Too short benchmarks may cause measurement error, due to low sampling rate of measurement instruments and too long benchmarks are unnecessary, because they prolong the overall experiments.

Benchmark classes:

- **Class S** – Very small, used for quick test purpose. Nowadays obsolete.
- **Class W** – The so-called ‘90’s workstation’ size, nowadays consisted small.
- **Classes A, B, C** – Standard test problems (~ 4 times size increase from each of the previous classes)
- **Classes D, E, F** – Large test problems (~ 16 times size increase from each of the previous classes)

2.3.1. NPB FOR CPU, C++ WITH OMP

NAS Parallel Benchmarks for CPU has been ported from Fortran to C++ – a programming language that has been developed for a long time as an Object-Oriented successor of a very popular and successful C language. It is a strongly typed, highly portable language with very high performance, compatibility and excellent yet difficult, manual memory management, which makes it a grea choice for writing code, that is meant to use in High Performance Computing, where fast execution times and low memory usages are among the most important factors. C++ offers also many well-optimised libraries to choose from, based on our goals, as well as extensive amount of documentations and books, explaining the principles of the language.

OpenMP (Open Multi-Processing) is an API (Application Programming Interface) and set of directives for parallel programming in shared-memory multiprocessing environments, primarily used for multi-threading. It’s designed to simplify the development of parallel applications by allowing developers to add parallelism to their code through compiler directives and library functions. OpenMP is particularly popular in scientific and high-performance computing applications where performance optimization is crucial.

Since NPB-OMP comes with up to eight benchmarks and each benchmark can be compiled with various class sizes, it creates a very diverse and flexible test cases to choose from, depending of our needs. More informations about choosing the correct configurations for our purposes are explained in Chapter 4: ‘Proposal of the Solution’

Another major advantage of these benchmarks are the fact, that after the compilation, one can be executed and pinned to specific logical processor using Linux’s *taskset* command [25], which is mainly used to set or retrieve the CPU affinity of a running process. In practice, it allows to set exact number of created benchmark processes and allows to manually pin them to the desired threads, which gives us the absolute control over choosing the benchmark configuration and gives us confidence, that they are correctly executed.

2.3.2. NPB FOR CPU, FORTRAN WITH MPI

Another suite of benchmarks created for execution on CPUs – an older set that was written in the world's first high level programming language, the Fortran. Fortran has been widely used in scientific, engineering and numerical computing applications for several decades. Its main advantages such as high performance, standardized libraries, support for complex numbers and parallel computing makes it a valuable and reliable choice in numerical and scientific computing to this day.

MPI (Message Passing Interface), is a standardized and portable message-passing system designed for parallel and distributed computing. It provides a set of routines and libraries for high-performance communication and coordination among processes or tasks in a parallel or distributed computing environment. MPI is particularly popular in high-performance computing (HPC) and scientific computing, where large-scale parallelism is essential for solving complex problems. In contrast to OMP, this implementation provides inter-node communication for execution of code on several nodes in parallel manner.

Similar to NPB-OMP, this benchmarks suite also comes with many kernels and class sizes to choose from, but it should be noted, that it does not offer complete flexibility in terms of choosing the number of processes – some kernels require a certain number of processes to run which are specified in the code documentation, and are caused by the way, the implementation was created.

2.3.3. NPB FOR GPU, CUDA

CUDA (Compute Unified Device Architecture), is a parallel computing platform and application programming interface (API) developed by NVIDIA Corporation. CUDA is specifically designed for accelerating general-purpose computational tasks on NVIDIA GPUs (Graphics Processing Units). It allows developers to harness the massive parallel processing power of GPUs to accelerate a wide range of applications, including scientific simulations, machine learning, image processing, and more.

Key features and concepts of CUDA include, but are not limited, to:

- **Parallel Computing Model:** CUDA enables developers to write parallel code that can be executed on GPUs.
- **GPU Acceleration:** CUDA provides a means to offload computationally intensive tasks from the CPU to the GPU.
- **CUDA C/C++ Language Extensions:** Developers can write CUDA code using C or C++ with CUDA-specific extensions. These extensions allow developers to define GPU kernels, which are functions that run in parallel on the GPU.
- **CUDA Toolkit:** The CUDA Toolkit includes the CUDA compiler, runtime libraries, and development tools. It enables developers to write, compile, and optimize CUDA applications.

As mentioned in the previously shown CPU implementations, this benchmarks suite also comes with various kernels and class sizes to choose from, according to our needs. One important thing should be noted, however – due to the nature of the way it has been implemented, the index number of the GPU, that will be used during tests must be specified in the configurations file during compilation. Therefore this benchmark sets are single-GPU only and require creation of many executable files with various configs, if the test server has many GPUs. This problem is explained in more detail and ultimately resolved in Chapter 4.

2.3.4. CUSTOM DEEP NEURAL NETWORKS MODEL

Last benchmark application used in this work fills the remaining gap for test suite – an multi-GPUs kernel that utilizes MPI and has potential for multi-node training. This application was created solely for the purpose of this paper and utilizes high-level deep learning frameworks, such as TensorFlow and Keras, widely known and used, general-purpose programming language such as Python, as well as Horovod [13] framework, that utilizes MPI and NCCL communication libraries. The mentioned components used in this benchmark are more precisely described below:

Python is a versatile and widely-used high-level programming language known for its simplicity, readability, and extensive standard libraries. The fact that it is a general-purpose language means, that it can be used in various fields, such as web development, data analysis, scientific computing and artificial intelligence. Its high-level abstraction provides a clean and readable syntax and the fact that its an interpreted language means that there is no need to compile the code before running it, shortening the time between each tests of new features. Moreover, Python comes with a large and comprehensive standard library for basic programming tasks such as file I/O, regular expressions, networking and data manipulation, and in case if someone needs more specific libraries, it features an easy solution for installing additional packages using PIP (Python Index Package).

TensorFlow is a comprehensive deep learning framework that provides low-level and high-level APIs. While one can use TensorFlow for a wide range of machine learning and deep learning tasks, it also includes the Keras API as a high-level component. TensorFlow provides more low-level control, allowing developers to customize their models and training loops extensively. This is useful for researchers and engineers who need fine-grained control over model architecture and training procedures.

Keras is an open-source high-level neural networks API written in Python. It provides a user-friendly, high-level interface for building and training neural networks. Keras was designed with simplicity and ease of use in mind.

Both TensorFlow and Keras offers extensive documentations that allows user to quickly learn the basics of theirs modules and libraries. Moreover, organizations such as Google [16] or Kaggle [20] offers a practical approach in learning such concepts, especially Kaggle comes with a wide array of many datasets and exemplary solutions, created by the community members.

Horovod [44] is an open-source distributed deep learning framework designed to make it easier to train machine learning models on distributed computing environments, such as clusters or cloud-based setups. It was developed by Uber Engineering and is designed to provide efficient and scalable distributed training for deep learning models. Horovod is particularly well-suited for distributed training with popular deep learning libraries like TensorFlow, PyTorch, and MXNet. Main advantages of Horovod framework are: large scale distributed training, that allows multi-GPU and multi-node support, efficient data parallelism, compatibility with deep learning frameworks, such as TensorFlow and Keras, resource awariness and dynamic scaling.

The custom deep learning benchmark, created for the purpose of this Master's Thesis, tackles the problem of image classification. It involves training a neural network to categorize or classify images into predefined classes or categories. The goal is to teach the model to recognize patterns and features in images that distinguish one category from another. This problem is commonly addressed using supervised learning, where the neural network learns from a labeled dataset containing examples of images and their corresponding class labels. The dataset used has been taken from the Kaggle website [35] classes of the dataset represents various species of birds, that are labeled. The birds in nearly all images occupy over 50% of the image, the photos

themselves are of good quality – no images are blurry, done in bad lighting or mislabeled. All images are $224 \times 224 \times 3$ color images in .jpg format. Considering all the informations about the dataset mentioned above, one can expect a good results when training deep neural networks models on it. During the tests and fine-tuning phase of the creation of the model, a very high accuracy of over 90% has been achieved in relatively small number of iteration. What is more important in the terms of this work, the utilization of the GPUs and their power draw has been very high, thus meeting the requirement for a good benchmark application, to use possibly as much of the devices resources during measurements as possible. Main features in this benchmark, that allowed to achieve such great results are: usage of Horovod innate function to automatically allocate memory growth based on chosen number of devices, Horovod handling the deep learning distributed optimizer among the devices via callbacks, TensorFlow's `.cache()` and `.prefetch()` utilizing maximum of GPUs VRAM and finally, usage of Xception Model – a deep convolutional neural network architecture that involves Depthwise Separable Convolutions, that has been developed mainly for the purpose of solving the images classification problems.

2.4. RELATED WORKS

2.4.1. 'A COMPARATIVE STUDY OF METHODS FOR MEASUREMENT OF ENERGY OF COMPUTING'

Authors of this work [15] investigated the accuracy of measurement of energy consumption during an application execution in order to ensure the application-level energy minimization techniques. They mentioned three most popular methods of measurement: (a) System-level physical measurements using external power meters; (b) Measurements using on-chip power sensors and (c) Energy predictive models. Later, they presented a comparison of the state-of-the-art on-chip power sensors as well as energy predictive models against system-level physical measurements using external power meters, which played the role of a credible measurement appliance.

The methodology is as follows: The ground truth tests were performed at first, using WattsUp Pro Meter [14]. The group of components that were running the benchmark kernel were defined as abstract processor – a whole that consists of multicore CPU processor, consisting of a certain number of physical cores and DRAM. In order to perform meaningful measurements, only a certain configuration of application was run, that executed solely on an abstract processor, without need of using any other system resources, such as solid state drives or network interface cards and so on. The result of such an approach is that HCLWattsUp reflects solely the power drawn from CPU and DRAM.

The researchers took other important precautions during the experiment. At first, all the resources they used have been reserved and fully dedicated to the experiment, making sure that no unwanted, third-party applications are being run in the background. During the tests, they actively monitored the disk consumption and network to ensure that there is no interference in the measurements. Another important factor of the testbed that draws power and generates heat are the fans. In default configuration, the fans speed is dependent on the increasing temperature of the CPUs, which rises as the time of the training goes on. That generates a dynamic energy consumption that impacts negatively on the outcomes of the experiment. To rule this phenomenon out, all the fans in the entire testbed are set at full speed, therefore they draw the same amount of power regardless of the actual strain put on CPUs and their temperature. All the procedures mentioned above ensure that the dynamic energy consumption obtained using HCLWattsUp, reflects the contribution solely by the abstract processor executing the given application kernel.

After determining the “ground truth” measurements using the configuration with HCLWattsUp,

the researchers conducted a series of tests that determined the dynamic energy consumption by the given application using RAPL. At first the Intel PCM/PAPI was used to obtain the base power of CPUs, both core and uncore as well as DRAM, with no straining workload applied. Then, in the next phase, using HCLWattsUp API the execution time of the given application has been obtained. After that, the Intel PCM/PAPI has been used in order to obtain the total consumption of the CPUs and DRAM, all within the execution time of a given benchmark. Lastly, the researchers responsible for the experiment calculated the dynamic energy consumption of the abstract processor by subtracting the base energy from the total energy drawn during the kernel execution. To determine the dynamic energy consumption using HCLWattsUp, all the steps mentioned before has been repeated, but using the HCLWattsUp software instead of the Intel RAPL.

The execution time of the benchmark kernels were the same for both of the power draw measurement tools, so any difference between the energy readings of the tools comes solely from their power readings. Finally, tests were conducted on three different sets of experiments in order to receive three different types of patterns.

In the first set of experiments, the FFTW and MKL-FFT energy consumption has been explored, by using a given workload size. For many tests on various problem sizes, the Intel RAPL reports showed less dynamic energy consumption for all application configurations than HCLWattsUp, but it follows the same pattern as HCLWattsUp for most of the data points. Therefore it is possible to calibrate the RAPL readings, which resulted in significant decrease of average error for the dynamic energy profile.

Second set of tests was conducted using OpenBLAS DGEMM. Executions were, again, performed using various configurations of data sizes, but results were less satisfactory than in the first tests. Like the first set of experiments, RAPL profiles lag behind the HCLWattsUp profiles. Unlike the first set of experiments where the error between both the profiles could be reduced significantly by calibration, the reduction of the average error for most of the application configurations was only as half as effective contrary to the first set of tests. This calibration, however, is again not the same for all the application configurations.

In the third and last set of experiments the team studied the dynamic energy behavior of FFTW as a function of problem size $N \times N$. The tests were performed in different problem ranges. Researchers claim that for many data points, RAPL reports an increase in dynamic energy consumption with respect to the previous data point in the profile whereas HCLWattsUp reports a decrease and vice versa. Therefore it is impossible to use calibration to reduce the average error between the profiles because of their interlacing behavior.

As a conclusion, readings from Intel RAPL and HCLWattsUp differ strongly based on executed benchmark and data size. In the first set of experiments, the FFTW and MKL-FFT energy consumption test, the Intel RAPL readings followed the pattern of HCLWattsUp readings, being even more accurate after calibrations. The second test however, showed that RAPL does not follow most of the energy consumption pattern of the power meter. This could be tuned to some extent by calibration, but not as good as in the first test case. In the last experiment, however, the RAPL does not follow the energy consumption pattern of the power meter and can not be calibrated, leaving the readings quite troublesome.

Next experiment conducted in this paper was the comparison of measurements by GPU and Xeon Phi Sensors with HCLWattsUp. The methodology of work is similar to the one explained before – the entire testbed is reserved solely for the purpose of the experiment, the fans are set to the maximum speed and only the abstract processor is measured. To strain the hardware, two applications were used:

The first one was the matrix multiplication (DGEMM), the second one was 2D-FFT. Tests were

performed on two NVIDIA GPUs: K40c [22] and P100, and one Intel coprocessor, the Intel Xeon Phi 3120P [39]. To obtain the power values from on-chip sensors on NVIDIA GPUs, the dedicated libraries were used, called NVIDIA NVML [30] and to obtain power values from Intel Xeon Phi, the Intel System Management Controller chip (SMC) [17] was used. Values from the Intel Xeon Phi can be programmatically obtained using Intel manycore platform software stack (MPSS) [17]. The methodology taken to compare the measurements using GPU and Xeon Phi sensors and HCLWattsUp is similar to this for RAPL. Briefly, HCLWattsUp API provides the dynamic energy consumption of an application using both CPU and an accelerator (GPU or Xeon Phi) instead of the components involved in its execution. Execution of an application using GPU/Xeon Phi involves the CPU host-core, DRAM and PCIe to copy the data between CPU host-core and GPU/Intel Xeon Phi. On-chip power sensors (NVML and MPSS) only provide the power consumption of GPU or Xeon Phi only. Therefore, to obtain the dynamic energy profiles of applications, the Intel RAPL was used to determine the energy contribution of CPU and DRAM. Energy contributions from data transfers using PCIe were considered as not significant.

At first, the DGEMM was used as the test benchmark with various workload sizes. The energy readings from the GPU NVIDIA K40c sensors exhibit a linear profile whereas HCLWattsUp does not. Moreover, the sensor does not follow the application behavior exhibited by HCLWattsUp for approximately two-thirds of the data points. In the case of the Intel Xeon Phi coprocessor, the results seemed to be better – sensors follow the trend exhibited by HCLWattsUp for third-fourth of the data points. However, sensors report higher dynamic energy than HCLWattsUp, but that can be reduced significantly using calibration.

In the case of the second benchmark, the 2D-FFT, the measurements by NVML follow the same trend for the majority of the data points, compared to the results from NCLWattsUp. The sensor of the Intel Xeon Phi followed the trend of HCLWattsUp for over 90% of all data points, which is a good result. Overall, Intel RAPL and NVML both exhibit the same trend for FFT. Therefore, the difference with HCLWattsUp comes from both sensors collectively.

The results of this test allows to draw several conclusions. First, the average error between measurements using sensors and HCLWattsUp can be reduced using calibration, which is, nevertheless, specific for an application configuration. Another important finding is that CPU host-core and DRAM consume equal or more dynamic energy than the accelerator for FFT applications (FFTW 2D and MKL FFT 2D), which means that data transfers (between CPU host-core and an accelerator) consume same amount of energy as that for computations on the accelerator for older generations of NVIDIA Tesla GPUs such as K40c and Intel Xeon Phi such as 3120P. However, for newer generations of Nvidia Tesla GPUs such as P100, the data transfers consume more dynamic energy than computations. It suggests that optimizing the data transfers for dynamic energy consumption is important.

2.4.2. 'VERIFIED INSTRUCTION-LEVEL ENERGY CONSUMPTION MEASUREMENT FOR NVIDIA GPUS'

Authors of this research paper [1] investigated the actual cost of the power/energy overhead of the internal microarchitecture of various NVIDIA GPUs from four different generations. In order to do so, they compared over 40 different PTX instructions and showed the effect of the CUDA compiler optimizations on the energy consumption of each instruction. To measure the power consumption, they used three different software techniques to read the GPU on-chip power sensors and determined their precision by comparing them to custom-designed hardware power measurement. The motivation of their work comes from the fact that in order to increase the performance

of the GPUs, their power consumption must be correctly and reliably measured, because it serves as a primary metric of performance evaluation. This issue is proven even more challenging, since the GPU vendors never publish the data on the actual energy cost of their GPUs' microarchitecture, therefore the independent research should be conducted in order to verify the power measurement software they provide.

The authors of the research paper prepared a set of special micro-benchmarks to stress the GPU, in order to capture the power usage of each PTX instruction [32], so the instructions were written in PTX as well. PTX is a virtual-assembly language used in NVIDIA's CUDA programming environment whose purpose is to control the exact sequence of instructions executing without any overhead. The researchers prepared two kernels for the purpose of this work – first one is tasked with adding integers and second one responsible for dividing variables with unsigned values. Since it is impossible to capture power draw of an execution of a single instruction, a different approach was proposed: the same instruction has been repeated millions of times and the power drawn during the entire test case has been measured. Then the amount of power reported by the measuring system was divided by the total number of instructions, giving the power consumed by a single PTX instruction. It is worth noting that GPUs drain power as static power and dynamic power. The static power is a constant power that the GPU consumes to maintain its operation. To eliminate the static power and any overhead dynamic power other than the instruction power consumption, the kernel's was computed twice and the energy consumption was measured both times. First, the kernel was run in a configuration to measure the total energy drawn for the operation. In the second run the back-to-back instructions were omitted and the energy measured was defined as overhead energy. Energy used on instruction was defined as subtraction of total energy and overhead energy, divided by the total number of instructions.

In this experiment, the ground truth of energy drawn by the GPUs was set by the external power meter. It was pointed out that the GPUs have two power sources: one is direct DC power, provided by a PSU, another one is the PCI-E power source, provided through the motherboard. In order to capture the total power, the measurement of current and voltage has been done for each power source simultaneously. A clamp meter and a shunt series resistor were used for the current measurement. For voltage measurement, a direct probe on the voltage line using an oscilloscope has been used. In case of measurements of current and voltage on the direct DC power supply source, everything was measured using an oscilloscope, therefore the power draw calculations were performed using a certain formula. The measurement of the PCI-E power source was more difficult. Since there wasn't any possibility to directly receive current or voltage, the authors of this paper decided to set up a special PCI-E riser board that measures the power supplied through the motherboard. Two in-series shunt resistors are used as a power sensing technique. Using the series property, the current that flows through the riser is the same current that goes to the graphics card, same with the voltages.

The experiment has been conducted for four NVIDIA GPUs from four different generations/architectures: GTX TITAN X from Maxwell architecture, GTX 1080 Ti from Pascal architecture, TITAN V from Volta architecture and TITAN RTX from Turing architecture. To compile and run the previously prepared benchmarks, the CUDA NVCC compiler [28] has been used. The results of the tests show that NVIDIA TITAN V has the lowest energy consumption per instruction among all the tested GPUs. Additionally the tests were performed on both CUDA optimized and non-optimized versions of code, and overall the optimized versions of instruction proved to be less energy hungry than the non-optimized ones. In terms of differences between various software power measuring techniques, namely PAPI versus MTSM, The dominant tendency of the results is that PAPI readings are always more than the MTSM. Although the difference is not significant

for small kernels, it can be up to 1 μ J for bigger kernels like Floating Single and Double Precision div instructions. Different software techniques (MTSM and PAPI) have been compared against the hardware setup on Volta TITAN V GPU. Compared to the ground truth hardware measurements, for all the instructions, the average Mean Absolute Percentage Error (MAPE) of MTSM Energy is 6.39 and the mean Root Mean Square Error (RMSE) is 3.97. In contrast, PAPI average MAPE is 10.24 and the average RMSE is 5.04. The results prove that MTSM is more accurate than PAPI as it is closer to what has been measured using the hardware.

2.4.3. 'MEASURING GPU POWER WITH THE K20 BUILT-IN SENSOR'

Authors of this research paper [9] investigated accurate profiling of the power consumption of GPU code when using the on-board power sensor of NVIDIA K20 GPUs. Moreover, two major anomalies that happened during the tests were more thoroughly analyzed – the first one being related to the doubling a benchmark kernel's runtime resulted with more than double energy usage, the second indicated that running two kernels in close temporal proximity inflates the energy consumption of the later kernel. Based on previous work in a similar field and set of preliminary tests, a new, reliable methodology [7] has been proposed as the conclusion of this experiment.

GPUs used in this project are NVIDIA Tesla K20, equipped with on-board sensors for querying the power consumption at runtime. As noted by the authors of the work, measurement of the power draw of the GPU using its built-in sensor is more complex than it would seem at first glance. The straightforward approach of sampling the power, computing the average, and multiplying by the runtime of the GPU code is likely to yield large errors and nonsensical results, hence the anomalies related to more energy used than expected due to increase of kernel's runtime or kernel's energy consumption increase after consecutive runs. Therefore another approach must be adopted. Methodology of the experiment is as follows: a number of unexpected behaviors when measuring a GPU's power consumption have been noted for further investigation, various observations has been noted during the tests runs conducted on the NVIDIA K20 GPUs and based on those observations and other related work, a correct way of measuring the power and energy consumption using sensor has been created. Later on it was validated for reliability by performing it multiple ways on many GPUs based on Kepler architecture, equipped with power sensor, such as the NVIDIA K20c, K20m, and K20x. The custom tool, created by the authors of the work, has been published for future use by other scientists, as an open source code.

Benchmark applications used in this paper solved two different n-body problem implementations. The algorithm models the simulation of gravity-induced motion of stars in a star cluster. The first kernel, called NB (N-Body), performs precise pairwise force calculations, which means that the same operations are performed for all n bodies, leading to a very regular implementation that maps well to GPUs. Moreover, the force calculations are independent, resulting in large amounts of parallelism. The second code, called BH, uses the Barnes-Hut algorithm to approximately compute the forces [12] [45]. It hierarchically partitions the volume around the n bodies into successively smaller cubes, called cells, until there is just one body per innermost cell. The resulting spatial hierarchy is recorded in an unbalanced octree. Each cell summarizes information about the bodies it contains. The NB code is relatively straightforward, has a high computational density, and only accesses main memory infrequently due to excellent caching in shared memory. In contrast, the BH code is quite complex, has a low computational density, performs mostly irregular [8] pointer-chasing memory accesses, and consists of multiple different kernels. Nevertheless, because of its lower time complexity, it is about 33 times faster than the NB code when

simulating one million stars.

In order to conduct the energy measurement from the GPU power sensor, the authors of the work wrote their own tool to query the sensor via the NVIDIA Management Library (NVML) interface, which returns the power readings in milliwatts. The sampling intervals of the measurement are lowest possible – 15 ms between measurements. At first, during the tests, there was a noticeable power lag and measurement distortion – power profiling tends to lag behind the kernel activity and shape of the profile does not match the kernel activity in both shape (minor difference) and time (major difference). The key insight in creating a model of correct measurement is the fact that the power sensor gradually approaches the true power level rather than doing so instantly. Since the ‘curved’ power readings between time when kernel start running and time when the power curve stabilizes reminded the authors of this work of capacitors charging and discharging, they tested whether the power profiles can be described by the same formulas. This turned out to work very well in the end. It is assumed that this is the case because the power sensor hardware uses a capacitor of some sort. After this revelation, the authors determined the ‘capacitance’ of the power sensor by using a single capacitor function to approximate the curve between the kernel start time and kernel stop time. After that, they determined the value of the capacitance that minimizes the sum of the differences between the measured values and the function values. As the capacitance is constant, it only needs to be established once for a given GPU, which is $C = 833.333$ on all tested K20 GPUs. Computing the true power draw value then become a single function of the slope of power profile derived in time domain and is shown in a function below:

$$P_{true}(t_i) = P_{meas}(t_i) + C \times \left(\frac{P_{meas}(t_{i+1}) - P_{meas}(t_{i-1})}{t_{i+1} - t_{i-1}} \right) [W] \quad (2.1)$$

Moving back onto the recommended steps for this experiment, following assumption should be considered: highest possible sample rate for NVML (which is 66.7 Hz / 15 ms between intervals) as well as including the time stamps, removal of consecutive samples of the same value that are no more than 4 ms apart of each other, computation of true power with the equation mentioned above and finally, computation of the true energy consumption by integrating the true power, using the time stamps, over all intervals where the power level is above the ‘active idle’ threshold of 52.5 W.

After incorporating the steps mentioned above in the tests, the authors of the research paper validated their results. To do so, they checked if the computed power profile follows the GPU kernel activity and also they revisited anomalies that they encountered before in order to check if their new approach eliminates them. In the end, the profiling almost instantly shoots up when the kernel starts, stays at a (more or less) constant level during execution, and almost instantly drops to the aforementioned 52.5 W after the kernel stops. Importantly, the power level during execution coincides with the asymptotic power between kernel start and kernel stop, which verifies the above hypothesis. This observation gives insight that power should be integrated from the time point of kernel start to the point of kernel end. Any energy consumption by the GPU before or after the kernel execution is due to idling (at different power levels) and is a function of time but independent of the kernel. In case of anomalies, the first one was regarding the kernel runtime changes and unintuitive increase of measured power draw. In the early tests that were based solely on readings from NVML, after increasing the kernel’s runtime by two times, the power draw readings were not increased proportionally as well, they were higher by approximately 8% than expected. The corrected power profile, however, indicates that, in fact, the energy consumption increase follows the total runtime one-to-one, thus resolving this particular anomaly. The second anomaly was that running the same kernel twice in close temporal proximity inflates the energy

consumption of the second invocation. Once again, the power profiling proposed by the author clearly resolves this anomaly as the two corrected profiles are now at the same level (and the power level between the kernel runs is at the active-idle level). In other words, the computed power profile of a kernel is unaffected by prior kernel runs, which is an important advantage of this approach. This means that there is no need to have to delay kernel runs until the GPU reaches its idle power level before one can measure the energy consumption of the next kernel. Those tests were also validated on other GPUs and the results showed that they behave in a similar manner to the first one. The only notable difference is that all of the measurements are a few watts higher on the second GPU. The difference is, however, within the 5 W absolute accuracy of the sensor. The profiled power obtained from tests on other GPUs are all very similar to each other, which further validates the used methodology.

As a conclusion of this work, many results and insights were obtained, such as: Power profile is distorted with respect to the kernel activity; the measured power lags behind the kernel activity; running multiple kernels one after another inflates the power draw of the later kernels; after a short-running kernel, the power draw can even increase for a while; integrating the power to a discernable time after a kernel stops does not correctly compensate for the power lag; the sampling interval lengths vary greatly; the GPU sensor only performs power measurements once in a while; the true sampling rate may be too low to accurately measure short-running kernels and the PCI-bus activity is not included in the sensor's measurements. This paper proposes and evaluates a power- and energy-measurement methodology that is accurate even in the presence of the above problems. It computes the true instant power consumption from the measured power samples, accounts for variations in the sampling frequency, and correctly identifies when kernels are running and when the GPU is in active-idle mode waiting for the next kernel launch.

3. DEFINITION OF EXPERIMENTS

3.1. CHOOSING THE CORRECT CONFIGURATIONS

In order to receive a wide set of data from the measurements, tests should be conducted using broad range of benchmark kernels with various classes sizes. To break down the hierarchy of the test cases, it is a sound strategy to start from the biggest piece – the server. Computational nodes used for the experiments are server-grade machines, codenamed ‘sanna.kask’ and ‘vinnana.kask’ respectively. Next step are the devices that the individual benchmarks use during runs. Each node utilize three setting: ‘CPUs’, ‘GPUs’ and the ‘Hybrid’ benchmarks, which are combination of previous two. After choosing the device, the correct implementation is specified, based both on the currently used server and device. This way, it has been decided, that in order to diversify the tests, implementation on the servers should vary. Experiments on ‘sanna.kask’ are utilizing the implementation based on OpenMP, where the ‘vinnana.kask’ server uses the solutions with MPI, TensorFlow and Horovod. Based on the device utilized, implementation are either OMP-C++ or MPI-Fortran for using on the CPUs and OMP-CUDA or Horovod-Python on GPUs. Another parameter of the tests are the benchmarks, which are the short applications or kernels, that solves various problems based on computational fluid dynamics. The main goal of the benchmark is to evaluate the performance of workstations by utilizing its targeted resources to the maximum. Properly conducted tests should include various benchmarks in order to avoid bias on the collected measurements data. The sets of used benchmarks are more complex that the previously mentioned parts of the configuration and will be explained graphically on the charts in the next section, but it should be mentioned, that all implementation except Horovod-Python uses three different benchmarks with various class sized. The Horovod-Python implementation use only one application, the custom deep neural networks model. Finally, based on both implementation and benchmark, the configuration is defined. By configuration we should expect the resources used – for the CPUs benchmark, such resources are the CPUs themselves and their logical processors, and for the GPUs benchmarks – the GPUs themselves and reserved for them CPUs threads. Again, the amount of configurations set for the purpose of the experiments are significant, therefore this part will also be represented in an easier to understand, graphical way.

All the mentioned benchmarks fulfill the most important requirement, that is, they utilize the targeted devices resources to the maximum or near-maximum. Therefore, the next factor should be the execution time. It should be noted, that this parameter is important, because of two reasons – first, too short benchmark kernel exposes the measurements to inaccuracy. To put this in a practical way, we should at first find the measurements tool with lowest sampling rate, that is possible due to their implementation. In our case it is the Yokogawa WT310E Power Meter – its software, the Yokotool, enables the measurements with the shortest interval of 0.1 second. Other methods, such as Linux Perf and NVML should use this sampling frequency as well, in order to maintain cohesion. In one of the research papers [23] that covered the topic of energy awarness in parallel applications, authors discussed the sampling period of *nvidia-smi*, which is 1 second, and the execution times of various benchmarks. Those values varied between 20 and 200 seconds, which gives the maximal error of GPU power draw reading less than 5% and the minimal error for the longest test runs is less than 0.5%.

This insight gives us a simple formula:

$$E = \frac{f}{T}[\%] \quad (3.1)$$

Where:

- E** – Error of reading [%]
- f** – Lowest sampling rate of the measurements tool [s]
- T** – Execution time of benchmark [s]

Based on that information, it has been tentatively accepted, that the maximum error of power draw reading should be lower than 0.5%, to be considered satisfactory. With lowest possible sampling period of 0.1 second, the execution time of any benchmark should be at least 20 seconds. To make sure, that this constraint is satisfied at all times, the preliminary tests should be done utilizing all of the resources on given device, i.e. CPU benchmarks should be run on all logical processors available in a parallel manner. At last, benchmark kernel should be also short enough in order to avoid excessively long execution time.

3.2. PRELIMINARY TESTS

The goal of the set of preliminary tests is to choose three, best suitable benchmarks from each implementation, based on the metrics stated earlier.

- Implementation: **OMP-CPP**
 - Benchmarks available: IS, FT, EP, CG, MG, LU, SP & BT
 - Class sizes to choose from: B, C & D
 - 24 different combinations of benchmarks and class sizes in total
 - All tests were run on 2 CPUs and 40 logical processors in parallel
- Implementation: **OMP-CUDA**
 - Benchmarks available: IS, FT, EP, CG, MG, LU, SP & BT
 - Class sizes to choose from: C, D & E
 - 24 different combinations of benchmarks and class sizes in total
 - All tests were run on a single GPU with the same grid sizes each time
- Implementation: **MPI-Fortran**
 - Benchmarks available: IS, FT, EP, CG, MG & LU
 - Class sizes to choose from: B, C & D
 - 18 different combinations of benchmarks and class sizes in total
 - All tests were run on 2 CPUs and 32 logical processors in parallel
- Implementation: **Horovod-Python**
 - Benchmark available: Xception
 - Number of training epochs: 1, 3 & 5
 - 3 different combinations of model training parameters to choose from
 - All tests were run in configurations with 1, 2 & 4 GPUs, to check training behavior

Every configuration stated above has been conducted 3 times in order to obtain an average results, which are presented in the tables below.

Table 3.1. Execution times of OMP-CPP benchmarks

Benchmark	Class size	Execution time [s]	Benchmark	Class size	Execution time [s]
IS	B	0.35	FT	B	2.92
IS	C	1.25	FT	C	16.14
IS	D	37.53	FT	D	371
EP	B	3.06	CG	B	7.77
EP	C	11.96	CG	C	27.2
EP	D	177	CG	D	(freezed)
MG	B	2.32	LU	B	9.88
MG	C	17.459	LU	C	34.88
MG	D	185	LU	D	1111
SP	B	14.48	BT	B	15.82
SP	C	(freezed)	BT	C	55.182
SP	D	(core dumped)	BT	D	1185

According to tests, three CPUs benchmarks has been chosen as the most suitable for the main experiments on *sanna.kask* later. Those benchmarks are:

- **bt.C** – Block Tri-diagonal solver, class size ‘C’
- **is.D** – Integer Sort, class size ‘D’
- **lu.C** – Lower-Upper Gauss-Seidel solver, class size ‘C’

Table 3.2. Execution times of OMP-CUDA benchmarks

Benchmark	Class size	Execution time [s]	Benchmark	Class size	Execution time [s]
IS	C	(too short)	FT	C	6.63
IS	D	15.33	FT	D	(unsuccessful)
IS	E	(not defined)	FT	E	(unsuccessful)
EP	C	(too short)	CG	C	(too short)
EP	D	27.16	CG	D	(freezed)
EP	E	442.79	CG	E	(freezed)
MG	C	(too short)	LU	C	16.89
MG	D	(unsuccessful)	LU	D	300.74
MG	E	(unsuccessful)	LU	E	(unsuccessful)
SP	C	10.39	BT	C	12.19
SP	D	220.86	BT	D	(unsuccessful)
SP	E	(unsuccessful)	BT	E	(unsuccessful)

According to tests, three GPUs benchmarks has been chosen as the most suitable for the main experiments on *sanna.kask* later. Those benchmarks are:

- **lu.D** – Lower-Upper Gauss-Seidel solver, class size ‘D’
- **sp.D** – Scalar Penta-diagonal solver, class size ‘D’
- **ep.D** – Embarassingly Parallel, class size ‘D’

Table 3.3. Execution times of MPI-Fortran benchmarks

Benchmark	Class size	Execution time [s]	Benchmark	Class size	Execution time [s]
IS	B	0.18	FT	B	3.11
IS	C	0.8	FT	C	13.83
IS	D	17.02	FT	D	360.15
EP	B	1.11	CG	B	3.33
EP	C	4.39	CG	C	8.57
EP	D	71.16	CG	D	761.02
MG	B	0.48	LU	B	10.97
MG	C	4.5	LU	C	41.81
MG	D	88.34	LU	D	770.52

According to tests, three CPUs benchmarks has been chosen as the most suitable for the main experiments on *vinnana.kask* later. Those benchmarks are:

- **ep.D.x** – Embarassingly Parallel, class size ‘D’
- **lu.C.x** – Lower-Upper Gauss-Seidel solver, class size ‘C’
- **is.D.x** – Integer Sort, class size ‘D’

Table 3.4. Execution times of Horovod-Python benchmarks

Number of GPUs used	Number of training epochs	Execution time [s]	Test accuracy [%]
1	1	124.2	88.99
1	3	326.7	91.30
1	5	528.6	91.59
2	1	72.2	87.9
2	3	176.2	90.72
2	5	280.9	92.46
4	1	48.5	86.81
4	3	101.1	91.59
4	5	155.3	93.91

According to tests, the best number of epochs for the training of ‘Xception’ model is 1. Test accuracy of model is high and the execution time is satisfactory.

What is worth noting in terms of GPUs benchmarks is the fact, that while utilizing between 90~100% for the GPUs resources, various kernels results in different power draw values. The averages of those values are listed below and may provide an additional insight on the future experiments:

- **lu.D** – Power draw fluctuates between 200~230 [W]
- **sp.D** – Average power draw of 200~210 [W]
- **ep.D** – Average power draw of 155~160 [W]
- **Xception** – Average power draw of 180~185 [W]

3.3. OVERVIEW ON ALL FINAL CONFIGURATIONS

The measurements will be conducted on many different setups, therefore it is important to explain throughly the entire hierarchy of the configurations.

3.3.1. SERVERS

The tests will be conducted on two computational nodes: *sanna.kask* and *vinnana.kask*. Both of those machines have two Intel Xeon CPUs and several NVIDIA Quadro RTX GPUs. Paired with huge amount of RAM, these servers are suitable choices in terms of choosing the testbed machines.



Fig. 3.1. sanna.kask




Fig. 3.2. vinnana.kask

Despite the similarity between those two servers, their components slightly differs from each other. A detailed comparison of their parts is listed in the table below:

Table 3.5. Overview on components of the servers used in experiments

	sanna.kask	vinnana.kask
OS	Linux Ubuntu 22.04 LTS	Linux Ubuntu 22.04 LTS
CPUs	2 × Intel(R) Xeon(R) Silver 4210 CPU 2.20GHz, TDP: 10/10 85 [W]	2 × Intel(R) Xeon(R) Silver 4210 CPU 2.20GHz, TDP: 10/10 85 [W]
GPUs	8 × NVIDIA Quadro RTX 6000, VRAM: 24 GB, power draw range: 100-260 [W]	4 × NVIDIA Quadro RTX 5000, VRAM: 16 GB, power draw range: 100-230 [W]
Motherboard	Inspur YZMB01130107	Inspur YZMB01130107
RAM	384 GB, DDR4, 2400 (MHz)	384 GB, DDR4, 2400 (MHz)
PSUs	4 × Delta Electronics DPS-2200AB-2	4 × Delta Electronics DPS-2200AB-2

3.3.2. DEVICES

Next, benchmarks will be organized based on devices, that will be utilized during tests. On this level, we distinguish three configurations: *CPUs*, *GPUs* and the *Hybrid*, when both are utilized during experiments.



Fig. 3.3. CPUs benchmarks, utilizing the Intel Xeon Silver CPUs



Fig. 3.4. GPUs benchmarks, utilizing the NVIDIA Quadro RTX GPUs



Fig. 3.5. Hybrid configuration, utilizing both CPUs and GPUs

3.3.3. IMPLEMENTATIONS

After receiving information about used server and devices, the correct implementations are used. For the first node, *sanna.kask* we distinguish *OMP-CPP* for CPUs, *OMP-CUDA* for GPUs and *OMP-CPP+OMP-CUDA* for Hybrid. For the second node, *vinnana.kask* we have *MPI-Fortran* for CPUs, *Horovod-Python* for GPUs and *MPI-Fortran+Horovod-Python* for Hybrid.



Fig. 3.6. CPU benchmarks implementation used on *sanna.kask*



Fig. 3.7. GPU benchmarks implementation used on *sanna.kask*



Fig. 3.8. CPU benchmarks implementation used on *vinnana.kask*



Fig. 3.9. GPU benchmarks implementation used on *vinnana.kask*

Table 3.6. Overview on implementations used in tests

	sanna.kask	vinnana.kask
CPUs	OMP-CPP	MPI-Fortran
GPUs	OMP-CUDA	Horovod-Python
Hybrid	OMP-CPP + OMP-CUDA	MPI-Fortran + Horovod-Python

3.3.4. BENCHMARKS

Based on the server and implementation currently used, benchmarks are chosen by a scheduler script to run. The main task the benchmarks are responsible is to utilize the devices' resources for the purpose of the tests.

**Fig. 3.10.** Benchmarks evaluation example

In the tables below are listed the benchmarks chosen for final tests:

Table 3.7. Overview on benchmarks used in tests on sanna.kask

	OMP-CPP	OMP-CUDA	OMP-CPP + OMP-CUDA
Benchmark #1	bt.C	lu.D	bt.C + lu.D
Benchmark #2	is.D	sp.D	is.D + sp.D
Benchmark #3	lu.C	ep.D	lu.D + ep.D

Table 3.8. Overview on benchmarks used in tests on vinnana.kask

	MPI-Fortran	Horovod-Python	MPI-Fortran + Horovod-Python
Benchmark #1	ep.D.x	XCeption	ep.D.x + XCeption
Benchmark #2	is.D.x	—	is.D.x + XCeption
Benchmark #3	lu.C.x	—	lu.C.x + XCeption

3.3.5. CONFIGURATIONS

Configurations are based on currently used implementation and defines, which devices' resources should be allocated for the purpose of the current benchmark kernel. Configuration vary heavily between each other and the comprehensive data sheet is listed in the tables below:

Table 3.9. Configurations used on node *sanna.kask*

sanna.kask				
Device	Implementation			
	OMP-CPP			
1 CPU	1 Thread	5 Threads	10 Threads	20 Threads
2 CPUs	2 Threads	10 Threads	20 Threads	40 Threads
	OMP-CUDA			
GPUs	1 GPU & 1 Thread	2 GPUs & 2 Threads	4 GPUs & 4 Threads	8 GPUs & 8 Threads
	OMP-CPP + OMP-CUDA			
Hybrid	1 CPU & 4 Threads 1 GPU & 1 Thread	1 CPU & 8 Threads 2 GPUs & 2 Threads	2 CPUs & 16 Threads 4 GPUs & 4 Threads	2 CPUs & 32 Threads 8 GPUs & 8 Threads

Table 3.10. Configurations used on node *vinnana.kask*

vinnana.kask				
Device	Implementation			
	MPI-Fortran			
CPUs	4 Processes	8 Processes	16 Processes	32 Processes
	Horovod-Python			
GPUs	—	1 GPU	2 GPUs	4 GPUs
	MPI-Fortran + Horovod-Python			
Hybrid	—	8 Processes & 1 GPU	16 Processes & 2 GPUs	32 Processes & 4 GPUs

It is worth noting, that the OMP-based implementations that has been prepared to run on *sanna.kask* node has been pinned carefully to respective CPUs logical processors and GPUs, based on input configuration values. Therefore, using a CPUs monitoring tools, such as *top* or *htop*, one can see the utilization of individual logical processors. Such graphical representation has been collected and prepared for inspection in the set of figures below. The MPI-based implementations were run without manual pinning of the processes to threads, leaving load balancing to the framework itself.

Overview on logical processors utilization in CPUs benchmarks




Fig. 3.11. 1 CPU, 1 Thread




Fig. 3.12. 1 CPU, 5 Threads




Fig. 3.13. 1 CPU, 10 Threads




Fig. 3.14. 1 CPU, 20 Threads




Fig. 3.15. 2 CPUs, 2 Threads




Fig. 3.16. 2 CPUs, 10 Threads




Fig. 3.17. 2 CPUs, 20 Threads




Fig. 3.18. 2 CPUs, 40 Threads

Overview on logical processors utilization in GPUs and Hybrid benchmarks




Fig. 3.19. 1 GPU, 1 Thread




Fig. 3.20. 2 GPUs, 2 Threads




Fig. 3.21. 4 GPUs, 4 Threads




Fig. 3.22. 8 GPUs, 8 Threads




Fig. 3.23. 1 CPU, 4 Threads + 1 GPU, 1 Thread




Fig. 3.24. 1 CPU, 8 Threads + 2 GPUs, 2 Threads




Fig. 3.25. 2 CPUs, 16 Threads + 4 GPU, 4 Threads




Fig. 3.26. 2 CPUs, 32 Threads + 8 GPUs, 8 Threads

4. EXPERIMENTS

4.1. PROPOSED WORKFLOW AND METHODOLOGY

In order to obtain the details of how system-level physical measurement estimates energy consumption by a component (such as a CPU or a GPU) during an application execution, several steps must be taken:

1. Exclusive reservation of the entire computational node.
2. Observation of the disk consumption and network usage before and during tests.
3. Monitoring of the CPUs and GPUs utilization before and during tests.
4. Running the benchmark kernels on an abstract processor only. Abstract processor comprises of the multicore CPU processor consisting of a certain number of physical cores and DRAM.
5. Gathering of the power measurements.
6. Verification of the accuracy and reliability of the software measurements tools, based on ground truth results.

One of the notable mentions that could be done in order to reduce the amount of uncertain power draw measurements done by the background components is setting the fans to full speed. This solution have a potential of reducing power draw fluctuations, especially during higher workloads, ièwhen running benchmarks kernels utilizing maximum amount of GPUs or running Hybrid configuration, where power draw of the entire nodes are very high. This could not be implemented, however, as the administrator of the department's servers stated, that interference in servers fans could be crucial for the nodes stability.

For Intel RAPL / NVIDIA Management Library:

1. Obtain base power of idle CPUs / GPUs.
2. Obtain execution time of benchmark application.
3. Obtain total energy consumption of the CPUs / GPUs during tests.
4. Calculate dynamic energy consumption by subtracting base energy from total energy used during run.

For Yokotool:

1. Obtain base power of idle CPUs.
2. Obtain execution time of benchmark application.
3. Obtain total energy consumption of the CPUs during tests.
4. Calculate dynamic energy consumption by subtracting base power from total energy used during run

In addition to the main experiments workflow, another methodology must be adapted – the data collection methodology. In order for the results to be properly comparable, several point have to be met:

1. Tests environment must be identical in every case, to eliminate discrepancy of the results.
2. The results of the power draw reading must be properly compared for the device only measurements (Intel RAPL / NVML) and the measurements of the entire node (Yokotool).
3. Experiments should be conducted on different nodes that utilizes different hardware, in order to state repeatability of tests.
4. Experiments should be conducted, using different benchmark kernels or application, to remove the possibility of bias of the results, due to poor diversification of test cases.
5. Test runs must be repeated many times.

4.2. MAIN TESTS

4.2.1. OVERVIEW ON THE SCHEDULER SCRIPT

In order to automate the experiments, an entire scheduler script had to be created. Its has two main tasks: first is to store the information about the configurations and run the benchmarks according to the presets. The second is to run the measurements softwares and save all the results in the ordered manner.

Both of these goals are reached by using dictionaries with key-value pairs. That created a tree-like dependencies between the corresponding layers of configurations. Finally, that solution works for both choosing the right config and providing the path to save measurements.

4.2.2. MAIN AUTOMATION FUNCTION

The entire scheduler script consists of classes and functions that are explicitly designated for their purposes in the code. Overview of them is as follows:

- **class ‘Config’** – Contains all the informations about the servers, devices, implementations, benchmarks and configurations, handles the relations between them and provides correct pathes to the corresponding measurements directories.
- **class ‘Benchmark’** – Defines functions responsible for executing CPUs and GPUs benchmarks, as well as the measurements softwares: Linux Perf, NVIDIA Management Library and Yokotool
- **class ‘Execution’** – This class contains functions tasked with calling the benchmark kernels and checking their status, if they are still running, for the purpose of ending the measurements. Since the measurements softwares are highly dependent on benchmarks being run, they are executed directly from the *main()* function and the *Execution* class only has the functions tasked with the proper termination of them.
- **function ‘scheduler()’** – This major function triggers secondary functions from *Execution* class and watched their status.
- **function ‘main()’** – Runs every configuration sequentially, based on the lists of presets. Additionally, repeats every ten times in order to achieve repeatability of the experiments.

In order to visualize the entire workflow of the scheduler, as well as the workings of the individual processes, two charts has been created:

- **General Flowchart** – This chart (**Fig. 4.1**) describes the relations between the currently used configurations and the instructions executed based on those conditions.
- **Processes Flowchart** – This chart (**Fig. 4.2**) shows the working principles of the *runner()* subfunction, which shows the benchmarks and measurements softwares are started on a high level of abstraction.

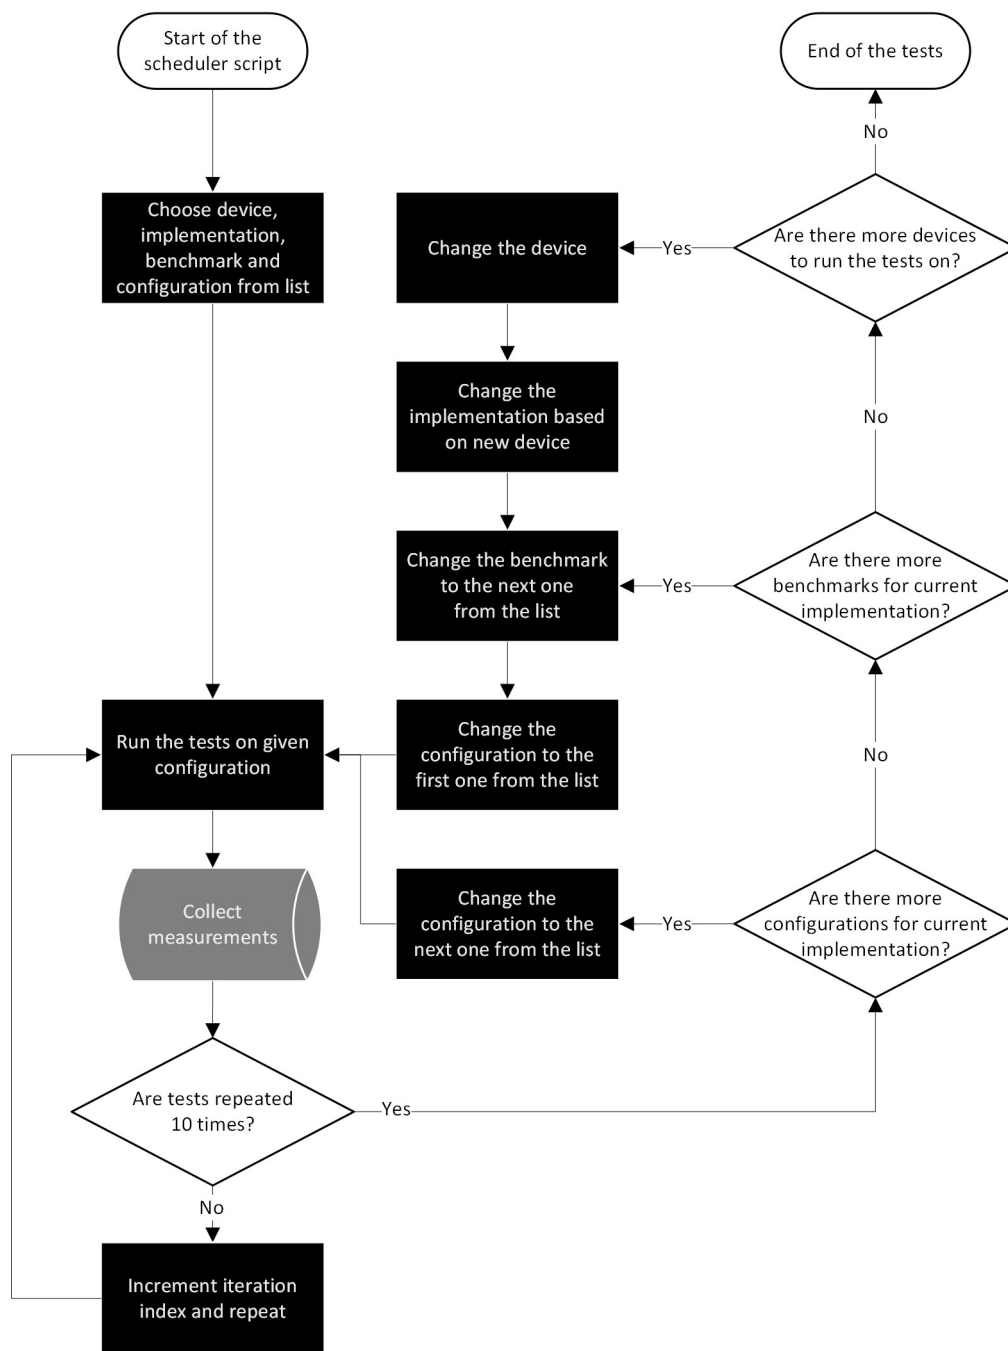


Fig. 4.1. General Flowchart




Fig. 4.2. Processes Flowchart

4.2.3. THREADS PINNING AND KERNELS EXECUTION

This function is responsible of executing CPUs benchmarks, based on the given configuration. It creates a separate processes by utilizing the Python *subprocess* module.

```
cpu_benchmark = subprocess.Popen(
    [
        "taskset --cpu-list <T> <P> <B> > /dev/null 2>&1"
    ],
    shell=True
)
return cpu_benchmark.pid
```

Here is an explanation of every part of the command:

- **cpu_benchmark** – A variable of type *<class 'subprocess.Popen'>* is created mainly in order to retrieve PID later on.
- **subprocess.Popen** – The underlying process creation and management is handled by the Popen class. Its function is to execute a child program in a new process.
- **taskset** – This command is used to set or retrieve the CPU affinity of a running process given its pid, or to launch a new command with a given CPU affinity.
- **--cpu-list** – This option interprets mask as numerical list of processors instead of a bitmask. Numbers are separated by commas and may include ranges. For example: 0,5,8-11.
- Variables that dynamically changes based on the configurations:
 - T** – Logical processors indexes
 - P** – Specified absolute path to the correct measurements folder
 - B** – Currently used benchmark kernel
- **> /dev/null 2>&1** – Redirecting *stderr* containing error messages from the executed command or script to *stdout* to the output of the command. Both are, in fact, redirected then to the so-called *null device*. The result of that action is suppression of all messages printed by the benchmark kernels. It is useful when collecting logs from the terminal, that is running the entire scheduler script, without unnecessary messages.
- **shell=True** – Invokes the program as 'shell'

Finally, the function returns the PID of newly created process as an integer value. It is done for the purpose of terminating the benchmark in a Hybrid configuration.

4.2.4. GPUS AND THREADS MANAGEMENT

This function consists of two parts: first part is responsible for executing the Horovod-Python benchmark and the second part is responsible for running the OMP-CUDA benchmark.

```
gpu_benchmark = subprocess.Popen([
    "mpirun -np <N> --map-by socket -x NCCL_DEBUG=INFO \
    python3 <P>+\"Xception.py > /dev/null 2>&1\"",
],
    shell=True
)
return gpu_benchmark.pid
```

Here is an explanation of every part of the command:

- **gpu_benchmark** – A variable of type `<class 'subprocess.Popen'>` is created mainly in order to retrieve PID later on.
- **subprocess.Popen** – The underlying process creation and management is handled by the Popen class. Its function is to execute a child program in a new process.
- **mpirun** – This command is used to execute serial and parallel jobs. It will run X copies of specified program in the current run-time environment and scheduling (by default) in a round-robin fashion by CPU slot.
- **-np** – This option specifies, how many processes will be started.
- Variables that dynamically changes based on the configurations:
 - N** – Number of GPUs used in training
 - P** – Specified absolute path to the correct measurements folder
- **--map-by socket** – Map to the specified object, such as slot, hwthread, core, socket, numa, board, node and more. In this particular case, benchmark application is mapped by 'socket', allowing to utilize multiple GPUs for training in a distributed manner.
- **-x** – Export the specified environment variables to the remote nodes before executing the program.
- **NCCL_DEBUG=INFO** – This flag is used for debugging. In case of NCCL failure, you can set NCCL_DEBUG=INFO to print an explicit warning message as well as basic NCCL initialization information.
- **python3** – Specify the use of Python interpreter when executing the script of deep neural networks model training.
- **Xception.py** – Name of the script file.
- **> /dev/null 2>&1** – As mentioned in the previous subsection, this command suppresses the output from the terminal, in order to avoid the unnecessary messages.
- **shell=True** – Invokes the program as 'shell'

Finally, the function returns the PID of newly created process as an integer value. It is done for the purpose of terminating the benchmark in a Hybrid configuration.

The second part of the function that runs the OMP-CUDA benchmarks is as follows:

```
list_of_gpu_benchmarks = []
value = Config.taskset_gpu[configuration_gpu]
for i in range(0, len(value), 1):
    gpu_benchmark = subprocess.Popen(
        [
            "taskset --cpu-list <I> <P+I> <B> > /dev/null 2>&1"
        ],
        shell=True,
    )
    list_of_gpu_benchmarks.append(gpu_benchmark.pid)
return list_of_gpu_benchmarks
```

Here is an explanation of every part of the command:

- **gpu_benchmark** – A variable of type *<class 'subprocess.Popen'>* is created mainly in order to retrieve PID later on.
- **subprocess.Popen** – The underlying process creation and management is handled by the Popen class. Its function is to execute a child program in a new process.
- **taskset** – This command is used to set or retrieve the CPU affinity of a running process given its pid, or to launch a new command with a given CPU affinity.
- **--cpu-list** – This option interprets mask as numerical list of processors instead of a bitmask. Numbers are separated by commas and may include ranges.
- Variables that dynamically changes based on the configurations:
 - I** – Allocation of individual logical processors to the GPUs, based on the index number from special dictionary
 - P+I** – Specified absolute path to the correct measurements folder, modified by the number of total GPUs used
 - B** – Currently used benchmark kernel
- **> /dev/null 2>&1** – As mentioned in the previous subsection, this command suppresses the output from the terminal, in order to avoid the unnecessary messages.
- **shell=True** – Invokes the program as 'shell'
- **list_of_gpu_benchmarks.append(gpu_benchmark.pid)** – Filling the list with PIDs of benchmarks for later termination.

Finally, the function returns the PID of newly created process as an integer value. In this particular case, if there are more than one GPUs used in tests, the PIDs of all spawned processes are parsed as a list to the function that is responsible in orderly terminating all the kernels.

4.2.5. MEASUREMENTS WITH YOKOTOOL SOFTWARE

This function utilizes the high-level Python wrapper for Yokogawa WT310E Power Meter, the *Yokotool*.

Syntax is similar to the previous examples:

```
yokotool = subprocess.Popen(  
    [  
        "yokotool read T,P -o <P+N> > /dev/null 2>&1 &"  
    ],  
    shell=True  
)  
return yokotool.pid
```

Here is an explanation of every part of the command:

- **yokotool** – Yokotool's command line interface is based on commands and sub-commands, similar to git and many other tools. This invokes the Yokotool wrapper for use.
- **read** – Read measurements data.
- **T,P** – Specifies, what data we want to read. In this particular case it is time **T** measured from the start of the epoch (on 'UNIX time' it starts at 00:00:00 UTC on 1 January 1970) and power **P**, measured in Watts [W]. Output is separated with comma for easy manipulation of data after tests.
- **-o** – This flag redirects the output from the measurements to a file for later analysis.
- Variables that dynamically changes based on the configurations:

P+N – Specified absolute path to the correct measurements folder, modified by the current number of iterations

- **> /dev/null 2>&1** – As mentioned in the previous subsection, this command suppresses the output from the terminal, in order to avoid the unnecessary messages. In this case, and additional & is placed at the end, which means that the entire command is put as a background process.
- **shell=True** – Invokes the program as 'shell'

Finally, the function returns the PID of newly created process as an integer value. It is done for the purpose of terminating the measurements by the designated function.

4.2.6. MEASUREMENTS WITH LINUX PERF SOFTWARE

Linux Perf is a lightweight profiling tool with performance counters. It utilizes Intel RAPL for measurements of pre-defined events on CPUs. To obtain more informations about what can be measured by Linux Perf, one can use *perf list* command in the terminal.

```
list_of_perf_pids = []
pin_to_cpus = ("0", "10")
idx_names = {"0": "0", "10": "1"}
for i in pin_to_cpus:
    perf = subprocess.Popen(
        [
            "perf stat --event=power/energy-pkg/ \
            --cpu=<C> --delay 100 --interval-print 100 \
            --summary --field-separator , \
            --output <P+N> > /dev/null 2>&1 &"
        ],
        shell = True,
    )
    list_of_perf_pids.append(perf.pid)
return list_of_perf_pids
```

Here is an explanation of every part of the command:

- **perf** – Invocation of measurements tool.
- **stat** – Run a command and gather performance counter statistics.
- **--event=power/energy-pkg/** – Event selector, in this case, the measured physical quantity is the energy usage of selected CPUs, during the benchmark kernel execution.
- **--cpu=<C>** – Targeting of a specific CPU. Correct integer value is based on the output of *lscpu* command and the information about *NUMA node(s)*.
- **--delay 100** – A small delay of 100 [ms] is introduced in order to offset the slight delay of measurements of Yokogawa power meter. This solution has been introduced as a result of an observation during the preliminary tests.
- **--interval-print 100** – Measurements are performed with the same time interval of 100 [ms] set as two others measurements methods.
- **--summary** – At the end of the measurements, an additional information about total energy used during tests, as well as the total measurements time is given. Mainly used during preliminary tests as an additional insight on gathered results.
- **--field-separator** – Sets the output delimiter for easier access in softwares like LibreOffice – every printed value is separated by commas ‘,’
- **--output** – Saves measurements to file. In this case, file names matches the consecutive runs.
- Variables that dynamically changes based on the configurations:
 - P+N** – Specified absolute path to the correct measurements folder, modified by the current number of iterations

- **> /dev/null 2>&1** – Silences the output, redirecting it to null device and puts the process in the background.
- **shell=True** – Invokes the program as ‘shell’

Finally, the function returns the PID of newly created process as an integer value. It is done for the purpose of terminating the measurements by the designated function.

4.2.7. MEASUREMENTS WITH NVML HANDLING FUNCTION

The method of gathering the measurements of power draw of GPUs, using NVIDIA Management Library is a little different than in previously shown implementations. No new processes are spawned, due to the fact, that the measurements are handled by a special function already implemented in the scheduler script. This function has two core parts: the first one is responsible of precise executing the measurements every 100 [ms] – it works as an built-in scheduler, and the second part handled the invocation of NVML-specific functions and saving the results to the file.

A short code-snippet below illustrates the practical usage of NVML related function

```
import py3nvml
nvmlInit()
handle_idx0 = nvmlDeviceGetHandleByIndex(0)
# based on number of GPUs used, there are more
# variables of that kind
measure_idx0 = nvmlDeviceGetPowerUsage(handle_idx0) / 1000.0
# rest of code contains the sub-scheduler routine
# and saves the output to file.
```

- **import py3nvml** – Import of a module, that function as a high level Python wrapper for NVML.
- **nvmlInit()** – NVIDIA Management Library initialization. It is mandatory to run before calling any other methods.
- **handle_idx0 = nvmlDeviceGetHandleByIndex(0)** – create variable *handle_idx0* and assign the return value of data type *<class 'py3nvml.py3nvml.LP_struct_c_nvmlDevice_t'>*, which is, in fact, a pointer to a memory register containing data about specific GPU.
- **measure_idx0 = nvmlDeviceGetPowerUsage(handle_idx0) / 1000.0** – This function return a value of power draw of a specified GPU, in Watts [W].

In terms of proper ending the measurements process, this task is handled quite differently than in previous implementations. The *nvml()* function is started as a parallel thread, using *multi-processing.Process()* module. The thread is then started as a daemon, collecting measurements and saving them in the background. After the end of benchmark kernel, the function is ended by using innate multiprocessing function – *.terminate()*.

4.2.8. TERMINATION OF BENCHMARKS IN HYBRID CONFIGURATION

In order to make sure that the Hybrid benchmarks starts and ends at the same time, an additional subroutine has been implemented. At first, the condition is checked if used *devices* in tests are currently set to *Hybrid*. Then, a special ‘while’ loop is started, in which a special functions checks, if CPUs and GPUs benchmarks are still running. They work by periodically checking, if PID associated with respective kernels still exists. If the benchmark on one of the devices has ended, then the second one should be terminated as well. The flag is passed to a function that, based on PID or PIDs of the second benchmark, starts the cleaning process. Major facilitation comes from the utilities of *psutil* module. This solution allows to easily find all the child processes spawned by the parent benchmark process, terminate them iteratively in a ordered manner, and finally, terminate the parent process. After that, the measurements threads are terminated via another, designated for that function and the current run is ended.

4.2.9. CLEANUP OF MEASUREMENTS DAEMONS

In terms of the termination of measurements daemons, a single function handles Linux Perf, NVIDIA Management Library and Yokotool. It utilizes two popular Python modules: *os* and *signal*. The first module provides a portable way of using operating system dependent functionality [37], while the second allows defining custom handlers to be executed when a signal is received [38]. In case of Perf and Yokotool, at first, the PIDs of those processes are being received, either by using *pidof* or *pigrep* of those daemons. Then, they are terminated using following, quite self-explanatory command:

```
os.kill(<PID>, signal.SIGTERM)
```

Since measurements of GPUs power draw that uses NVML are incorporated directly into the scheduler script as an function, the entire thread is being handled and terminated by using the function, that comes from *multiprocessing* module – *.terminate()*.

4.3. ANALYSIS OF THE RESULTS AND DISCUSSION

This section contains a summary of tables with measurements results as well as the analysis of the dependencies between individual results.

The first set of tables contains the information about the average power draw and average energy used for every configuration chosen during the preliminary tests. The tables themselves are self-explanatory, the captions associated with them gives hints about used server, device, implementation, benchmark and finally, the data displayed.

Two things should be noted, however. First, the physical quantities with label (Yokogawa) means that the measurements has been done using Yokogawa WT310E Power Meter and are related to the entire node, while those with labels (CPU: <index>) and (GPU: <index>) are related to measurements done using Linux Perf and NVML respectively. Second, the configurations execution times vary – majority of the tests have shorter execution times the more resources of used devices are dedicated. Such configurations are: OMP-CPP benchmarks for *sanna.kask*, MPI-Fortran and Horovod-Python for *vinnana.kask* and their hybrid variants. This is due to the fact, those implementations utilize parallel programmings paradigms, therefore the execution speed-up can be observed. In case of the OMP-CUDA benchmarks, situation is different – those benchmarks were prepared for single GPU tests only. Since the main purpose of the benchmarks is to simply strain and fully utilize the devices resources, rather than compute meaningful research problems, the OMP-CUDA benchmarks on multi-GPUs are simply many instances of the same

benchmark, running on different GPUs. Since those kernels are independent from each other, the average execution times between configurations are roughly the same. This also explains the fact, that the Hybrid benchmarks that were run on *sanna.kask* has sometimes shorter execution times the more resources are dedicated and sometimes they are similar. The Hybrid benchmarks works in a way, that the listener function from the scheduler script checks, whether the benchmarks are running or not. If the CPUs benchmarks ends, the termination signal is send to GPUs benchmarks, resulting with varying execution times. If the GPUs benchmarks in given configuration ends faster, the CPUs benchmarks are terminated in return, resulting in similar execution times.

Table 4.1. server: **sanna.kask**, device: **1 CPU**, implementation: **OMP-CPP**,
benchmark: **bt.C**, data displayed: **power draw**

	1 CPU			
Results from 10 runs	1 Thread	5 Threads	10 Threads	20 Threads
Avg. Exec. time [s]	1054.795	216.315	114.315	101.372
Std. dev. of time [-]	0.966	0.243	0.121	0.158
(Yokogawa) Avg. power draw [W]	379.962	402.881	432.171	445.752
(Yokogawa) Std. dev. of avg. power draw [-]	0.605	0.225	0.382	1.007
(CPU: 0) Avg. power draw [W]	33.717	51.274	70.433	77.958
(CPU: 0) Std. dev. of avg. power draw [-]	0.102	0.13	0.085	0.089
(CPU: 1) Avg. power draw [W]	28.98	28.887	28.871	28.874
(CPU: 1) Std. dev. of avg. power draw [-]	0.125	0.062	0.067	0.055
(GPU: 0) Avg. power draw [W]	21.755	21.673	21.604	21.634
(GPU: 0) Std. dev. of avg. power draw [-]	0.343	0.061	0.2	0.05
(GPU: 1) Avg. power draw [W]	25.522	25.534	25.516	25.524
(GPU: 1) Std. dev. of avg. power draw [-]	0.023	0.032	0.052	0.029
(GPU: 2) Avg. power draw [W]	13.96	14.024	13.744	13.642
(GPU: 2) Std. dev. of avg. power draw [-]	0.117	0.114	0.243	0.157
(GPU: 3) Avg. power draw [W]	17.196	17.21	17.207	17.209
(GPU: 3) Std. dev. of avg. power draw [-]	0.016	0.022	0.015	0.012
(GPU: 4) Avg. power draw [W]	4.109	4.122	4.119	4.096
(GPU: 4) Std. dev. of avg. power draw [-]	0.01	0.018	0.016	0.012
(GPU: 5) Avg. power draw [W]	18.456	18.372	18.297	18.334
(GPU: 5) Std. dev. of avg. power draw [-]	0.04	0.097	0.107	0.111
(GPU: 6) Avg. power draw [W]	6.285	6.3	6.351	6.449
(GPU: 6) Std. dev. of avg. power draw [-]	0.075	0.08	0.146	0.075
(GPU: 7) Avg. power draw [W]	12.174	12. 163	12.071	12.134
(GPU: 7) Std. dev. of avg. power draw [-]	0.313	0.14	0.316	0.311

Table 4.2. server: **sanna.kask**, device: **1 CPU**, implementation: **OMP-CPP**,
benchmark: **bt.C**, data displayed: **energy used**

	1 CPU			
Results from 10 runs	1 Thread	5 Threads	10 Threads	20 Threads
Avg. Exec. time [s]	1054.795	216.315	114.315	101.372
Std. dev. of time [-]	0.966	0.243	0.121	0.158
(Yokogawa) Avg. energy used [kJ]	400.782	87.149	49.404	45.187
(Yokogawa) Std. dev. of avg. energy used [-]	0.778	0.104	0.067	0.123
(CPU: 0) Avg. energy used [kJ]	35.555	11.078	8.033	7.883
(CPU: 0) Std. dev. of avg. energy used [-]	0.123	0.032	0.012	0.006
(CPU: 1) Avg. energy used [kJ]	30.56	6.241	3.293	2.92
(CPU: 1) Std. dev. of avg. energy used [-]	0.145	0.015	0.009	0.006
(GPU: 0) Avg. energy used [kJ]	22.947	4.688	2.469	2.186
(GPU: 0) Std. dev. of avg. energy used [-]	0.368	0.013	0.023	0.005
(GPU: 1) Avg. energy used [kJ]	26.921	5.523	2.917	2.587
(GPU: 1) Std. dev. of avg. energy used [-]	0.025	0.01	0.006	0.004
(GPU: 2) Avg. energy used [kJ]	14.725	3.034	1.571	1.383
(GPU: 2) Std. dev. of avg. energy used [-]	0.126	0.025	0.028	0.016
(GPU: 3) Avg. energy used [kJ]	18.139	3.723	1.967	1.744
(GPU: 3) Std. dev. of avg. energy used [-]	0.023	0.006	0.002	0.003
(GPU: 4) Avg. energy used [kJ]	4.334	0.892	0.471	0.415
(GPU: 4) Std. dev. of avg. energy used [-]	0.011	0.005	0.002	0.001
(GPU: 5) Avg. energy used [kJ]	19.467	3.974	2.092	1.859
(GPU: 5) Std. dev. of avg. energy used [-]	0.043	0.021	0.013	0.011
(GPU: 6) Avg. energy used [kJ]	6.63	1.363	0.726	0.654
(GPU: 6) Std. dev. of avg. energy used [-]	0.078	0.016	0.017	0.008
(GPU: 7) Avg. energy used [kJ]	12.841	2.631	1.38	1.23
(GPU: 7) Std. dev. of avg. energy used [-]	0.334	0.031	0.036	0.032

Table 4.3. server: **sanna.kask**, device: **2 CPUs**, implementation: **OMP-CPP**,
benchmark: **bt.C**, data displayed: **power draw**

	2 CPUs			
Results from 10 runs	2 Threads	10 Threads	20 Threads	40 Threads
Avg. Exec. time [s]	955.334	114.56	69.86	64.855
Std. dev. of time [-]	1.891	1.843	0.395	0.979
(Yokogawa) Avg. power draw [W]	382.827	432.746	477.802	500.032
(Yokogawa) Std. dev. of avg. power draw [-]	1.469	0.575	1.18	2.708
(CPU: 0) Avg. power draw [W]	33.397	51.044	65.793	71.718
(CPU: 0) Std. dev. of avg. power draw [-]	0.336	0.188	0.175	0.372
(CPU: 1) Avg. power draw [W]	31.426	49.813	65.932	71.781
(CPU: 1) Std. dev. of avg. power draw [-]	1.188	0.239	0.166	0.358
(GPU: 0) Avg. power draw [W]	21.634	21.663	21.595	21.714
(GPU: 0) Std. dev. of avg. power draw [-]	0.105	0.115	0.064	0.068
(GPU: 1) Avg. power draw [W]	25.556	25.553	25.577	25.549
(GPU: 1) Std. dev. of avg. power draw [-]	0.039	0.029	0.047	0.025
(GPU: 2) Avg. power draw [W]	14.01	13.63	13.701	13.732
(GPU: 2) Std. dev. of avg. power draw [-]	0.172	0.153	0.134	0.132
(GPU: 3) Avg. power draw [W]	17.213	17.181	17.207	17.204
(GPU: 3) Std. dev. of avg. power draw [-]	0.012	0.012	0.021	0.022
(GPU: 4) Avg. power draw [W]	4.091	4.11	4.108	4.105
(GPU: 4) Std. dev. of avg. power draw [-]	0.006	0.025	0.01	0.017
(GPU: 5) Avg. power draw [W]	18.414	18.358	18.574	18.639
(GPU: 5) Std. dev. of avg. power draw [-]	0.058	0.062	0.127	0.089
(GPU: 6) Avg. power draw [W]	6.437	6.457	6.484	6.507
(GPU: 6) Std. dev. of avg. power draw [-]	0.189	0.11	0.103	0.17
(GPU: 7) Avg. power draw [W]	12.281	12.136	11.891	11.98
(GPU: 7) Std. dev. of avg. power draw [-]	0.362	0.253	0.158	0.135

Table 4.4. server: **sanna.kask**, device: **2 CPUs**, implementation: **OMP-CPP**,
benchmark: **bt.C**, data displayed: **energy used**

	2 CPUs			
Results from 10 runs	2 Threads	10 Threads	20 Threads	40 Threads
Avg. Exec. time [s]	955.334	114.56	69.86	64.855
Std. dev. of time [-]	1.891	1.843	0.395	0.979
(Yokogawa) Avg. energy used [kJ]	365.73	49.575	33.379	32.428
(Yokogawa) Std. dev. of avg. energy used [-]	1.935	0.739	0.188	0.39
(CPU: 0) Avg. energy used [kJ]	28.779	5.834	4.578	4.63
(CPU: 0) Std. dev. of avg. energy used [-]	3.718	0.084	0.018	0.05
(CPU: 1) Avg. energy used [kJ]	26.958	5.693	4.587	4.634
(CPU: 1) Std. dev. of avg. energy used [-]	2.456	0.07	0.024	0.05
(GPU: 0) Avg. energy used [kJ]	20.668	2.482	1.509	1.408
(GPU: 0) Std. dev. of avg. energy used [-]	0.108	0.05	0.007	0.022
(GPU: 1) Avg. energy used [kJ]	24.415	2.927	1.787	1.657
(GPU: 1) Std. dev. of avg. energy used [-]	0.055	0.046	0.011	0.024
(GPU: 2) Avg. energy used [kJ]	13.384	1.562	0.957	0.891
(GPU: 2) Std. dev. of avg. energy used [-]	0.163	0.034	0.012	0.016
(GPU: 3) Avg. energy used [kJ]	16.444	1.968	1.202	1.116
(GPU: 3) Std. dev. of avg. energy used [-]	0.032	0.031	0.006	0.017
(GPU: 4) Avg. energy used [kJ]	3.908	0.471	0.287	0.266
(GPU: 4) Std. dev. of avg. energy used [-]	0.006	0.006	0.002	0.004
(GPU: 5) Avg. energy used [kJ]	17.591	2.103	1.298	1.209
(GPU: 5) Std. dev. of avg. energy used [-]	0.064	0.036	0.007	0.021
(GPU: 6) Avg. energy used [kJ]	6.149	0.74	0.453	0.422
(GPU: 6) Std. dev. of avg. energy used [-]	0.179	0.018	0.007	0.008
(GPU: 7) Avg. energy used [kJ]	11.733	1.39	0.831	0.777
(GPU: 7) Std. dev. of avg. energy used [-]	0.351	0.028	0.011	0.014

Table 4.5. server: **sanna.kask**, device: **1 CPU**, implementation: **OMP-CPP**,
benchmark: **is.D**, data displayed: **power draw**

	1 CPU			
Results from 10 runs	1 Thread	5 Threads	10 Threads	20 Threads
Avg. Exec. time [s]	649	143.321	76.25	57.389
Std. dev. of time [-]	0.277	0.062	0.26	0.056
(Yokogawa) Avg. power draw [W]	378.666	394.555	413.561	426.861
(Yokogawa) Std. dev. of avg. power draw [-]	0.076	0.278	0.28	0.464
(CPU: 0) Avg. power draw [W]	32.558	42.525	56.774	63.154
(CPU: 0) Std. dev. of avg. power draw [-]	0.032	5.14	0.01	0.048
(CPU: 1) Avg. power draw [W]	28.52	28.607	29.026	29.05
(CPU: 1) Std. dev. of avg. power draw [-]	0.135	1.489	0.077	0.062
(GPU: 0) Avg. power draw [W]	21.648	21.72	21.707	21.791
(GPU: 0) Std. dev. of avg. power draw [-]	0.056	0.063	0.04	0.076
(GPU: 1) Avg. power draw [W]	25.545	25.563	25.548	25.504
(GPU: 1) Std. dev. of avg. power draw [-]	0.024	0.028	0.046	0.026
(GPU: 2) Avg. power draw [W]	13.861	13.718	13.501	13.791
(GPU: 2) Std. dev. of avg. power draw [-]	0.161	0.142	0.148	0.166
(GPU: 3) Avg. power draw [W]	17.227	17.235	17.209	17.225
(GPU: 3) Std. dev. of avg. power draw [-]	0.022	0.023	0.013	0.017
(GPU: 4) Avg. power draw [W]	4.104	4.112	4.103	4.084
(GPU: 4) Std. dev. of avg. power draw [-]	0.016	0.013	0.017	0.02
(GPU: 5) Avg. power draw [W]	18.345	18.483	18.409	18.436
(GPU: 5) Std. dev. of avg. power draw [-]	0.068	0.103	0.088	0.131
(GPU: 6) Avg. power draw [W]	6.488	6.221	6.268	6.298
(GPU: 6) Std. dev. of avg. power draw [-]	0.157	0.074	0.095	0.116
(GPU: 7) Avg. power draw [W]	12.351	12.097	12.134	12.576
(GPU: 7) Std. dev. of avg. power draw [-]	0.119	0.098	0.175	0.138

Table 4.6. server: **sanna.kask**, device: **1 CPU**, implementation: **OMP-CPP**,
benchmark: **is.D**, data displayed: **energy used**

	1 CPU			
Results from 10 runs	1 Thread	5 Threads	10 Threads	20 Threads
Avg. Exec. time [s]	649	143.321	76.25	57.389
Std. dev. of time [-]	0.277	0.062	0.26	0.056
(Yokogawa) Avg. energy used [kJ]	245.754	56.548	31.534	24.497
(Yokogawa) Std. dev. of avg. energy used [-]	0.101	0.03	0.09	0.024
(CPU: 0) Avg. energy used [kJ]	21.122	7.194	4.314	3.607
(CPU: 0) Std. dev. of avg. energy used [-]	0.028	2.777	0.014	0.004
(CPU: 1) Avg. energy used [kJ]	18.502	5.063	2.206	1.659
(CPU: 1) Std. dev. of avg. energy used [-]	0.085	2.857	0.005	0.004
(GPU: 0) Avg. energy used [kJ]	14.05	3.113	1.655	1.251
(GPU: 0) Std. dev. of avg. energy used [-]	0.034	0.009	0.008	0.004
(GPU: 1) Avg. energy used [kJ]	16.579	3.664	1.948	1.464
(GPU: 1) Std. dev. of avg. energy used [-]	0.014	0.004	0.007	0.001
(GPU: 2) Avg. energy used [kJ]	8.996	1.966	1.03	0.792
(GPU: 2) Std. dev. of avg. energy used [-]	0.103	0.021	0.013	0.009
(GPU: 3) Avg. energy used [kJ]	11.18	2.47	1.312	0.988
(GPU: 3) Std. dev. of avg. energy used [-]	0.016	0.004	0.004	0.001
(GPU: 4) Avg. energy used [kJ]	2.663	0.589	0.313	0.234
(GPU: 4) Std. dev. of avg. energy used [-]	0.01	0.002	0.002	0.001
(GPU: 5) Avg. energy used [kJ]	11.906	2.649	1.403	1.058
(GPU: 5) Std. dev. of avg. energy used [-]	0.04	0.015	0.008	0.008
(GPU: 6) Avg. energy used [kJ]	4.211	0.892	0.478	0.362
(GPU: 6) Std. dev. of avg. energy used [-]	0.103	0.011	0.007	0.007
(GPU: 7) Avg. energy used [kJ]	8.016	1.734	0.925	0.722
(GPU: 7) Std. dev. of avg. energy used [-]	0.078	0.014	0.014	0.008

Table 4.7. server: **sanna.kask**, device: **2 CPUs**, implementation: **OMP-CPP**,
benchmark: **is.D**, data displayed: **power draw**

	2 CPUs			
Results from 10 runs	2 Threads	10 Threads	20 Threads	40 Threads
Avg. Exec. time [s]	393.075	86.848	47.879	38.223
Std. dev. of time [-]	7.964	2.783	1.263	1.377
(Yokogawa) Avg. power draw [W]	383.887	414.22	447.238	470.243
(Yokogawa) Std. dev. of avg. power draw [-]	0.326	0.689	1.146	2.189
(CPU: 0) Avg. power draw [W]	32.17	43.594	54.718	60.222
(CPU: 0) Std. dev. of avg. power draw [-]	0.18	0.266	0.449	0.635
(CPU: 1) Avg. power draw [W]	32.6	42.639	53.835	58.977
(CPU: 1) Std. dev. of avg. power draw [-]	0.235	0.237	0.284	0.623
(GPU: 0) Avg. power draw [W]	21.695	21.572	21.577	21.557
(GPU: 0) Std. dev. of avg. power draw [-]	0.083	0.06	0.072	0.072
(GPU: 1) Avg. power draw [W]	25.538	25.531	25.514	25.52
(GPU: 1) Std. dev. of avg. power draw [-]	0.026	0.029	0.025	0.027
(GPU: 2) Avg. power draw [W]	13.802	13.757	13.892	13.832
(GPU: 2) Std. dev. of avg. power draw [-]	0.167	0.152	0.138	0.167
(GPU: 3) Avg. power draw [W]	17.229	17.23	17.227	17.225
(GPU: 3) Std. dev. of avg. power draw [-]	0.013	0.016	0.017	0.021
(GPU: 4) Avg. power draw [W]	4.108	4.076	4.101	4.095
(GPU: 4) Std. dev. of avg. power draw [-]	0.014	0.017	0.023	0.005
(GPU: 5) Avg. power draw [W]	18.309	18.404	18.385	18.488
(GPU: 5) Std. dev. of avg. power draw [-]	0.124	0.102	0.093	0.077
(GPU: 6) Avg. power draw [W]	6.199	6.356	6.41	6.282
(GPU: 6) Std. dev. of avg. power draw [-]	0.053	0.067	0.056	0.077
(GPU: 7) Avg. power draw [W]	12.377	12.525	12.336	12.424
(GPU: 7) Std. dev. of avg. power draw [-]	0.183	0.183	0.21	0.163

Table 4.8. server: **sanna.kask**, device: **2 CPUs**, implementation: **OMP-CPP**,
benchmark: **is.D**, data displayed: **energy used**

	2 CPUs			
Results from 10 runs	2 Threads	10 Threads	20 Threads	40 Threads
Avg. Exec. time [s]	393.075	86.848	47.879	38.223
Std. dev. of time [-]	7.964	2.783	1.263	1.377
(Yokogawa) Avg. energy used [kJ]	150.896	35.972	21.412	17.972
(Yokogawa) Std. dev. of avg. energy used [-]	3.034	1.099	0.523	0.585
(CPU: 0) Avg. energy used [kJ]	12.799	3.774	2.603	2.279
(CPU: 0) Std. dev. of avg. energy used [-]	0.628	0.103	0.051	0.06
(CPU: 1) Avg. energy used [kJ]	12.968	3.691	2.561	2.232
(CPU: 1) Std. dev. of avg. energy used [-]	0.605	0.102	0.058	0.06
(GPU: 0) Avg. energy used [kJ]	8.528	1.873	1.033	0.824
(GPU: 0) Std. dev. of avg. energy used [-]	0.176	0.062	0.027	0.03
(GPU: 1) Avg. energy used [kJ]	10.038	2.217	1.222	0.976
(GPU: 1) Std. dev. of avg. energy used [-]	0.2	0.072	0.033	0.035
(GPU: 2) Avg. energy used [kJ]	5.426	1.195	0.665	0.529
(GPU: 2) Std. dev. of avg. energy used [-]	0.141	0.037	0.018	0.019
(GPU: 3) Avg. energy used [kJ]	6.772	1.496	0.825	0.658
(GPU: 3) Std. dev. of avg. energy used [-]	0.137	0.049	0.022	0.024
(GPU: 4) Avg. energy used [kJ]	1.615	0.354	0.196	0.157
(GPU: 4) Std. dev. of avg. energy used [-]	0.034	0.012	0.005	0.006
(GPU: 5) Avg. energy used [kJ]	7.198	1.598	0.88	0.707
(GPU: 5) Std. dev. of avg. energy used [-]	0.189	0.053	0.022	0.026
(GPU: 6) Avg. energy used [kJ]	2.437	0.552	0.307	0.24
(GPU: 6) Std. dev. of avg. energy used [-]	0.048	0.017	0.009	0.009
(GPU: 7) Avg. energy used [kJ]	4.865	1.088	0.591	0.475
(GPU: 7) Std. dev. of avg. energy used [-]	0.115	0.043	0.021	0.018

Table 4.9. server: **sanna.kask**, device: **1 CPU**, implementation: **OMP-CPP**,
benchmark: **lu.C**, data displayed: **power draw**

	1 CPU			
Results from 10 runs	1 Thread	5 Threads	10 Threads	20 Threads
Avg. Exec. time [s]	706.566	147.528	76.747	64.501
Std. dev. of time [-]	5.233	1.844	0.169	0.421
(Yokogawa) Avg. power draw [W]	383.095	417.551	456.064	470.299
(Yokogawa) Std. dev. of avg. power draw [-]	1.031	8.528	29.324	8.099
(CPU: 0) Avg. power draw [W]	34.852	54.436	74.175	82.297
(CPU: 0) Std. dev. of avg. power draw [-]	0.064	0.167	0.091	0.128
(CPU: 1) Avg. power draw [W]	28.756	29.406	30	29.246
(CPU: 1) Std. dev. of avg. power draw [-]	0.133	0.542	1.638	0.638
(GPU: 0) Avg. power draw [W]	21.946	26.237	31.805	27.565
(GPU: 0) Std. dev. of avg. power draw [-]	0.789	6.809	24.367	6.578
(GPU: 1) Avg. power draw [W]	25.496	25.515	25.487	25.474
(GPU: 1) Std. dev. of avg. power draw [-]	0.027	0.026	0.027	0.031
(GPU: 2) Avg. power draw [W]	13.741	13.592	13.864	13.625
(GPU: 2) Std. dev. of avg. power draw [-]	0.206	0.142	0.209	0.127
(GPU: 3) Avg. power draw [W]	17.222	17.26	17.271	17.28
(GPU: 3) Std. dev. of avg. power draw [-]	0.023	0.019	0.022	0.024
(GPU: 4) Avg. power draw [W]	4.108	4.101	4.116	4.127
(GPU: 4) Std. dev. of avg. power draw [-]	0.013	0.012	0.008	0.018
(GPU: 5) Avg. power draw [W]	18.188	18.372	18.366	18.332
(GPU: 5) Std. dev. of avg. power draw [-]	0.1	0.118	0.069	0.094
(GPU: 6) Avg. power draw [W]	6.44	6.528	6.402	6.444
(GPU: 6) Std. dev. of avg. power draw [-]	0.142	0.125	0.101	0.092
(GPU: 7) Avg. power draw [W]	12.915	12.736	12.772	12.863
(GPU: 7) Std. dev. of avg. power draw [-]	0.131	0.132	0.198	0.143

Table 4.10. server: **sanna.kask**, device: **1 CPU**, implementation: **OMP-CPP**,
benchmark: **lu.C**, data displayed: **energy used**

	1 CPU			
Results from 10 runs	1 Thread	5 Threads	10 Threads	20 Threads
Avg. Exec. time [s]	706.566	147.528	76.747	64.501
Std. dev. of time [-]	5.233	1.844	0.169	0.421
(Yokogawa) Avg. energy used [kJ]	270.681	61.597	35.005	30.335
(Yokogawa) Std. dev. of avg. energy used [-]	2.117	1.321	2.319	0.554
(CPU: 0) Avg. energy used [kJ]	24.511	8.016	5.673	5.286
(CPU: 0) Std. dev. of avg. energy used [-]	0.34	0.078	0.014	0.03
(CPU: 1) Avg. energy used [kJ]	20.224	4.33	2.295	1.879
(CPU: 1) Std. dev. of avg. energy used [-]	0.35	0.093	0.13	0.051
(GPU: 0) Avg. energy used [kJ]	15.507	3.868	2.444	1.778
(GPU: 0) Std. dev. of avg. energy used [-]	0.572	0.991	1.88	0.424
(GPU: 1) Avg. energy used [kJ]	18.015	3.764	1.956	1.643
(GPU: 1) Std. dev. of avg. energy used [-]	0.129	0.048	0.003	0.011
(GPU: 2) Avg. energy used [kJ]	9.709	2.005	1.064	0.879
(GPU: 2) Std. dev. of avg. energy used [-]	0.174	0.027	0.016	0.009
(GPU: 3) Avg. energy used [kJ]	12.168	2.546	1.325	1.115
(GPU: 3) Std. dev. of avg. energy used [-]	0.093	0.03	0.004	0.008
(GPU: 4) Avg. energy used [kJ]	2.903	0.605	0.316	0.266
(GPU: 4) Std. dev. of avg. energy used [-]	0.024	0.008	0.001	0.002
(GPU: 5) Avg. energy used [kJ]	12.851	2.71	1.41	1.182
(GPU: 5) Std. dev. of avg. energy used [-]	0.11	0.04	0.007	0.004
(GPU: 6) Avg. energy used [kJ]	4.55	0.963	0.492	0.415
(GPU: 6) Std. dev. of avg. energy used [-]	0.105	0.022	0.008	0.007
(GPU: 7) Avg. energy used [kJ]	9.125	1.879	0.98	0.83
(GPU: 7) Std. dev. of avg. energy used [-]	0.088	0.024	0.015	0.007

Table 4.11. server: **sanna.kask**, device: **2 CPUs**, implementation: **OMP-CPP**,
benchmark: **lu.C**, data displayed: **power draw**

	2 CPUs			
Results from 10 runs	2 Threads	10 Threads	20 Threads	40 Threads
Avg. Exec. time [s]	599.96	75.879	46.189	43.611
Std. dev. of time [-]	2.984	0.893	0.889	0.981
(Yokogawa) Avg. power draw [W]	385.213	437.88	484.794	502.661
(Yokogawa) Std. dev. of avg. power draw [-]	1.288	0.628	2.355	2.661
(CPU: 0) Avg. power draw [W]	32.9	52.479	67.311	71.26
(CPU: 0) Std. dev. of avg. power draw [-]	0.397	0.208	0.387	0.434
(CPU: 1) Avg. power draw [W]	33.085	51.276	67.364	71.366
(CPU: 1) Std. dev. of avg. power draw [-]	0.657	0.266	0.473	0.458
(GPU: 0) Avg. power draw [W]	21.581	21.577	21.593	21.517
(GPU: 0) Std. dev. of avg. power draw [-]	0.156	0.065	0.09	0.08
(GPU: 1) Avg. power draw [W]	25.514	25.488	25.473	25.461
(GPU: 1) Std. dev. of avg. power draw [-]	0.028	0.022	0.035	0.025
(GPU: 2) Avg. power draw [W]	13.522	13.952	13.853	14.04
(GPU: 2) Std. dev. of avg. power draw [-]	0.171	0.184	0.116	0.193
(GPU: 3) Avg. power draw [W]	17.235	17.224	17.25	17.225
(GPU: 3) Std. dev. of avg. power draw [-]	0.037	0.014	0.022	0.015
(GPU: 4) Avg. power draw [W]	4.095	4.085	4.09	4.083
(GPU: 4) Std. dev. of avg. power draw [-]	0.046	0.011	0.017	0.013
(GPU: 5) Avg. power draw [W]	18.365	18.626	18.685	18.529
(GPU: 5) Std. dev. of avg. power draw [-]	0.12	0.17	0.116	0.09
(GPU: 6) Avg. power draw [W]	6.19	6.092	6.173	6.306
(GPU: 6) Std. dev. of avg. power draw [-]	0.174	0.12	0.094	0.128
(GPU: 7) Avg. power draw [W]	12.276	12.53	12.253	12.635
(GPU: 7) Std. dev. of avg. power draw [-]	0.227	0.148	0.308	0.149

Table 4.12. server: **sanna.kask**, device: **2 CPUs**, implementation: **OMP-CPP**,
benchmark: **lu.C**, data displayed: **energy used**

	2 CPUs			
Results from 10 runs	2 Threads	10 Threads	20 Threads	40 Threads
Avg. Exec. time [s]	599.96	75.879	46.189	43.611
Std. dev. of time [-]	2.984	0.893	0.889	0.981
(Yokogawa) Avg. energy used [kJ]	231.113	33.225	22.391	21.919
(Yokogawa) Std. dev. of avg. energy used [-]	1.38	0.366	0.334	0.396
(CPU: 0) Avg. energy used [kJ]	19.668	3.968	3.089	3.084
(CPU: 0) Std. dev. of avg. energy used [-]	2.786	0.043	0.045	0.052
(CPU: 1) Avg. energy used [kJ]	19.754	3.877	3.09	3.088
(CPU: 1) Std. dev. of avg. energy used [-]	2.606	0.034	0.043	0.051
(GPU: 0) Avg. energy used [kJ]	12.948	1.637	0.997	0.938
(GPU: 0) Std. dev. of avg. energy used [-]	0.124	0.019	0.021	0.022
(GPU: 1) Avg. energy used [kJ]	15.307	1.934	1.176	1.11
(GPU: 1) Std. dev. of avg. energy used [-]	0.082	0.022	0.022	0.025
(GPU: 2) Avg. energy used [kJ]	8.113	1.059	0.64	0.612
(GPU: 2) Std. dev. of avg. energy used [-]	0.131	0.02	0.014	0.016
(GPU: 3) Avg. energy used [kJ]	10.34	1.307	0.797	0.751
(GPU: 3) Std. dev. of avg. energy used [-]	0.057	0.015	0.015	0.017
(GPU: 4) Avg. energy used [kJ]	2.457	0.31	0.189	0.178
(GPU: 4) Std. dev. of avg. energy used [-]	0.029	0.003	0.004	0.004
(GPU: 5) Avg. energy used [kJ]	11.018	1.413	0.863	0.808
(GPU: 5) Std. dev. of avg. energy used [-]	0.092	0.025	0.02	0.018
(GPU: 6) Avg. energy used [kJ]	3.714	0.462	0.285	0.275
(GPU: 6) Std. dev. of avg. energy used [-]	0.102	0.009	0.005	0.008
(GPU: 7) Avg. energy used [kJ]	7.365	0.951	0.566	0.551
(GPU: 7) Std. dev. of avg. energy used [-]	0.127	0.016	0.02	0.012

Table 4.13. server: **sanna.kask**, device: **GPUs**, implementation: **OMP-CUDA**,
benchmark: **ep.D**, data displayed: **power draw**

	GPUs			
Results from 10 runs	1 GPU	2 GPUs	4 GPUs	8 GPUs
Avg. Exec. time [s]	28.154	28.336	28.342	28.956
Std. dev. of time [-]	0.148	0.035	0.15	0.132
(Yokogawa) Avg. power draw [W]	547.031	703.097	1287.656	1946.823
(Yokogawa) Std. dev. of avg. power draw [-]	7.345	4.882	20.937	15.778
(CPU: 0) Avg. power draw [W]	33.478	38.279	38.121	45.37
(CPU: 0) Std. dev. of avg. power draw [-]	0.043	0.038	0.035	0.135
(CPU: 1) Avg. power draw [W]	28.434	28.669	36.312	44.512
(CPU: 1) Std. dev. of avg. power draw [-]	0.139	0.123	0.194	0.116
(GPU: 0) Avg. power draw [W]	158.814	159.994	155.383	152.417
(GPU: 0) Std. dev. of avg. power draw [-]	1.964	0.503	0.74	0.532
(GPU: 1) Avg. power draw [W]	25.339	163.717	160.382	157.506
(GPU: 1) Std. dev. of avg. power draw [-]	0.04	2.372	0.874	0.548
(GPU: 2) Avg. power draw [W]	14.108	14.047	163.474	160.36
(GPU: 2) Std. dev. of avg. power draw [-]	0.189	0.106	1.832	1.365
(GPU: 3) Avg. power draw [W]	17.211	17.21	176.651	175.437
(GPU: 3) Std. dev. of avg. power draw [-]	0.016	0.023	2.46	0.978
(GPU: 4) Avg. power draw [W]	4.114	4.155	4.234	154.105
(GPU: 4) Std. dev. of avg. power draw [-]	0.014	0.014	0.022	1.687
(GPU: 5) Avg. power draw [W]	18.475	18.543	19.306	168.136
(GPU: 5) Std. dev. of avg. power draw [-]	0.101	0.051	0.132	1.009
(GPU: 6) Avg. power draw [W]	6.194	6.378	7.139	159.608
(GPU: 6) Std. dev. of avg. power draw [-]	0.08	0.12	0.441	1.287
(GPU: 7) Avg. power draw [W]	12.047	12.126	13.2	176.508
(GPU: 7) Std. dev. of avg. power draw [-]	0.183	0.122	0.177	2.815

Table 4.14. server: **sanna.kask**, device: **GPUs**, implementation: **OMP-CUDA**, benchmark: **ep.D**, data displayed: **energy used**

	GPUs			
Results from 10 runs	1 GPU	2 GPUs	4 GPUs	8 GPUs
Avg. Exec. time [s]	28.154	28.336	28.342	28.956
Std. dev. of time [-]	0.148	0.035	0.15	0.132
(Yokogawa) Avg. energy used [kJ]	15.401	19.923	36.497	56.373
(Yokogawa) Std. dev. of avg. energy used [-]	0.192	0.132	0.74	0.596
(CPU: 0) Avg. energy used [kJ]	0.934	1.075	1.071	1.302
(CPU: 0) Std. dev. of avg. energy used [-]	0.005	0.001	0.006	0.006
(CPU: 1) Avg. energy used [kJ]	0.793	0.805	1.02	1.277
(CPU: 1) Std. dev. of avg. energy used [-]	0.008	0.004	0.009	0.005
(GPU: 0) Avg. energy used [kJ]	4.471	4.532	4.404	4.413
(GPU: 0) Std. dev. of avg. energy used [-]	0.078	0.011	0.041	0.01
(GPU: 1) Avg. energy used [kJ]	0.713	4.639	4.546	4.561
(GPU: 1) Std. dev. of avg. energy used [-]	0.004	0.067	0.047	0.022
(GPU: 2) Avg. energy used [kJ]	0.397	0.398	4.633	4.643
(GPU: 2) Std. dev. of avg. energy used [-]	0.006	0.003	0.033	0.041
(GPU: 3) Avg. energy used [kJ]	0.485	0.488	5.007	5.08
(GPU: 3) Std. dev. of avg. energy used [-]	0.003	0.001	0.05	0.046
(GPU: 4) Avg. energy used [kJ]	0.116	0.118	0.12	4.462
(GPU: 4) Std. dev. of avg. energy used [-]	0.001	0.001	0.001	0.046
(GPU: 5) Avg. energy used [kJ]	0.52	0.525	0.547	4.868
(GPU: 5) Std. dev. of avg. energy used [-]	0.005	0.002	0.005	0.031
(GPU: 6) Avg. energy used [kJ]	0.174	0.181	0.202	4.621
(GPU: 6) Std. dev. of avg. energy used [-]	0.002	0.004	0.013	0.042
(GPU: 7) Avg. energy used [kJ]	0.339	0.344	0.374	5.111
(GPU: 7) Std. dev. of avg. energy used [-]	0.005	0.004	0.006	0.09

Table 4.15. server: **sanna.kask**, device: **GPUs**, implementation: **OMP-CUDA**,
benchmark: **lu.D**, data displayed: **power draw**

	GPUs			
Results from 10 runs	1 GPU	2 GPUs	4 GPUs	8 GPUs
Avg. Exec. time [s]	303.873	303.892	303.94	304.673
Std. dev. of time [-]	0.07	0.089	0.073	0.177
(Yokogawa) Avg. power draw [W]	605.116	806.718	1456.046	2453.513
(Yokogawa) Std. dev. of avg. power draw [-]	1.115	4.324	34.057	27.56
(CPU: 0) Avg. power draw [W]	33.388	37.623	37.77	46.243
(CPU: 0) Std. dev. of avg. power draw [-]	0.038	0.317	0.44	0.149
(CPU: 1) Avg. power draw [W]	29.156	29.271	37.316	45.244
(CPU: 1) Std. dev. of avg. power draw [-]	0.135	0.081	0.089	0.071
(GPU: 0) Avg. power draw [W]	222.687	221.965	208.846	207.528
(GPU: 0) Std. dev. of avg. power draw [-]	1.532	0.87	1.528	0.83
(GPU: 1) Avg. power draw [W]	25.209	227.456	215.767	213.007
(GPU: 1) Std. dev. of avg. power draw [-]	0.04	1.971	0.837	0.207
(GPU: 2) Avg. power draw [W]	13.17	13.093	218.461	216.842
(GPU: 2) Std. dev. of avg. power draw [-]	0.126	0.099	0.602	0.545
(GPU: 3) Avg. power draw [W]	17.246	17.321	222.286	220.981
(GPU: 3) Std. dev. of avg. power draw [-]	0.025	0.043	0.449	0.55
(GPU: 4) Avg. power draw [W]	4.075	4.103	4.35	214.034
(GPU: 4) Std. dev. of avg. power draw [-]	0.015	0.023	0.043	0.409
(GPU: 5) Avg. power draw [W]	18.557	18.811	20.909	214.9
(GPU: 5) Std. dev. of avg. power draw [-]	0.11	0.117	0.144	0.215
(GPU: 6) Avg. power draw [W]	8.246	8.769	8.016	220.018
(GPU: 6) Std. dev. of avg. power draw [-]	0.198	0.087	0.13	0.542
(GPU: 7) Avg. power draw [W]	13.041	13.062	13.302	245.404
(GPU: 7) Std. dev. of avg. power draw [-]	0.141	0.162	0.171	1.431

Table 4.16. server: **sanna.kask**, device: **GPUs**, implementation: **OMP-CUDA**,
benchmark: **lu.D**, data displayed: **energy used**

	GPUs			
Results from 10 runs	1 GPU	2 GPUs	4 GPUs	8 GPUs
Avg. Exec. time [s]	303.873	303.892	303.94	304.673
Std. dev. of time [-]	0.07	0.089	0.073	0.177
(Yokogawa) Avg. energy used [kJ]	183.878	245.155	442.551	747.518
(Yokogawa) Std. dev. of avg. energy used [-]	0.355	1.311	10.395	8.234
(CPU: 0) Avg. energy used [kJ]	10.137	11.423	11.47	14.076
(CPU: 0) Std. dev. of avg. energy used [-]	0.013	0.096	0.132	0.044
(CPU: 1) Avg. energy used [kJ]	8.852	8.887	11.332	13.772
(CPU: 1) Std. dev. of avg. energy used [-]	0.041	0.025	0.028	0.026
(GPU: 0) Avg. energy used [kJ]	67.669	67.453	63.477	63.228
(GPU: 0) Std. dev. of avg. energy used [-]	0.467	0.266	0.466	0.271
(GPU: 1) Avg. energy used [kJ]	7.66	69.122	65.58	64.897
(GPU: 1) Std. dev. of avg. energy used [-]	0.012	0.597	0.257	0.048
(GPU: 2) Avg. energy used [kJ]	4.002	3.979	66.399	66.066
(GPU: 2) Std. dev. of avg. energy used [-]	0.039	0.03	0.176	0.182
(GPU: 3) Avg. energy used [kJ]	5.241	5.264	67.562	67.327
(GPU: 3) Std. dev. of avg. energy used [-]	0.007	0.014	0.147	0.167
(GPU: 4) Avg. energy used [kJ]	1.238	1.247	1.322	65.21
(GPU: 4) Std. dev. of avg. energy used [-]	0.004	0.007	0.013	0.116
(GPU: 5) Avg. energy used [kJ]	5.639	5.717	6.355	65.474
(GPU: 5) Std. dev. of avg. energy used [-]	0.033	0.035	0.044	0.068
(GPU: 6) Avg. energy used [kJ]	2.506	2.665	2.436	67.033
(GPU: 6) Std. dev. of avg. energy used [-]	0.06	0.027	0.039	0.156
(GPU: 7) Avg. energy used [kJ]	3.963	3.969	4.043	74.768
(GPU: 7) Std. dev. of avg. energy used [-]	0.043	0.049	0.052	0.437

Table 4.17. server: **sanna.kask**, device: **GPUs**, implementation: **OMP-CUDA**,
benchmark: **sp.D**, data displayed: **power draw**

	GPUs			
Results from 10 runs	1 GPU	2 GPUs	4 GPUs	8 GPUs
Avg. Exec. time [s]	222.319	222.368	222.136	222.79
Std. dev. of time [-]	0.128	0.122	0.073	0.215
(Yokogawa) Avg. power draw [W]	604.859	807.122	1488.43	2479.957
(Yokogawa) Std. dev. of avg. power draw [-]	3.025	4.083	33.983	20.783
(CPU: 0) Avg. power draw [W]	33.366	37.119	37.595	46.159
(CPU: 0) Std. dev. of avg. power draw [-]	0.056	0.204	0.106	0.1
(CPU: 1) Avg. power draw [W]	29.143	29.373	37.426	45.29
(CPU: 1) Std. dev. of avg. power draw [-]	0.094	0.056	0.1	0.068
(GPU: 0) Avg. power draw [W]	220.497	222.542	207.888	207.154
(GPU: 0) Std. dev. of avg. power draw [-]	2.315	1.705	0.74	0.613
(GPU: 1) Avg. power draw [W]	25.442	226.649	215.391	213.297
(GPU: 1) Std. dev. of avg. power draw [-]	0.038	3.073	0.363	1.211
(GPU: 2) Avg. power draw [W]	13.469	13.427	216.762	215.364
(GPU: 2) Std. dev. of avg. power draw [-]	0.183	0.188	1.363	0.905
(GPU: 3) Avg. power draw [W]	17.244	17.247	221.942	221.593
(GPU: 3) Std. dev. of avg. power draw [-]	0.025	0.025	0.962	0.575
(GPU: 4) Avg. power draw [W]	4.05	4.028	4.061	211.718
(GPU: 4) Std. dev. of avg. power draw [-]	0.012	0.022	0.014	0.595
(GPU: 5) Avg. power draw [W]	18.49	18.685	20.532	213.672
(GPU: 5) Std. dev. of avg. power draw [-]	0.167	0.078	0.147	0.188
(GPU: 6) Avg. power draw [W]	7.626	7.912	7.129	219.677
(GPU: 6) Std. dev. of avg. power draw [-]	0.181	0.089	0.173	1.284
(GPU: 7) Avg. power draw [W]	13.469	13.329	14.158	244.845
(GPU: 7) Std. dev. of avg. power draw [-]	0.237	0.152	0.158	1.533

Table 4.18. server: **sanna.kask**, device: **GPUs**, implementation: **OMP-CUDA**,
benchmark: **sp.D**, data displayed: **energy used**

	GPUs			
Results from 10 runs	1 GPU	2 GPUs	4 GPUs	8 GPUs
Avg. Exec. time [s]	222.319	222.368	222.136	222.79
Std. dev. of time [-]	0.128	0.122	0.073	0.215
(Yokogawa) Avg. energy used [kJ]	134.471	179.478	330.635	552.51
(Yokogawa) Std. dev. of avg. energy used [-]	0.599	0.912	7.618	4.649
(CPU: 0) Avg. energy used [kJ]	7.409	8.244	8.341	10.271
(CPU: 0) Std. dev. of avg. energy used [-]	0.016	0.045	0.023	0.023
(CPU: 1) Avg. energy used [kJ]	6.471	6.524	8.304	10.078
(CPU: 1) Std. dev. of avg. energy used [-]	0.024	0.014	0.023	0.019
(GPU: 0) Avg. energy used [kJ]	49.021	49.486	46.18	46.152
(GPU: 0) Std. dev. of avg. energy used [-]	0.54	0.386	0.172	0.139
(GPU: 1) Avg. energy used [kJ]	5.656	50.4	47.846	47.52
(GPU: 1) Std. dev. of avg. energy used [-]	0.011	0.696	0.089	0.266
(GPU: 2) Avg. energy used [kJ]	2.994	2.986	48.151	47.981
(GPU: 2) Std. dev. of avg. energy used [-]	0.04	0.041	0.298	0.212
(GPU: 3) Avg. energy used [kJ]	3.834	3.835	49.302	49.369
(GPU: 3) Std. dev. of avg. energy used [-]	0.007	0.005	0.224	0.145
(GPU: 4) Avg. energy used [kJ]	0.9	0.896	0.902	47.169
(GPU: 4) Std. dev. of avg. energy used [-]	0.002	0.005	0.003	0.119
(GPU: 5) Avg. energy used [kJ]	4.111	4.155	4.561	47.604
(GPU: 5) Std. dev. of avg. energy used [-]	0.036	0.017	0.034	0.032
(GPU: 6) Avg. energy used [kJ]	1.695	1.759	1.584	48.942
(GPU: 6) Std. dev. of avg. energy used [-]	0.04	0.02	0.038	0.296
(GPU: 7) Avg. energy used [kJ]	2.994	2.964	3.145	54.549
(GPU: 7) Std. dev. of avg. energy used [-]	0.051	0.033	0.036	0.353

Table 4.19. server: **sanna.kask**, device: **Hybrid**, implementation: **OMP-CPP+OMP-CUDA**, benchmark: **bt.C+lu.D**, data displayed: **power draw**

	Hybrid			
Results from 10 runs (GPU(s) + Threads)	1 + 4	2 + 8	4 + 16	8 + 32
Avg. Exec. time [s]	269.755	139.845	86.92	82.297
Std. dev. of time [-]	2.918	0.269	0.441	0.692
(Yokogawa) Avg. power draw [W]	627.96	849.046	1677.74	2594.956
(Yokogawa) Std. dev. of avg. power draw [-]	2.936	7.87	6.929	14.026
(CPU: 0) Avg. power draw [W]	51.318	70.893	66.181	73.708
(CPU: 0) Std. dev. of avg. power draw [-]	0.379	0.092	0.157	0.171
(CPU: 1) Avg. power draw [W]	29.779	29.812	66.803	74.48
(CPU: 1) Std. dev. of avg. power draw [-]	0.087	0.031	0.159	0.204
(GPU: 0) Avg. power draw [W]	219.962	218.631	204.386	203.447
(GPU: 0) Std. dev. of avg. power draw [-]	1.83	0.458	0.149	0.612
(GPU: 1) Avg. power draw [W]	25.054	223.668	212.302	209.995
(GPU: 1) Std. dev. of avg. power draw [-]	0.051	3.104	0.471	0.608
(GPU: 2) Avg. power draw [W]	13.377	13.227	214.191	213.455
(GPU: 2) Std. dev. of avg. power draw [-]	0.225	0.133	1.027	1.34
(GPU: 3) Avg. power draw [W]	17.233	17.291	221.448	219.9
(GPU: 3) Std. dev. of avg. power draw [-]	0.023	0.03	0.359	0.426
(GPU: 4) Avg. power draw [W]	4.059	4.119	4.562	209.92
(GPU: 4) Std. dev. of avg. power draw [-]	0.016	0.013	0.058	1.43
(GPU: 5) Avg. power draw [W]	18.469	18.811	21.059	212.269
(GPU: 5) Std. dev. of avg. power draw [-]	0.164	0.103	0.228	0.875
(GPU: 6) Avg. power draw [W]	7.945	8.335	8.52	216.049
(GPU: 6) Std. dev. of avg. power draw [-]	0.12	0.183	0.513	1.779
(GPU: 7) Avg. power draw [W]	12.656	12.513	13.283	239.658
(GPU: 7) Std. dev. of avg. power draw [-]	0.195	0.097	0.229	2.208

Table 4.20. server: **sanna.kask**, device: **Hybrid**, implementation: **OMP-CPP+OMP-CUDA**, benchmark: **bt.C+lu.D**, data displayed: **energy used**

	Hybrid			
Results from 10 runs (GPU(s) + Threads)	1 + 4	2 + 8	4 + 16	8 + 32
Avg. Exec. time [s]	269.755	139.845	86.92	82.297
Std. dev. of time [-]	2.918	0.269	0.441	0.692
(Yokogawa) Avg. energy used [kJ]	169.403	118.735	145.829	213.561
(Yokogawa) Std. dev. of avg. energy used [-]	2.612	1.185	0.977	2.49
(CPU: 0) Avg. energy used [kJ]	13.829	9.895	5.734	6.043
(CPU: 0) Std. dev. of avg. energy used [-]	0.087	0.015	0.021	0.043
(CPU: 1) Avg. energy used [kJ]	8.025	4.161	5.787	6.106
(CPU: 1) Std. dev. of avg. energy used [-]	0.069	0.01	0.022	0.04
(GPU: 0) Avg. energy used [kJ]	59.331	30.574	17.765	16.743
(GPU: 0) Std. dev. of avg. energy used [-]	0.176	0.077	0.092	0.126
(GPU: 1) Avg. energy used [kJ]	6.759	31.278	18.453	17.282
(GPU: 1) Std. dev. of avg. energy used [-]	0.069	0.413	0.092	0.176
(GPU: 2) Avg. energy used [kJ]	3.608	1.85	18.618	17.566
(GPU: 2) Std. dev. of avg. energy used [-]	0.054	0.02	0.165	0.137
(GPU: 3) Avg. energy used [kJ]	4.649	2.418	19.248	18.039
(GPU: 3) Std. dev. of avg. energy used [-]	0.055	0.004	0.097	0.165
(GPU: 4) Avg. energy used [kJ]	1.095	0.576	0.397	17.276
(GPU: 4) Std. dev. of avg. energy used [-]	0.012	0.002	0.005	0.18
(GPU: 5) Avg. energy used [kJ]	4.982	2.63	1.83	17.469
(GPU: 5) Std. dev. of avg. energy used [-]	0.082	0.018	0.018	0.188
(GPU: 6) Avg. energy used [kJ]	2.143	1.166	0.74	17.78
(GPU: 6) Std. dev. of avg. energy used [-]	0.049	0.026	0.043	0.174
(GPU: 7) Avg. energy used [kJ]	3.414	1.75	1.154	19.723
(GPU: 7) Std. dev. of avg. energy used [-]	0.056	0.014	0.017	0.228

Table 4.21. server: **sanna.kask**, device: **Hybrid**, implementation: **OMP-CPP+OMP-CUDA**, benchmark: **is.D+sp.D**, data displayed: **power draw**

	Hybrid			
Results from 10 runs (GPU(s) + Threads)	1 + 4	2 + 8	4 + 16	8 + 32
Avg. Exec. time [s]	177.298	92.61	57.096	43.868
Std. dev. of time [-]	0.426	0.051	1.986	0.606
(Yokogawa) Avg. power draw [W]	625.727	840.728	1666.487	2503.797
(Yokogawa) Std. dev. of avg. power draw [-]	4.184	11.648	17.348	15.631
(CPU: 0) Avg. power draw [W]	45.545	60.354	58.849	64.799
(CPU: 0) Std. dev. of avg. power draw [-]	0.049	0.116	0.464	0.23
(CPU: 1) Avg. power draw [W]	29.725	29.755	57.727	65.174
(CPU: 1) Std. dev. of avg. power draw [-]	0.078	0.054	0.433	0.213
(GPU: 0) Avg. power draw [W]	219.519	214.766	205.621	199.845
(GPU: 0) Std. dev. of avg. power draw [-]	2.656	1.592	0.436	1.919
(GPU: 1) Avg. power draw [W]	25.354	219.389	209.693	205.12
(GPU: 1) Std. dev. of avg. power draw [-]	0.044	3.004	0.836	1.347
(GPU: 2) Avg. power draw [W]	13.039	12.923	212.433	207.042
(GPU: 2) Std. dev. of avg. power draw [-]	0.185	0.144	0.886	1.318
(GPU: 3) Avg. power draw [W]	17.176	17.179	220.225	216.365
(GPU: 3) Std. dev. of avg. power draw [-]	0.012	0.016	0.637	1.066
(GPU: 4) Avg. power draw [W]	4.037	4.041	4.12	203.852
(GPU: 4) Std. dev. of avg. power draw [-]	0.021	0.01	0.024	1.066
(GPU: 5) Avg. power draw [W]	18.261	18.421	20.406	208.75
(GPU: 5) Std. dev. of avg. power draw [-]	0.107	0.078	0.296	0.487
(GPU: 6) Avg. power draw [W]	7.83	7.712	7.5	209.678
(GPU: 6) Std. dev. of avg. power draw [-]	0.443	0.148	0.19	1.71
(GPU: 7) Avg. power draw [W]	12.549	12.679	14.03	231.716
(GPU: 7) Std. dev. of avg. power draw [-]	0.296	0.191	0.153	2.379

Table 4.22. server: **sanna.kask**, device: **Hybrid**, implementation: **OMP-CPP+OMP-CUDA**, benchmark: **is.D+sp.D**, data displayed: **energy used**

	Hybrid			
Results from 10 runs (GPU(s) + Threads)	1 + 4	2 + 8	4 + 16	8 + 32
Avg. Exec. time [s]	177.298	92.61	57.096	43.868
Std. dev. of time [-]	0.426	0.051	1.986	0.606
(Yokogawa) Avg. energy used [kJ]	110.94	77.86	95.151	109.841
(Yokogawa) Std. dev. of avg. energy used [-]	0.8	1.084	3.509	1.936
(CPU: 0) Avg. energy used [kJ]	8.063	5.574	3.342	2.819
(CPU: 0) Std. dev. of avg. energy used [-]	0.022	0.01	0.094	0.032
(CPU: 1) Avg. energy used [kJ]	5.262	2.748	3.278	2.835
(CPU: 1) Std. dev. of avg. energy used [-]	0.014	0.006	0.091	0.034
(GPU: 0) Avg. energy used [kJ]	38.92	19.89	11.74	8.767
(GPU: 0) Std. dev. of avg. energy used [-]	0.468	0.153	0.41	0.174
(GPU: 1) Avg. energy used [kJ]	4.495	20.318	11.974	8.998
(GPU: 1) Std. dev. of avg. energy used [-]	0.012	0.276	0.451	0.132
(GPU: 2) Avg. energy used [kJ]	2.312	1.197	12.129	9.083
(GPU: 2) Std. dev. of avg. energy used [-]	0.033	0.014	0.423	0.161
(GPU: 3) Avg. energy used [kJ]	3.045	1.591	12.574	9.492
(GPU: 3) Std. dev. of avg. energy used [-]	0.008	0.002	0.458	0.125
(GPU: 4) Avg. energy used [kJ]	0.716	0.374	0.235	8.942
(GPU: 4) Std. dev. of avg. energy used [-]	0.004	0.001	0.008	0.112
(GPU: 5) Avg. energy used [kJ]	3.238	1.706	1.165	9.158
(GPU: 5) Std. dev. of avg. energy used [-]	0.02	0.008	0.037	0.141
(GPU: 6) Avg. energy used [kJ]	1.388	0.714	0.428	9.198
(GPU: 6) Std. dev. of avg. energy used [-]	0.078	0.014	0.016	0.114
(GPU: 7) Avg. energy used [kJ]	2.225	1.174	0.801	10.164
(GPU: 7) Std. dev. of avg. energy used [-]	0.054	0.018	0.025	0.143

Table 4.23. server: **sanna.kask**, device: **Hybrid**, implementation: **OMP-CPP+OMP-CUDA**, benchmark: **lu.C+ep.D**, data displayed: **power draw**

	Hybrid			
Results from 10 runs (GPU(s) + Threads)	1 + 4	2 + 8	4 + 16	8 + 32
Avg. Exec. time [s]	28.1	28.245	28.236	29.099
Std. dev. of time [-]	0.112	0.047	0.058	0.161
(Yokogawa) Avg. power draw [W]	577.48	749.685	1399.988	2031.805
(Yokogawa) Std. dev. of avg. power draw [-]	6.219	4.585	13.288	11.626
(CPU: 0) Avg. power draw [W]	53.683	73.773	68.65	71.074
(CPU: 0) Std. dev. of avg. power draw [-]	0.057	0.115	0.28	1.042
(CPU: 1) Avg. power draw [W]	29.696	29.57	69.472	71.574
(CPU: 1) Std. dev. of avg. power draw [-]	0.057	0.08	0.376	1.206
(GPU: 0) Avg. power draw [W]	157.086	156.864	155.278	152.018
(GPU: 0) Std. dev. of avg. power draw [-]	0.777	0.472	0.602	0.772
(GPU: 1) Avg. power draw [W]	25.395	160.977	159.713	156.332
(GPU: 1) Std. dev. of avg. power draw [-]	0.045	0.891	0.447	0.606
(GPU: 2) Avg. power draw [W]	13.055	13.084	160.666	158.022
(GPU: 2) Std. dev. of avg. power draw [-]	0.089	0.147	1.012	0.806
(GPU: 3) Avg. power draw [W]	17.235	17.221	175.985	172.448
(GPU: 3) Std. dev. of avg. power draw [-]	0.031	0.023	1.561	1.711
(GPU: 4) Avg. power draw [W]	4.046	4.057	4.085	152.676
(GPU: 4) Std. dev. of avg. power draw [-]	0.011	0.013	0.021	0.64
(GPU: 5) Avg. power draw [W]	18.334	18.328	19.013	166.557
(GPU: 5) Std. dev. of avg. power draw [-]	0.098	0.1	0.179	0.564
(GPU: 6) Avg. power draw [W]	6.381	6.478	7.533	158.864
(GPU: 6) Std. dev. of avg. power draw [-]	0.106	0.069	0.412	0.775
(GPU: 7) Avg. power draw [W]	12.034	12.483	13.108	173.497
(GPU: 7) Std. dev. of avg. power draw [-]	0.158	0.177	0.331	1.893

Table 4.24. server: **sanna.kask**, device: **Hybrid**, implementation: **OMP-CPP+OMP-CUDA**, benchmark: **lu.C+ep.D**, data displayed: **energy used**

	Hybrid			
Results from 10 runs (GPU(s) + Threads)	1 + 4	2 + 8	4 + 16	8 + 32
Avg. Exec. time [s]	28.1	28.245	28.236	29.099
Std. dev. of time [-]	0.112	0.047	0.058	0.161
(Yokogawa) Avg. energy used [kJ]	16.227	21.175	39.53	59.123
(Yokogawa) Std. dev. of avg. energy used [-]	0.15	0.116	0.379	0.509
(CPU: 0) Avg. energy used [kJ]	1.494	2.064	1.917	2.046
(CPU: 0) Std. dev. of avg. energy used [-]	0.006	0.006	0.007	0.036
(CPU: 1) Avg. energy used [kJ]	0.826	0.828	1.941	2.06
(CPU: 1) Std. dev. of avg. energy used [-]	0.004	0.002	0.009	0.04
(GPU: 0) Avg. energy used [kJ]	4.414	4.431	4.385	4.423
(GPU: 0) Std. dev. of avg. energy used [-]	0.037	0.016	0.016	0.015
(GPU: 1) Avg. energy used [kJ]	0.714	4.547	4.51	4.549
(GPU: 1) Std. dev. of avg. energy used [-]	0.004	0.027	0.016	0.022
(GPU: 2) Avg. energy used [kJ]	0.367	0.37	4.537	4.598
(GPU: 2) Std. dev. of avg. energy used [-]	0.003	0.004	0.029	0.032
(GPU: 3) Avg. energy used [kJ]	0.484	0.486	4.969	5.018
(GPU: 3) Std. dev. of avg. energy used [-]	0.003	0.001	0.04	0.049
(GPU: 4) Avg. energy used [kJ]	0.114	0.115	0.116	4.443
(GPU: 4) Std. dev. of avg. energy used [-]	0.001	0.001	0.001	0.03
(GPU: 5) Avg. energy used [kJ]	0.515	0.518	0.537	4.847
(GPU: 5) Std. dev. of avg. energy used [-]	0.003	0.003	0.005	0.037
(GPU: 6) Avg. energy used [kJ]	0.179	0.183	0.213	4.623
(GPU: 6) Std. dev. of avg. energy used [-]	0.003	0.002	0.012	0.043
(GPU: 7) Avg. energy used [kJ]	0.338	0.353	0.37	5.048
(GPU: 7) Std. dev. of avg. energy used [-]	0.004	0.005	0.01	0.076

Table 4.25. server: **vinnana.kask**, device: **CPUs**, implementation: **MPI-Fortran**,
benchmark: **ep.D.x**, data displayed: **power draw**

	CPUs			
Results from 10 runs	Processes: 4	Processes: 8	Processes: 16	Processes: 32
Avg. Exec. time [s]	379.715	193.936	99.675	71.101
Std. dev. of time [-]	0.925	4.385	2.088	0.2
(Yokogawa) Avg. power draw [W]	353.574	370.336	409.864	455.735
(Yokogawa) Std. dev. of avg. power draw [-]	5.535	0.238	0.788	2.487
(CPUs) Avg. power draw [W]	82.234	96.666	128.389	157.78
(CPUs) Std. dev. of avg. power draw [-]	4.737	0.092	0.227	0.206
(GPU: 0) Avg. power draw [W]	11.833	11.636	11.658	11.318
(GPU: 0) Std. dev. of avg. power draw [-]	0.946	0.04	0.154	0.055
(GPU: 1) Avg. power draw [W]	13.152	13.241	13.298	13.239
(GPU: 1) Std. dev. of avg. power draw [-]	0.111	0.048	0.065	0.057
(GPU: 2) Avg. power draw [W]	11.196	11.278	11.203	11.169
(GPU: 2) Std. dev. of avg. power draw [-]	0.048	0.061	0.062	0.026
(GPU: 3) Avg. power draw [W]	14.362	14.384	14.392	14.401
(GPU: 3) Std. dev. of avg. power draw [-]	0.038	0.023	0.036	0.033

Table 4.26. server: **vinnana.kask**, device: **CPUs**, implementation: **MPI-Fortran**,
benchmark: **ep.D.x**, data displayed: **energy used**

	CPUs			
Results from 10 runs	Processes: 4	Processes: 8	Processes: 16	Processes: 32
Avg. Exec. time [s]	379.715	193.936	99.675	71.101
Std. dev. of time [-]	0.925	4.385	2.088	0.2
(Yokogawa) Avg. energy used [kJ]	134.257	71.821	40.854	32.403
(Yokogawa) Std. dev. of avg. energy used [-]	2.081	1.612	0.903	0.213
(CPUs) Avg. energy used [kJ]	31.204	18.722	12.764	11.178
(CPUs) Std. dev. of avg. energy used [-]	1.792	0.437	0.287	0.036
(GPU: 0) Avg. energy used [kJ]	4.493	2.257	1.162	0.805
(GPU: 0) Std. dev. of avg. energy used [-]	0.355	0.05	0.022	0.004
(GPU: 1) Avg. energy used [kJ]	4.994	2.568	1.325	0.941
(GPU: 1) Std. dev. of avg. energy used [-]	0.045	0.065	0.027	0.006
(GPU: 2) Avg. energy used [kJ]	4.252	2.187	1.117	0.794
(GPU: 2) Std. dev. of avg. energy used [-]	0.026	0.05	0.026	0.003
(GPU: 3) Avg. energy used [kJ]	5.453	2.79	1.434	1.024
(GPU: 3) Std. dev. of avg. energy used [-]	0.009	0.063	0.028	0.004

Table 4.27. server: **vinnana.kask**, device: **CPUs**, implementation: **MPI-Fortran**,
benchmark: **is.D.x**, data displayed: **power draw**

	CPUs			
Results from 10 runs	Processes: 4	Processes: 8	Processes: 16	Processes: 32
Avg. Exec. time [s]	153.946	82.896	46.065	33.465
Std. dev. of time [-]	0.267	0.285	0.075	0.528
(Yokogawa) Avg. power draw [W]	350.311	367.628	402.346	443.225
(Yokogawa) Std. dev. of avg. power draw [-]	0.215	0.317	0.517	1.332
(CPUs) Avg. power draw [W]	77.438	90.402	114.976	138.504
(CPUs) Std. dev. of avg. power draw [-]	0.05	0.146	0.124	0.604
(GPU: 0) Avg. power draw [W]	11.371	11.286	11.318	11.403
(GPU: 0) Std. dev. of avg. power draw [-]	0.051	0.054	0.092	0.114
(GPU: 1) Avg. power draw [W]	13.185	13.225	13.349	13.328
(GPU: 1) Std. dev. of avg. power draw [-]	0.069	0.068	0.041	0.051
(GPU: 2) Avg. power draw [W]	11.303	11.238	11.266	11.35
(GPU: 2) Std. dev. of avg. power draw [-]	0.031	0.044	0.039	0.046
(GPU: 3) Avg. power draw [W]	14.482	14.401	14.439	14.463
(GPU: 3) Std. dev. of avg. power draw [-]	0.036	0.037	0.057	0.049

Table 4.28. server: **vinnana.kask**, device: **CPUs**, implementation: **MPI-Fortran**,
benchmark: **is.D.x**, data displayed: **energy used**

	CPUs			
Results from 10 runs	Processes: 4	Processes: 8	Processes: 16	Processes: 32
Avg. Exec. time [s]	153.946	82.896	46.065	33.465
Std. dev. of time [-]	0.267	0.285	0.075	0.528
(Yokogawa) Avg. energy used [kJ]	53.929	30.475	18.534	14.832
(Yokogawa) Std. dev. of avg. energy used [-]	0.08	0.124	0.041	0.203
(CPUs) Avg. energy used [kJ]	11.901	7.47	5.266	4.6
(CPUs) Std. dev. of avg. energy used [-]	0.022	0.036	0.009	0.089
(GPU: 0) Avg. energy used [kJ]	1.75	0.936	0.521	0.382
(GPU: 0) Std. dev. of avg. energy used [-]	0.009	0.007	0.004	0.008
(GPU: 1) Avg. energy used [kJ]	2.3	1.096	0.615	0.446
(GPU: 1) Std. dev. of avg. energy used [-]	0.011	0.006	0.002	0.007
(GPU: 2) Avg. energy used [kJ]	1.74	0.932	0.519	0.38
(GPU: 2) Std. dev. of avg. energy used [-]	0.006	0.004	0.002	0.007
(GPU: 3) Avg. energy used [kJ]	2.229	1.194	0.665	0.484
(GPU: 3) Std. dev. of avg. energy used [-]	0.006	0.006	0.003	0.008

Table 4.29. server: **vinnana.kask**, device: **CPUs**, implementation: **MPI-Fortran**,
benchmark: **lu.C.x**, data displayed: **power draw**

	CPUs			
Results from 10 runs	Processes: 4	Processes: 8	Processes: 16	Processes: 32
Avg. Exec. time [s]	167.99	92.314	48.768	43.598
Std. dev. of time [-]	0.711	0.738	0.596	0.365
(Yokogawa) Avg. power draw [W]	358.608	382.562	431.436	460.828
(Yokogawa) Std. dev. of avg. power draw [-]	0.189	2.012	1.193	1.996
(CPUs) Avg. power draw [W]	83.263	100.595	134.568	152.584
(CPUs) Std. dev. of avg. power draw [-]	0.057	0.159	0.942	0.296
(GPU: 0) Avg. power draw [W]	11.363	11.939	11.317	11.348
(GPU: 0) Std. dev. of avg. power draw [-]	0.095	1.472	0.075	0.138
(GPU: 1) Avg. power draw [W]	13.169	13.194	13.161	13.19
(GPU: 1) Std. dev. of avg. power draw [-]	0.078	0.13	0.107	0.042
(GPU: 2) Avg. power draw [W]	11.273	11.2	11.214	11.141
(GPU: 2) Std. dev. of avg. power draw [-]	0.03	0.055	0.048	0.039
(GPU: 3) Avg. power draw [W]	14.455	14.434	14.416	14.401
(GPU: 3) Std. dev. of avg. power draw [-]	0.048	0.039	0.031	0.034

Table 4.30. server: **vinnana.kask**, device: **CPUs**, implementation: **MPI-Fortran**,
benchmark: **lu.C.x**, data displayed: **energy used**

	CPUs			
Results from 10 runs	Processes: 4	Processes: 8	Processes: 16	Processes: 32
Avg. Exec. time [s]	167.99	92.314	48.768	43.598
Std. dev. of time [-]	0.711	0.738	0.596	0.365
(Yokogawa) Avg. energy used [kJ]	60.242	35.316	21.04	20.092
(Yokogawa) Std. dev. of avg. energy used [-]	0.255	0.331	0.238	0.226
(CPUs) Avg. energy used [kJ]	13.966	9.26	6.528	6.613
(CPUs) Std. dev. of avg. energy used [-]	0.059	0.066	0.124	0.063
(GPU: 0) Avg. energy used [kJ]	1.909	1.102	0.552	0.495
(GPU: 0) Std. dev. of avg. energy used [-]	0.022	0.138	0.008	0.009
(GPU: 1) Avg. energy used [kJ]	2.212	1.218	0.642	0.575
(GPU: 1) Std. dev. of avg. energy used [-]	0.014	0.017	0.009	0.005
(GPU: 2) Avg. energy used [kJ]	1.894	1.034	0.547	0.486
(GPU: 2) Std. dev. of avg. energy used [-]	0.01	0.007	0.008	0.004
(GPU: 3) Avg. energy used [kJ]	2.428	1.332	0.703	0.628
(GPU: 3) Std. dev. of avg. energy used [-]	0.014	0.012	0.008	0.006

Table 4.31. server: **vinnana.kask**, device: **GPUs**, implementation: **Horovod-Python**,
benchmark: **Xception**, data displayed: **power draw**

	GPUs		
Results from 10 runs	1 GPU	2 GPUs	4 GPUs
Avg. Exec. time [s]	133.363	80.633	56.574
Std. dev. of time [-]	0.403	0.36	0.35
(Yokogawa) Avg. power draw [W]	511.645	665.957	905.241
(Yokogawa) Std. dev. of avg. power draw [-]	1.12	2.813	2.909
(CPUs) Avg. power draw [W]	55.354	67.363	81.37
(CPUs) Std. dev. of avg. power draw [-]	0.074	0.203	0.283
(GPU: 0) Avg. power draw [W]	169.563	156.733	135.377
(GPU: 0) Std. dev. of avg. power draw [-]	0.878	0.832	1.089
(GPU: 1) Avg. power draw [W]	14.416	159.318	138.517
(GPU: 1) Std. dev. of avg. power draw [-]	0.097	1.436	1.123
(GPU: 2) Avg. power draw [W]	11.06	10.967	133.84
(GPU: 2) Std. dev. of avg. power draw [-]	0.054	0.12	1.084
(GPU: 3) Avg. power draw [W]	14.637	14.873	137.662
(GPU: 3) Std. dev. of avg. power draw [-]	0.046	0.091	1.053

Table 4.32. server: **vinnana.kask**, device: **GPUs**, implementation: **Horovod-Python**,
benchmark: **Xception**, data displayed: **energy used**

	GPUs		
Results from 10 runs	1 GPU	2 GPUs	4 GPUs
Avg. Exec. time [s]	133.363	80.633	56.574
Std. dev. of time [-]	0.403	0.36	0.35
(Yokogawa) Avg. energy used [kJ]	68.234	53.697	51.213
(Yokogawa) Std. dev. of avg. energy used [-]	0.268	0.18	0.268
(CPUs) Avg. energy used [kJ]	7.368	5.414	4.582
(CPUs) Std. dev. of avg. energy used [-]	0.023	0.028	0.028
(GPU: 0) Avg. energy used [kJ]	22.614	12.638	7.659
(GPU: 0) Std. dev. of avg. energy used [-]	0.178	0.036	0.044
(GPU: 1) Avg. energy used [kJ]	1.923	12.846	7.836
(GPU: 1) Std. dev. of avg. energy used [-]	0.018	0.092	0.037
(GPU: 2) Avg. energy used [kJ]	1.475	0.884	7.572
(GPU: 2) Std. dev. of avg. energy used [-]	0.007	0.01	0.07
(GPU: 3) Avg. energy used [kJ]	1.952	1.199	7.788
(GPU: 3) Std. dev. of avg. energy used [-]	0.006	0.008	0.079

Table 4.33. server: **vinnana.kask**, device: **Hybrid**, implementation: **MPI-Fortran+Horovod-Python**, benchmark: **ep.D.x+XCeption**, data displayed: **power draw**

	Hybrid		
Results from 10 runs (GPU(s) + Processes)	1 + 8	2 + 16	4 + 32
Avg. Exec. time [s]	131.454	80.068	73.33
Std. dev. of time [-]	0.552	1.59	0.355
(Yokogawa) Avg. power draw [W]	561.049	741.187	915.957
(Yokogawa) Std. dev. of avg. power draw [-]	1.145	6.528	2.154
(CPUs) Avg. power draw [W]	99.713	135.506	163.343
(CPUs) Std. dev. of avg. power draw [-]	0.123	0.217	0.221
(GPU: 0) Avg. power draw [W]	171.163	156.899	112.283
(GPU: 0) Std. dev. of avg. power draw [-]	0.676	2.686	0.845
(GPU: 1) Avg. power draw [W]	14.324	159.284	115.304
(GPU: 1) Std. dev. of avg. power draw [-]	0.087	2.51	0.371
(GPU: 2) Avg. power draw [W]	11.058	11.096	111.088
(GPU: 2) Std. dev. of avg. power draw [-]	0.052	0.068	0.945
(GPU: 3) Avg. power draw [W]	14.67	14.96	115.475
(GPU: 3) Std. dev. of avg. power draw [-]	0.034	0.064	1.035

Table 4.34. server: **vinnana.kask**, device: **Hybrid**, implementation: **MPI-Fortran+Horovod-Python**, benchmark: **ep.D.x+Xception**, data displayed: **energy used**

	Hybrid		
Results from 10 runs (GPU(s) + Processes)	1 + 8	2 + 16	4 + 32
Avg. Exec. time [s]	131.454	80.068	73.33
Std. dev. of time [-]	0.552	1.59	0.355
(Yokogawa) Avg. energy used [kJ]	73.752	59.337	67.166
(Yokogawa) Std. dev. of avg. energy used [-]	0.343	0.737	0.229
(CPUs) Avg. energy used [kJ]	13.081	10.813	11.935
(CPUs) Std. dev. of avg. energy used [-]	0.056	0.204	0.063
(GPU: 0) Avg. energy used [kJ]	22.5	12.559	8.234
(GPU: 0) Std. dev. of avg. energy used [-]	0.164	0.052	0.056
(GPU: 1) Avg. energy used [kJ]	1.883	12.75	8.455
(GPU: 1) Std. dev. of avg. energy used [-]	0.015	0.123	0.04
(GPU: 2) Avg. energy used [kJ]	1.454	0.888	8.146
(GPU: 2) Std. dev. of avg. energy used [-]	0.01	0.021	0.064
(GPU: 3) Avg. energy used [kJ]	1.928	1.198	8.468
(GPU: 3) Std. dev. of avg. energy used [-]	0.011	0.023	0.078

Table 4.35. server: **vinnana.kask**, device: **Hybrid**, implementation: **MPI-Fortran+Horovod-Python**, benchmark: **is.D.x+XCeption**, data displayed: **power draw**

	Hybrid		
Results from 10 runs (GPU(s) + Processes)	1 + 8	2 + 16	4 + 32
Avg. Exec. time [s]	82.853	47.997	35.723
Std. dev. of time [-]	0.279	0.502	0.194
(Yokogawa) Avg. power draw [W]	547.1	697.771	720.439
(Yokogawa) Std. dev. of avg. power draw [-]	1.914	3.633	11.668
(CPUs) Avg. power draw [W]	94.054	123.423	147.79
(CPUs) Std. dev. of avg. power draw [-]	0.133	0.22	0.322
(GPU: 0) Avg. power draw [W]	158.512	138.724	71.034
(GPU: 0) Std. dev. of avg. power draw [-]	0.822	1.105	2.74
(GPU: 1) Avg. power draw [W]	14.358	141.267	73.891
(GPU: 1) Std. dev. of avg. power draw [-]	0.356	1.547	2.972
(GPU: 2) Avg. power draw [W]	11.212	11.13	69.662
(GPU: 2) Std. dev. of avg. power draw [-]	0.143	0.04	2.393
(GPU: 3) Avg. power draw [W]	14.924	14.884	74.406
(GPU: 3) Std. dev. of avg. power draw [-]	0.311	0.047	2.716

Table 4.36. server: **vinnana.kask**, device: **Hybrid**, implementation: **MPI-Fortran+Horovod-Python**, benchmark: **is.D.x+XCeption**, data displayed: **energy used**

	Hybrid		
Results from 10 runs (GPU(s) + Processes)	1 + 8	2 + 16	4 + 32
Avg. Exec. time [s]	82.853	47.997	35.723
Std. dev. of time [-]	0.279	0.502	0.194
(Yokogawa) Avg. energy used [kJ]	45.329	33.491	25.736
(Yokogawa) Std. dev. of avg. energy used [-]	0.245	0.395	0.437
(CPUs) Avg. energy used [kJ]	7.768	5.891	5.241
(CPUs) Std. dev. of avg. energy used [-]	0.029	0.063	0.024
(GPU: 0) Avg. energy used [kJ]	13.133	6.659	2.538
(GPU: 0) Std. dev. of avg. energy used [-]	0.088	0.106	0.101
(GPU: 1) Avg. energy used [kJ]	1.19	6.78	2.639
(GPU: 1) Std. dev. of avg. energy used [-]	0.03	0.095	0.106
(GPU: 2) Avg. energy used [kJ]	0.929	0.534	2.488
(GPU: 2) Std. dev. of avg. energy used [-]	0.012	0.006	0.087
(GPU: 3) Avg. energy used [kJ]	1.236	0.714	2.658
(GPU: 3) Std. dev. of avg. energy used [-]	0.026	0.008	0.097

Table 4.37. server: **vinnana.kask**, device: **Hybrid**, implementation: **MPI-Fortran+Horovod-Python**, benchmark: **lu.C.x+XCeption**, data displayed: **power draw**

	Hybrid		
Results from 10 runs (GPU(s) + Processes)	1 + 8	2 + 16	4 + 32
Avg. Exec. time [s]	92.854	57.626	52.643
Std. dev. of time [-]	0.297	0.266	0.329
(Yokogawa) Avg. power draw [W]	562.155	731.711	850.486
(Yokogawa) Std. dev. of avg. power draw [-]	2.046	3.034	10.848
(CPUs) Avg. power draw [W]	104.009	140.296	156.675
(CPUs) Std. dev. of avg. power draw [-]	0.097	0.152	0.243
(GPU: 0) Avg. power draw [W]	161.692	143.748	97.141
(GPU: 0) Std. dev. of avg. power draw [-]	0.719	0.98	2.403
(GPU: 1) Avg. power draw [W]	14.369	147.024	100.048
(GPU: 1) Std. dev. of avg. power draw [-]	0.081	1.612	2.599
(GPU: 2) Avg. power draw [W]	11.101	11.052	96.436
(GPU: 2) Std. dev. of avg. power draw [-]	0.126	0.094	2.452
(GPU: 3) Avg. power draw [W]	14.789	14.979	101.289
(GPU: 3) Std. dev. of avg. power draw [-]	0.114	0.06	2.684

Table 4.38. server: **vinnana.kask**, device: **Hybrid**, implementation: **MPI-Fortran+Horovod-Python**, benchmark: **lu.C.x+XCeption**, data displayed: **energy used**

	Hybrid		
Results from 10 runs (GPU(s) + Processes)	1 + 8	2 + 16	4 + 32
Avg. Exec. time [s]	92.854	57.626	52.643
Std. dev. of time [-]	0.297	0.266	0.329
(Yokogawa) Avg. energy used [kJ]	52.199	42.166	44.773
(Yokogawa) Std. dev. of avg. energy used [-]	0.244	0.342	0.692
(CPUs) Avg. energy used [kJ]	9.63	8.047	8.207
(CPUs) Std. dev. of avg. energy used [-]	0.032	0.037	0.062
(GPU: 0) Avg. energy used [kJ]	15.014	8.284	5.114
(GPU: 0) Std. dev. of avg. energy used [-]	0.102	0.084	0.14
(GPU: 1) Avg. energy used [kJ]	1.334	8.472	5.267
(GPU: 1) Std. dev. of avg. energy used [-]	0.007	0.12	0.15
(GPU: 2) Avg. energy used [kJ]	1.031	0.637	5.077
(GPU: 2) Std. dev. of avg. energy used [-]	0.011	0.006	0.143
(GPU: 3) Avg. energy used [kJ]	1.373	0.863	5.332
(GPU: 3) Std. dev. of avg. energy used [-]	0.011	0.005	0.155

In order to obtain information about the changes of power draw measurements behaviour during the execution of benchmarks, a base values of idle nodes power draw should be gathered as a reference for further investigation. Idle nodes power draw has been measured as follows: each server has been reserved for the duration of tests, with no other processes running in the background. Then, using the Yokogawa WT310E Power Meter, the power draw of entire servers has been measured independently for the period of one hour. Additionally, the power draw of each CPU and each GPU has been measured as well, using Linux Perf and NVML respectively. Finally, the base value for further analysis has been calculated by subtracting the sum of idle CPUs and GPUs power draw from the idle node power draw. These values are presented in the tables below as a reference:

Table 4.39. Power draw measurements of idle *sanna.kask* server

	Idle power draw
Measured components	Power draw readings [W]
Sum of entire node idle power draw (Yokogawa)	357.575
Sum of CPUs idle power draw (Linux Perf / Intel RAPL)	44.908
Sum of GPUs idle power draw (NVML)	120.604
Difference between idle node power draw and sum of CPUs and GPUs power draw	192.064

Table 4.40. Power draw measurements of idle *vinnana.kask* server

	Idle power draw
Measured components	Power draw readings [W]
Sum of entire node idle power draw (Yokogawa)	314.693
Sum of CPUs idle power draw (Linux Perf / Intel RAPL)	46.967
Sum of GPUs idle power draw (NVML)	50.255
Difference between idle node power draw and sum of CPUs and GPUs power draw	217.471

Last set of tables presented on the next pages serve as a summary of differences in power draw between the values measured during the execution of benchmark kernels and the base value of idle node power draw. Explanation of data presented in summary tables is shown below:

- **Difference** – subtraction of active node and the sum of active CPUs and GPUs power draw [W]
- **Offset** – difference between the subtraction of active node and the sum of active CPUs and GPUs power draw [W] and subtraction of idle node and the sum of idle CPUs and GPUs power draw [W]
- **Increase** – percentage increase in node power draw during tests, compared to idle node power draw [%]

Table 4.41. server: **sanna.kask**, device: **1 CPU**, implementation: **OMP-CPP**,
benchmark: **bt.C**, data displayed: **differences in power draw**

	1 CPU			
Configurations	(Yokogawa) Measured power draw of entire node [W]	Difference (Base value: 192.064) [W]	Offset [W]	Increase [%]
1 CPU 1 Thread	379.962	197.808	5.744	6.26
1 CPU 5 Threads	402.881	203.322	11.258	12.67
1 CPU 10 Threads	432.171	213.958	21.894	20.86
1 CPU 20 Threads	445.752	219.968	27.904	24.66

Table 4.42. server: **sanna.kask**, device: **1 CPU**, implementation: **OMP-CPP**,
benchmark: **is.D**, data displayed: **differences in power draw**

	1 CPU			
Configurations	(Yokogawa) Measured power draw of entire node [W]	Difference (Base value: 192.064) [W]	Offset [W]	Increase [%]
1 CPU 1 Thread	378.666	198.019	5.955	5.90
1 CPU 5 Threads	394.555	204.274	12.21	10.34
1 CPU 10 Threads	413.561	208.882	16.818	15.66
1 CPU 20 Threads	426.861	214.952	22.888	19.38

Table 4.43. server: **sanna.kask**, device: **1 CPU**, implementation: **OMP-CPP**,
benchmark: **lu.C**, data displayed: **differences in power draw**

	1 CPU			
Configurations	(Yokogawa) Measured power draw of entire node [W]	Difference (Base value: 192.064) [W]	Offset [W]	Increase [%]
1 CPU 1 Thread	383.095	199.431	7.367	7.14
1 CPU 5 Threads	417.551	209.368	17.304	16.77
1 CPU 10 Threads	456.064	221.806	29.742	27.54
1 CPU 20 Threads	470.299	233.046	40.982	31.52

Table 4.44. server: **sanna.kask**, device: **2 CPUs**, implementation: **OMP-CPP**,
benchmark: **bt.C**, data displayed: **differences in power draw**

	2 CPUs			
Configurations	(Yokogawa) Measured power draw of entire node [W]	Difference (Base value: 192.064) [W]	Offset [W]	Increase [%]
2 CPUs 2 Threads	382.827	198.368	6.304	7.06
2 CPUs 10 Threads	432.746	212.801	20.737	21.02
2 CPUs 20 Threads	477.802	226.94	34.876	33.62
2 CPUs 40 Threads	500.032	237.103	45.039	39.84

Table 4.45. server: **sanna.kask**, device: **2 CPUs**, implementation: **OMP-CPP**,
benchmark: **is.D**, data displayed: **differences in power draw**

	2 CPUs			
Configurations	(Yokogawa) Measured power draw of entire node [W]	Difference (Base value: 192.064) [W]	Offset [W]	Increase [%]
2 CPUs 2 Threads	383.887	199.86	7.796	7.36
2 CPUs 10 Threads	414.22	208.536	16.472	15.84
2 CPUs 20 Threads	447.238	219.243	27.179	25.08
2 CPUs 40 Threads	470.243	231.621	39.557	31.51

Table 4.46. server: **sanna.kask**, device: **2 CPUs**, implementation: **OMP-CPP**,
benchmark: **lu.C**, data displayed: **differences in power draw**

	2 CPUs			
Configurations	(Yokogawa) Measured power draw of entire node [W]	Difference (Base value: 192.064) [W]	Offset [W]	Increase [%]
2 CPUs 2 Threads	385.213	200.45	8.386	7.73
2 CPUs 10 Threads	437.88	214.551	22.487	22.46
2 CPUs 20 Threads	484.794	230.749	38.685	35.58
2 CPUs 40 Threads	502.661	240.239	48.175	40.57

Table 4.47. server: **sanna.kask**, device: **GPUs**, implementation: **OMP-CUDA**, benchmark: **ep.D**, data displayed: **differences in power draw**

	GPUs			
Configurations	(Yokogawa) Measured power draw of entire node [W]	Difference (Base value: 192.064) [W]	Offset [W]	Increase [%]
1 GPU	547.031	228.817	36.753	52.98
2 GPUs	703.097	240.029	47.965	96.63
4 GPUs	1287.656	513.454	321.39	260.11
8 GPUs	1946.823	552.864	360.8	444.45

Table 4.48. server: **sanna.kask**, device: **GPUs**, implementation: **OMP-CUDA**, benchmark: **lu.D**, data displayed: **differences in power draw**

	GPUs			
Configurations	(Yokogawa) Measured power draw of entire node [W]	Difference (Base value: 192.064) [W]	Offset [W]	Increase [%]
1 GPU	605.116	220.341	28.277	69.23
2 GPUs	806.718	215.244	23.18	125.61
4 GPUs	1456.046	469.023	276.959	307.20
8 GPUs	2453.513	609.312	417.248	586.15

Table 4.49. server: **sanna.kask**, device: **GPUs**, implementation: **OMP-CUDA**, benchmark: **sp.D**, data displayed: **differences in power draw**

	GPUs			
Configurations	(Yokogawa) Measured power draw of entire node [W]	Difference (Base value: 192.064) [W]	Offset [W]	Increase [%]
1 GPU	604.859	222.063	29.999	69.16
2 GPUs	807.122	216.811	24.747	125.72
4 GPUs	1488.43	505.546	313.482	316.26
8 GPUs	2479.957	641.188	449.124	593.55

Table 4.50. server: **sanna.kask**, device: **Hybrid**, implementation: **OMP-CPP+OMP-CUDA**, benchmarks: **bt.C+lu.D**, data displayed: **differences in power draw**

	Hybrid			
Configurations	(Yokogawa) Measured power draw of entire node [W]	Difference (Base value: 192.064) [W]	Offset [W]	Increase [%]
1 GPU + 4 Threads	627.96	228.108	36.044	75.62
2 GPUs + 8 Threads	849.046	231.746	39.682	137.45
4 GPUs + 16 Threads	1677.74	645.005	452.941	369.20
8 GPUs + 32 Threads	2594.956	722.785	530.721	625.71

Table 4.51. server: **sanna.kask**, device: **Hybrid**, implementation: **OMP-CPP+OMP-CUDA**, benchmarks: **is.D+sp.D**, data displayed: **differences in power draw**

	Hybrid			
Configurations	(Yokogawa) Measured power draw of entire node [W]	Difference (Base value: 192.064) [W]	Offset [W]	Increase [%]
1 GPU + 4 Threads	625.727	232.692	40.628	74.99
2 GPUs + 8 Threads	840.728	243.509	51.445	135.12
4 GPUs + 16 Threads	1666.487	655.883	463.819	366.05
8 GPUs + 32 Threads	2503.797	691.456	499.392	600.22

Table 4.52. server: **sanna.kask**, device: **Hybrid**, implementation: **OMP-CPP+OMP-CUDA**, benchmarks: **lu.C+ep.D**, data displayed: **differences in power draw**

	Hybrid			
Configurations	(Yokogawa) Measured power draw of entire node [W]	Difference (Base value: 192.064) [W]	Offset [W]	Increase [%]
1 GPU + 4 Threads	577.48	240.535	48.471	61.50
2 GPUs + 8 Threads	749.685	256.85	64.786	109.66
4 GPUs + 16 Threads	1399.988	566.485	374.421	291.52
8 GPUs + 32 Threads	2031.805	598.743	406.679	468.22

Table 4.53. server: **vinnana.kask**, device: **CPUs**, implementation: **MPI-Fortran**, benchmark: **ep.D.x**, data displayed: **differences in power draw**

	CPUs			
Configurations	(Yokogawa) Measured power draw of entire node [W]	Difference (Base value: 217.471) [W]	Offset [W]	Increase [%]
Processes: 4	353.715	220.797	3.326	12.40
Processes: 8	370.336	223.131	5.66	17.68
Processes: 16	409.864	230.924	13.453	30.24
Processes: 32	455.735	247.828	30.357	44.82

Table 4.54. server: **vinnana.kask**, device: **CPUs**, implementation: **MPI-Fortran**, benchmark: **is.D.x**, data displayed: **differences in power draw**

	CPUs			
Configurations	(Yokogawa) Measured power draw of entire node [W]	Difference (Base value: 217.471) [W]	Offset [W]	Increase [%]
Processes: 4	350.311	222.532	5.061	11.32
Processes: 8	367.628	227.076	9.605	16.82
Processes: 16	402.346	236.998	19.527	27.85
Processes: 32	443.225	254.177	36.706	40.84

Table 4.55. server: **vinnana.kask**, device: **CPUs**, implementation: **MPI-Fortran**, benchmark: **lu.C.x**, data displayed: **differences in power draw**

	CPUs			
Configurations	(Yokogawa) Measured power draw of entire node [W]	Difference (Base value: 217.471) [W]	Offset [W]	Increase [%]
Processes: 4	358.608	225.085	7.614	13.95
Processes: 8	382.562	231.2	13.729	21.57
Processes: 16	431.436	246.76	29.289	37.10
Processes: 32	460.828	258.164	40.693	46.44

Table 4.56. server: **vinnana.kask**, device: **GPUs**, implementation: **Horovod-Python**, benchmark: **Xception**, data displayed: **differences in power draw**

	GPUs			
Configurations	(Yokogawa) Measured power draw of entire node [W]	Difference (Base value: 217.471) [W]	Offset [W]	Increase [%]
1 GPU	511.645	246.615	29.144	62.59
2 GPUs	665.957	256.703	39.232	111.62
4 GPUs	905.241	278.475	61.004	187.66

Table 4.57. server: **vinnana.kask**, device: **Hybrid**, implementation: **MPI-Fortran+Horovod-Python**, benchmark: **ep.D.x+Xception**, data displayed: **differences in power draw**

	Hybrid			
Configurations	(Yokogawa) Measured power draw of entire node [W]	Difference (Base value: 217.471) [W]	Offset [W]	Increase [%]
1 GPU + Processes: 8	561.049	250.121	232.65	78.28
2 GPUs + Processes: 16	741.187	263.442	45.971	135.53
4 GPUs + Processes: 32	915.957	298.464	80.993	191.06

Table 4.58. server: **vinnana.kask**, device: **Hybrid**, implementation: **MPI-Fortran+Horovod-Python**, benchmark: **is.D.x+Xception**, data displayed: **differences in power draw**

	Hybrid			
Configurations	(Yokogawa) Measured power draw of entire node [W]	Difference (Base value: 217.471) [W]	Offset [W]	Increase [%]
1 GPU + Processes: 8	547.1	254.04	36.569	73.85
2 GPUs + Processes: 16	697.771	268.343	50.872	121.73
4 GPUs + Processes: 32	720.439	283.656	66.185	128.93

Table 4.59. server: **vinnana.kask**, device: **Hybrid**, implementation: **MPI-Fortran+Horovod-Python**, benchmark: **lu.C.x+Xception**, data displayed: **differences in power draw**

	Hybrid			
Configurations	(Yokogawa) Measured power draw of entire node [W]	Difference (Base value: 217.471) [W]	Offset [W]	Increase [%]
1 GPU + Processes: 8	562.155	256.195	38.724	78.64
2 GPUs + Processes: 16	731.711	274.612	57.141	132.52
4 GPUs + Processes: 32	850.486	298.897	81.426	170.26

In this section, the conclusion are drawn based on the interpretation of summary tables of differences in power draw during various tests configurations.

- **CPUs benchmarks on *sanna.kask*** – based on the observation of power draw offsets, the conclusion is that the power draw of entire node grows almost linearly, the more logical processors are used for running the benchmarks. Biggest change occurs when the tests are started at all, in the *1 CPU 1 Thread* and *2 CPUs 2 Threads* configurations, and gets smaller, the more resources are dedicated to tests.
- **GPUs benchmarks on *sanna.kask*** – in this set of tests one can observe an interesting phenomenon. At first, situation seems similar to CPUs benchmarks in terms of near linear offset changes, when 1 or 2 GPUs are used for the tests. When benchmarks are executed on 4 or 8 GPUs, however, we can observe a significant surge in power draw of the entire node. This could be related to the increased activity of other components of the server, such as increased power draw of motherboard, increased fans speed or greater RAM usage.
- **Hybrid benchmarks on *sanna.kask*** – this set gives similar results to the GPUs-only set. The increase in power draw of the whole server, when executing benchmarks on 4 GPUs + 16 Threads and 8 GPUs + 32 Threads is even more visible. Again, the conclusions are similar to the analysis of the GPUs benchmarks.
- **CPUs benchmarks on *vinnana.kask*** – analysis of the results from this configurations confirm the observation of CPUs benchmarks on *sanna.kask*. This time, the offset of node power draw seems to be near linear as well, again with slightly bigger difference at lower processes count and decreasing the more processes are being spawned.
- **GPUs benchmarks on *vinnana.kask*** – interesting difference between these GPU benchmarks and the benchmarks executed on *sanna.kask* is fact, that the results of the tests performed on *vinnana.kask* are more similar to the CPU benchmarks. The power draw surge is not observed anymore, when executing the training on the configuration with 4 GPUs, leading to a more linear changes in offset and percentage increase of node power draw compared to idle values.
- **Hybrid benchmarks on *vinnana.kask*** – finally, the last set of configurations. While the values of offsets and percentage increases are higher than those in GPUs benchmarks, overall the change is without major changes between configurations, unlike the tests performed on *vinnana.kask*.

5. SUMMARY AND FUTURE WORK

5.1. SUMMARY

The goal of this master thesis was to conduct experiments and compare various methods of power draw measurements of computational nodes. The idea came from the fact, that the ICT sector holds a fair share in global electricity usage and this contribution will rapidly grow in the near future. Therefore, data centers responsible with major energy usage are more prone to shift their focus on the so-called *green computing*, where the performance of the computations meets the lower energy consumption. In order to achieve this, a precise power measurements tools are necessary, so any changes such as kernels optimization or power-capping can be verified as a proper solutions for lowering the energy usage.

For the purpose of the tests, two software measurements tools provided by the major CPUs and GPUs manufacturers were compared against external power meter in order to check their reliability. The software tools used are Intel RAPL, wrapped conveniently in Linux Perf toolkit for the measurements of the Intel CPUs and NVIDIA Management Library for the measurements of the NVIDIA GPUs. As a ground truth, the Yokogawa WT310E Power Meter has been used – certified and reliable source of the power measurements of the entire computational nodes.

Later in this work, the benchmarks were discussed – set of programs that are designed to help evaluate the performance of parallel supercomputers. Their main goal is to put the tested devices under strain and fully utilize their resources. After set of preliminary tests, benchmarks chosen for the final tests uses different implementations and solves various problems with different input sizes, all in order to recreate diverse computational problems that could be faced during the tests in data centers.

Next, the three research papers are presented as a part of the ‘related works’ section. This gives the insight on methodology used by the scientists that conducted similar experiments. Such research allows to learn more about the topic and gives hints on how the tests should be correctly executed.

For the clarity of the work, an exhausting overview on all of the configurations used in experiments was given. Every setup has been sorted in ordered manner: server, device, implementation, benchmark and configuration. This approach helps to get grasp on the used solutions and helps in organizing and writing the *scheduler* script, that was responsible for the automation of the experiments. The individual parts of the script has been then thoroughly explained, both the workflows of the application and the tasks performed by the implemented functions.

Finally, the measurements data has been presented in form of convenient tables. The first set of tables shows the power draw and energy consumption of both individual devices and the entire node as a whole. It gives a very precise informations about every components used during tests, with values accompanied with the standard deviations, that helps to describe how dispersed these sets of data are. For a better overview on the differences between power draw of active node during during tests and base values of power draw from idle node, a set of summary tables has been created. Based on the information they include, the analysis of the entire experiment has been done. As a final remark, the best solution for measurements of the power draw of the entire server is to use external power meter – this way one can be sure that estimates of the energy usage are as close to the real values as possible, due to the fact, that the power draw of components other than CPUs and GPUs are also included. The devices such as RAM sticks, motherboard, fans tends to increase their power usage, the more of the CPUs and GPUs

resources are being utilized during tests. In case of the software measurements tools, they seem to be very useful for gathering information about device specific power draw and by analysing the mentioned in previous chapter *offset* and *percentage increase in power draw during tests*, they behave predictably, showing rather linear increases, based on the more resources used. The only configurations that show concerns are the GPUs and Hybrid benchmarks executed on *sanna.kask*, where the surges of power draw could be observed in configurations with 4 and 8 GPUs. As of the reason for what causes such behaviour, further investigation would be required in the future works.

5.2. FUTURE WORK

In the terms of future work, this master thesis gives a strong insights for further expansion of the experiments.

One of the ideas is to perform tests on different node with more diverse components, such as different CPUs and GPUs, than those tested, to check if the results can be repeated and compared with already performed experiments. Other idea is to perform tests on both *sanna.kask* and *vinnana.kask* nodes simultaneously, using the MPI and Horovod based benchmarks, due to their inter-node communication. This approach would simulate running benchmarks on a HPC cluster. For the configurations of this setup, it is a sound strategy to run the tests utilising significant amount of the resources, such as tests on half of the devices and then, utilising all of them on both servers. This will avoid the repetition of already performed experiments and focus solely on researching the behaviour of the nodes under high computational strain.

Another good idea is to use Linux Perf to gather the information about energy consumption of the RAM of the servers. This would reduce the uncertainty of the measurements done by Yokogawa WT310E on the entire nodes, giving more informations about the behaviour of the servers power draw during tests, therefore allowing to receive more precise results and form even better conclusions.

Due to the fact, that many tests using different benchmarkshas has been performed on various configurations, starting from the single thread benchmarks, through the multi-CPU and multi-GPU and the finally hybrid benchmarks utilizing all the resources available on the servers, another solution has been proposed. Future goal is to create an empirical model, that will give insight about how much will increase the difference between readings from the Yokogawa external power meter and software solutions, such as Linux Perf and NVML. The purpose of such model would be rapidly checking the link of the entire node power draw and the sum of power draws of active CPUs and GPUs.

Other suggestion is to give more attentions to the PSUs used in order to power the servers. Many PSUs have different certifications, such as *Bronze*, *Silver*, *Gold*, *Platinum* and *Titanium*. Their main difference is their energy efficiencies at various load levels. The least efficient are the PSUs with *Bronze* certification, while the best are *Titanium* ones. This could discover potential differences in the readings of power draw of the servers and could also benefit in fine-tuning the previously mentioned empirical model.

As the culmination of this work, there is a high potential in writing a research paper that would contain the most insightful informations and conclusions made after the tests. As a motivation for this idea it should be noted, that this master thesis contains very exhausting experiments part, with various implementations, benchmarks and class sizes tested as well as meaningful conclusions drawn from the tests, than paired with previously mentioned ideas for the expansion of the experiments, can form a strong base for the research paper. The topic of energy aware computing

is currently very popular, due to increase of electricity use by the ICT sector world-wide, and the rising demand of the *green computing*, that has in mind finding the best ratio of performance to energy use. Research paper based on this work could contribute to this idea, due to verification of the software measurements method upon which the power draw of the computational nodes is estimated.

REFERENCES

- [1] Yehia Arafa et al. “Verified Instruction-Level Energy Consumption Measurement for NVIDIA GPUs”. In: *Proceedings of the 17th ACM International Conference on Computing Frontiers*. CF ’20. New York, NY, USA: Association for Computing Machinery, 2020, pp. 60–70. ISBN: 9781450379564. DOI: 10.1145/3387902.3392613. URL: <https://doi.org/10.1145/3387902.3392613>.
- [2] Gabriell Araujo et al. “NAS Parallel Benchmarks with CUDA and beyond”. In: *Software: Practice and Experience* 53.1 (2023), pp. 53–80. DOI: <https://doi.org/10.1002/spe.3056>. URL: <https://onlinelibrary.wiley.com/doi/abs/10.1002/spe.3056>.
- [3] Gabriell Alves de Araujo et al. “Efficient NAS Parallel Benchmark Kernels with CUDA”. In: *2020 28th Euromicro International Conference on Parallel, Distributed and Network-Based Processing (PDP)*. 2020, pp. 9–16. DOI: 10.1109/PDP50117.2020.00009.
- [4] Maria Avgerinou, Paolo Bertoldi, and Luca Castellazzi. “Trends in Data Centre Energy Consumption under the European Code of Conduct for Data Centre Energy Efficiency”. In: *Energies* 10.10 (2017). ISSN: 1996-1073. DOI: 10.3390/en10101470. URL: <https://www.mdpi.com/1996-1073/10/10/1470>.
- [5] David H. Bailey. “NAS Parallel Benchmarks”. In: *Encyclopedia of Parallel Computing*. Ed. by David Padua. Boston, MA: Springer US, 2011, pp. 1254–1259. ISBN: 978-0-387-09766-4. DOI: 10.1007/978-0-387-09766-4_133. URL: https://doi.org/10.1007/978-0-387-09766-4_133.
- [6] Artem Bityutskiy, Antti Laakso, and Helia Correia. *Yokotool*. URL: <https://github.com/intel/yoko-tool>.
- [7] Martin Burtscher. *K20Power v1.1*. URL: <https://userweb.cs.txstate.edu/~burtscher/research/K20power/>.
- [8] Martin Burtscher, Rupesh Nasre, and Keshav Pingali. “A quantitative study of irregular programs on GPUs”. In: *2012 IEEE International Symposium on Workload Characterization (IISWC)*. Nov. 2012, pp. 141–151. DOI: 10.1109/IISWC.2012.6402918.
- [9] Martin Burtscher, Ivan Zecena, and Ziliang Zong. “Measuring GPU Power with the K20 Built-in Sensor”. In: *Proceedings of Workshop on General Purpose Processing Using GPUs. GPGPU-7*. New York, NY, USA: Association for Computing Machinery, 2014, pp. 28–36. ISBN: 9781450327664. DOI: 10.1145/2588768.2576783. URL: <https://doi.org/10.1145/2588768.2576783>.
- [10] International Electrotechnical Commission. *IEC 62301:2011 | IEC Webstore | energy efficiency, smart city, standby power*. URL: <https://webstore.iec.ch/publication/6789>.
- [11] Howard David et al. “RAPL: Memory power estimation and capping”. In: *2010 ACM/IEEE International Symposium on Low-Power Electronics and Design (ISLPED)*. Aug. 2010, pp. 189–194. DOI: 10.1145/1840845.1840883.
- [12] Walter Dehnen. “A Hierarchical (N) Force Calculation Algorithm”. In: *Journal of Computational Physics* 179.1 (June 2002), pp. 27–42. DOI: 10.1006/jcph.2002.7026. URL: <https://doi.org/10.1006/jcph.2002.7026>.

- [13] The Institute for Development and Resources in Intensive Scientific Computing. *Horovod: Multi-GPU and multi-node data parallelism*. URL: <http://www.idris.fr/eng/jean-zay/gpu/jean-zay-gpu-hvd-tf-multi-eng.html>.
- [14] Electronic Educational Devices. *WattsUp? Operators Manual*. URL: <https://arcb.csc.ncsu.edu/~mueller/cluster/arc/wattsup/metertools-1.0.0/docs/meters/wattsup/manual.pdf>.
- [15] Muhammad Fahad et al. "A Comparative Study of Methods for Measurement of Energy of Computing". In: *Energies* 12.11 (2019). ISSN: 1996-1073. DOI: 10.3390/en12112204. URL: <https://www.mdpi.com/1996-1073/12/11/2204>.
- [16] Google. *Machine Learning Crash Course with TensorFlow APIs*. URL: <https://developers.google.com/machine-learning/crash-course>.
- [17] Intel. *Intel® Xeon Phi™ Coprocessor DEVELOPER'S QUICK START GUIDE*. URL: <https://www.intel.com/content/dam/develop/external/us/en/documents/intel-xeon-phi-coprocessor-quick-start-developers-guide.pdf>.
- [18] Intel. *Running Average Power Limit Energy Reporting*. URL: <https://www.intel.com/content/www/us/en/developer/articles/technical/software-security-guidance/advisory-guidance/running-average-power-limit-energy-reporting.html>.
- [19] Chao Jin et al. "A survey on software methods to improve the energy efficiency of parallel computing". In: *The International Journal of High Performance Computing Applications* 31.6 (2017), pp. 517–549. DOI: 10.1177/1094342016665471.
- [20] Kaggle. *Kaggle: Your Machine Learning and Data Science Community*. URL: <https://www.kaggle.com/>.
- [21] Michael Kerrisk. *Pthreads - Linux manual page*. URL: <https://man7.org/linux/man-pages/man7/pthreads.7.html>.
- [22] Hamidreza Khaleghzadeh et al. "Out-of-Core Implementation for Accelerator Kernels on Heterogeneous Clouds". In: *J. Supercomput.* 74.2 (Feb. 2018), pp. 551–568. ISSN: 0920-8542. DOI: 10.1007/s11227-017-2141-4. URL: <https://doi.org/10.1007/s11227-017-2141-4>.
- [23] Adam Krzywaniak and Pawel Czarnul. "Performance/Energy Aware Optimization of Parallel Applications on GPUs Under Power Capping". In: *Parallel Processing and Applied Mathematics*. Ed. by Roman Wyrzykowski et al. Cham: Springer International Publishing, 2020, pp. 123–133. ISBN: 978-3-030-43222-5. URL: https://link.springer.com/100005d4u0098.han.bg.pg.edu.pl/chapter/10.1007/978-3-030-43222-5_11.
- [24] Adam Krzywaniak, Jerzy Proficz, and Pawel Czarnul. "Analyzing energy/performance trade-offs with power capping for parallel applications on modern multi and many core processors". In: Sept. 2018, pp. 339–346. DOI: 10.15439/2018F177.
- [25] Linux. *"taskset" - Linux manual page*. URL: <https://man7.org/linux/man-pages/man1/taskset.1.html>.

- [26] Júnior Löff et al. "The NAS Parallel Benchmarks for evaluating C++ parallel programming frameworks on shared-memory architectures". In: *Future Generation Computer Systems* 125 (2021), pp. 743–757. ISSN: 0167-739X. DOI: <https://doi.org/10.1016/j.future.2021.07.021>. URL: <https://www.sciencedirect.com/science/article/pii/S0167739X21002831>.
- [27] NASA. *NAS Parallel Benchmarks*. URL: <https://www.nas.nasa.gov/software/npb.html>.
- [28] NVIDIA. *NVIDIA CUDA Compiler Driver NVCC*. URL: <https://docs.nvidia.com/cuda/cuda-compiler-driver-nvcc/>.
- [29] NVIDIA. *NVIDIA Management Library (NVML)*. URL: <https://developer.nvidia.com/nvidia-management-library-nvml>.
- [30] NVIDIA. *NVML API Reference Guide*. URL: <https://docs.nvidia.com/deploy/nvml-api/index.html>.
- [31] NVIDIA. *PAPI CUDA Component*. URL: <https://developer.nvidia.com/papi-cuda-component>.
- [32] NVIDIA. *Parallel Thread Execution ISA Version 8.2*. URL: <https://docs.nvidia.com/cuda/parallel-thread-execution/index.html>.
- [33] NVIDIA. *System Management Interface SMI*. URL: <https://developer.nvidia.com/nvidia-system-management-interface>.
- [34] OpenMP. *OpenMP Reference Guides*. URL: <https://www.openmp.org/resources/refguides/>.
- [35] Gerry Piosenka. *Birds 525 Species - Image Classification*. May 2023. URL: <https://www.kaggle.com/datasets/gpiosenska/100-bird-species>.
- [36] ICR Polska. *EnergyStar - ICR POLAND - testing and certification*. URL: <https://icrpolska.com/en/energystar/>.
- [37] Python. *"os" — Miscellaneous operating system interfaces*. 2023. URL: <https://docs.python.org/3/library/signal.html?highlight=signal#module-signal>.
- [38] Python. *"signal" - Set handlers for asynchronous events*. URL: <https://docs.python.org/3/library/os.html>.
- [39] Rezaur. Rahman and SpringerLink (Online service). *Intel® Xeon Phi™ Coprocessor Architecture and Tools*. Berkeley, CA : Apress : 2013. URL: <https://doi.org/10.1007/978-1-4302-5927-5>.
- [40] Efraim Rotem et al. "Power-Management Architecture of the Intel Microarchitecture Code-Named Sandy Bridge". In: *IEEE Micro* 32.2 (Mar. 2012), pp. 20–27. ISSN: 1937-4143. DOI: 10.1109/MM.2012.12.
- [41] The Parliamentary Office of Science and Technology. *Energy Consumption of ICT*. URL: <https://researchbriefings.files.parliament.uk/documents/POST-PN-0677/POST-PN-0677.pdf>.
- [42] SPEC. *SPECpower and Performance Committee*. URL: <https://www.spec.org/power/>.
- [43] iTeh Standards. *Electrical and electronic household and office equipment - Measurement of low power consumption*. URL: <https://standards.iteh.ai/catalog/standards/clc/371d2d67-a439-4f20-96f0-02675496fd03/en-50564-2011>.

- [44] Uber Technologies. *Horovod documentation*. URL: <https://horovod.readthedocs.io/en/latest/index.html/>.
- [45] ISS Group at the University of Texas. *LonestarGPU*. URL: <https://iss.oden.utexas.edu/?p=projects/galois/lonestargpu>.
- [46] Yokogawa. *WT300E Digital Power Analyzer*. URL: <https://tmi.yokogawa.com/solutions/products/power-analyzers/digital-power-meter-wt300e/#Documents-Downloads>.
- [47] Yokogawa. *WT300E Series Digital Power Meter*. URL: <https://cdn.tmi.yokogawa.com/1/2562/files/BUWT300E-01EN.pdf>.

LIST OF FIGURES

3.1. sanna.kask	24
3.2. vinnana.kask	24
3.3. CPUs benchmarks, utilizing the Intel Xeon Silver CPUs	25
3.4. GPUs benchmarks, utilizing the NVIDIA Quadro RTX GPUs	25
3.5. Hybrid configuration, utilizing both CPUs and GPUs	25
3.6. CPU benchmarks implementation used on <i>sanna.kask</i>	25
3.7. GPU benchmarks implementation used on <i>sanna.kask</i>	25
3.8. CPU benchmarks implementation used on <i>vinnana.kask</i>	25
3.9. GPU benchmarks implementation used on <i>vinnana.kask</i>	25
3.10. Benchmarks evaluation example	26
3.11.1 CPU, 1 Thread	28
3.12.1 CPU, 5 Threads	28
3.13.1 CPU, 10 Threads	28
3.14.1 CPU, 20 Threads	28
3.15.2 CPUs, 2 Threads	28
3.16.2 CPUs, 10 Threads	28
3.17.2 CPUs, 20 Threads	28
3.18.2 CPUs, 40 Threads	28
3.19.1 GPU, 1 Thread	29
3.20.2 GPUs, 2 Threads	29
3.21.4 GPUs, 4 Threads	29
3.22.8 GPUs, 8 Threads	29
3.23.1 CPU, 4 Threads + 1 GPU, 1 Thread	29
3.24.1 CPU, 8 Threads + 2 GPUs, 2 Threads	29
3.25.2 CPUs, 16 Threads + 4 GPU, 4 Threads	29
3.26.2 CPUs, 32 Threads + 8 GPUs, 8 Threads	29
4.1. General Flowchart	33
4.2. Processes Flowchart	34

LIST OF TABLES

3.1. Execution times of OMP-CPP benchmarks	22
3.2. Execution times of OMP-CUDA benchmarks	22
3.3. Execution times of MPI-Fortran benchmarks	23
3.4. Execution times of Horovod-Python benchmarks	23
3.5. Overview on components of the servers used in experiments	24
3.6. Overview on implementations used in tests	26
3.7. Overview on benchmarks used in tests on <i>sanna.kask</i>	26
3.8. Overview on benchmarks used in tests on <i>vinnana.kask</i>	26
3.9. Configurations used on node <i>sanna.kask</i>	27
3.10. Configurations used on node <i>vinnana.kask</i>	27
4.1. server: sanna.kask , device: 1 CPU , implementation: OMP-CPP , benchmark: bt.C , data displayed: power draw	43
4.2. server: sanna.kask , device: 1 CPU , implementation: OMP-CPP , benchmark: bt.C , data displayed: energy used	44
4.3. server: sanna.kask , device: 2 CPUs , implementation: OMP-CPP , benchmark: bt.C , data displayed: power draw	45
4.4. server: sanna.kask , device: 2 CPUs , implementation: OMP-CPP , benchmark: bt.C , data displayed: energy used	46
4.5. server: sanna.kask , device: 1 CPU , implementation: OMP-CPP , benchmark: is.D , data displayed: power draw	47
4.6. server: sanna.kask , device: 1 CPU , implementation: OMP-CPP , benchmark: is.D , data displayed: energy used	48
4.7. server: sanna.kask , device: 2 CPUs , implementation: OMP-CPP , benchmark: is.D , data displayed: power draw	49
4.8. server: sanna.kask , device: 2 CPUs , implementation: OMP-CPP , benchmark: is.D , data displayed: energy used	50
4.9. server: sanna.kask , device: 1 CPU , implementation: OMP-CPP , benchmark: lu.C , data displayed: power draw	51
4.10. server: sanna.kask , device: 1 CPU , implementation: OMP-CPP , benchmark: lu.C , data displayed: energy used	52
4.11. server: sanna.kask , device: 2 CPUs , implementation: OMP-CPP , benchmark: lu.C , data displayed: power draw	53
4.12. server: sanna.kask , device: 2 CPUs , implementation: OMP-CPP , benchmark: lu.C , data displayed: energy used	54
4.13. server: sanna.kask , device: GPUs , implementation: OMP-CUDA , benchmark: ep.D , data displayed: power draw	55
4.14. server: sanna.kask , device: GPUs , implementation: OMP-CUDA , benchmark: ep.D , data displayed: energy used	56
4.15. server: sanna.kask , device: GPUs , implementation: OMP-CUDA , benchmark: lu.D , data displayed: power draw	57
4.16. server: sanna.kask , device: GPUs , implementation: OMP-CUDA , benchmark: lu.D , data displayed: energy used	58

4.17.server: sanna.kask , device: GPUs , implementation: OMP-CUDA , benchmark: sp.D , data displayed: power draw	59
4.18.server: sanna.kask , device: GPUs , implementation: OMP-CUDA , benchmark: sp.D , data displayed: energy used	60
4.19.server: sanna.kask , device: Hybrid , implementation: OMP-CPP+OMP-CUDA , benchmark: bt.C+lu.D , data displayed: power draw	61
4.20.server: sanna.kask , device: Hybrid , implementation: OMP-CPP+OMP-CUDA , benchmark: bt.C+lu.D , data displayed: energy used	62
4.21.server: sanna.kask , device: Hybrid , implementation: OMP-CPP+OMP-CUDA , benchmark: is.D+sp.D , data displayed: power draw	63
4.22.server: sanna.kask , device: Hybrid , implementation: OMP-CPP+OMP-CUDA , benchmark: is.D+sp.D , data displayed: energy used	64
4.23.server: sanna.kask , device: Hybrid , implementation: OMP-CPP+OMP-CUDA , benchmark: lu.C+ep.D , data displayed: power draw	65
4.24.server: sanna.kask , device: Hybrid , implementation: OMP-CPP+OMP-CUDA , benchmark: lu.C+ep.D , data displayed: energy used	66
4.25.server: vinnana.kask , device: CPUs , implementation: MPI-Fortran , benchmark: ep.D.x , data displayed: power draw	67
4.26.server: vinnana.kask , device: CPUs , implementation: MPI-Fortran , benchmark: ep.D.x , data displayed: energy used	68
4.27.server: vinnana.kask , device: CPUs , implementation: MPI-Fortran , benchmark: is.D.x , data displayed: power draw	69
4.28.server: vinnana.kask , device: CPUs , implementation: MPI-Fortran , benchmark: is.D.x , data displayed: energy used	70
4.29.server: vinnana.kask , device: CPUs , implementation: MPI-Fortran , benchmark: lu.C.x , data displayed: power draw	71
4.30.server: vinnana.kask , device: CPUs , implementation: MPI-Fortran , benchmark: lu.C.x , data displayed: energy used	72
4.31.server: vinnana.kask , device: GPUs , implementation: Horovod-Python , bench- mark: Xception , data displayed: power draw	73
4.32.server: vinnana.kask , device: GPUs , implementation: Horovod-Python , bench- mark: Xception , data displayed: energy used	74
4.33.server: vinnana.kask , device: Hybrid , implementation: MPI-Fortran+Horovod- Python , benchmark: ep.D.x+Xception , data displayed: power draw	75
4.34.server: vinnana.kask , device: Hybrid , implementation: MPI-Fortran+Horovod- Python , benchmark: ep.D.x+Xception , data displayed: energy used	76
4.35.server: vinnana.kask , device: Hybrid , implementation: MPI-Fortran+Horovod- Python , benchmark: is.D.x+Xception , data displayed: power draw	77
4.36.server: vinnana.kask , device: Hybrid , implementation: MPI-Fortran+Horovod- Python , benchmark: is.D.x+Xception , data displayed: energy used	78
4.37.server: vinnana.kask , device: Hybrid , implementation: MPI-Fortran+Horovod- Python , benchmark: lu.C.x+Xception , data displayed: power draw	79
4.38.server: vinnana.kask , device: Hybrid , implementation: MPI-Fortran+Horovod- Python , benchmark: lu.C.x+Xception , data displayed: energy used	80
4.39.Power draw measurements of idle <i>sanna.kask</i> server	81
4.40.Power draw measurements of idle <i>vinnana.kask</i> server	81

4.41.server: sanna.kask , device: 1 CPU , implementation: OMP-CPP , benchmark: bt.C , data displayed: differences in power draw	82
4.42.server: sanna.kask , device: 1 CPU , implementation: OMP-CPP , benchmark: is.D , data displayed: differences in power draw	82
4.43.server: sanna.kask , device: 1 CPU , implementation: OMP-CPP , benchmark: lu.C , data displayed: differences in power draw	82
4.44.server: sanna.kask , device: 2 CPUs , implementation: OMP-CPP , benchmark: bt.C , data displayed: differences in power draw	83
4.45.server: sanna.kask , device: 2 CPUs , implementation: OMP-CPP , benchmark: is.D , data displayed: differences in power draw	83
4.46.server: sanna.kask , device: 2 CPUs , implementation: OMP-CPP , benchmark: lu.C , data displayed: differences in power draw	83
4.47.server: sanna.kask , device: GPUs , implementation: OMP-CUDA , benchmark: ep.D , data displayed: differences in power draw	84
4.48.server: sanna.kask , device: GPUs , implementation: OMP-CUDA , benchmark: lu.D , data displayed: differences in power draw	84
4.49.server: sanna.kask , device: GPUs , implementation: OMP-CUDA , benchmark: sp.D , data displayed: differences in power draw	84
4.50.server: sanna.kask , device: Hybrid , implementation: OMP-CPP+OMP-CUDA , benchmarks: bt.C+lu.D , data displayed: differences in power draw	85
4.51.server: sanna.kask , device: Hybrid , implementation: OMP-CPP+OMP-CUDA , benchmarks: is.D+sp.D , data displayed: differences in power draw	85
4.52.server: sanna.kask , device: Hybrid , implementation: OMP-CPP+OMP-CUDA , benchmarks: lu.C+ep.D , data displayed: differences in power draw	85
4.53.server: vinnana.kask , device: CPUs , implementation: MPI-Fortran , benchmark: ep.D.x , data displayed: differences in power draw	86
4.54.server: vinnana.kask , device: CPUs , implementation: MPI-Fortran , benchmark: is.D.x , data displayed: differences in power draw	86
4.55.server: vinnana.kask , device: CPUs , implementation: MPI-Fortran , benchmark: lu.C.x , data displayed: differences in power draw	86
4.56.server: vinnana.kask , device: GPUs , implementation: Horovod-Python , bench- mark: Xception , data displayed: differences in power draw	87
4.57.server: vinnana.kask , device: Hybrid , implementation: MPI-Fortran+Horovod- Python , benchmark: ep.D.x+Xception , data displayed: differences in power draw	87
4.58.server: vinnana.kask , device: Hybrid , implementation: MPI-Fortran+Horovod- Python , benchmark: is.D.x+Xception , data displayed: differences in power draw	87
4.59.server: vinnana.kask , device: Hybrid , implementation: MPI-Fortran+Horovod- Python , benchmark: lu.C.x+Xception , data displayed: differences in power draw	87